

Everything You Always Wanted to Know about Synchronization but Were Afraid to Ask

Tudor David, Rachid Guerraoui, Vasileios Trigonakis

The multi-core revolution



- **Big challenges in hardware & software**
- Software
 - **scalability**: \uparrow performance by \uparrow number of cores

Synchronization is one of the biggest scalability bottlenecks

Synchronization

- Cannot always be avoided
 - not all applications embarrassingly parallel
- Synchronization is just an overhead
 - but guarantees correctness
- Scalability of synchronization
 - do not ↓ performance as the number of cores ↑



Scalability of synchronization is key to application scalability

Synchronization is difficult

- Tons of work
 - design of synchronization schemes [ISCA'89, TPDS'90, ASPLOS'91, TOCS'91, PPOPP'01, PODC'95, ICPP'06, SPAA'10, IPDPS'11, PPOPP'12, ATC'12, ...]
 - fix synchronization bottlenecks [SOSP'89, HPCA'07, OSDI'99, OSDI'08, SOSP'09, APLOS'09, OSR'09, OSDI'10, ...]
- Scalability issues?
 - hardware
 - usage of specific atomic operations
 - synchronization algorithm
 - application context
 - workload



Limited understanding of the behavior of synchronization

Take a step back

and perform a thorough analysis of synchronization on modern hardware

What is the main source of scalability problems in synchronization?

Answer

Scalability of synchronization is mainly a property of the hardware

Key observations


1. Crossing sockets is a killer



2. Sharing within a socket is necessary but not sufficient



3. Intra-socket uniformity matters



4. Loads & stores can be as expensive as atomic operations
5. Simple locks are powerful

...

Disclaimer ☺

We do not claim

“Bad synchronization” in software
will scale **well** due to hardware



We claim


“Good synchronization” in software
might **not scale as expected** due to hardware



Analysis method


Hardware processors

Multi-sockets



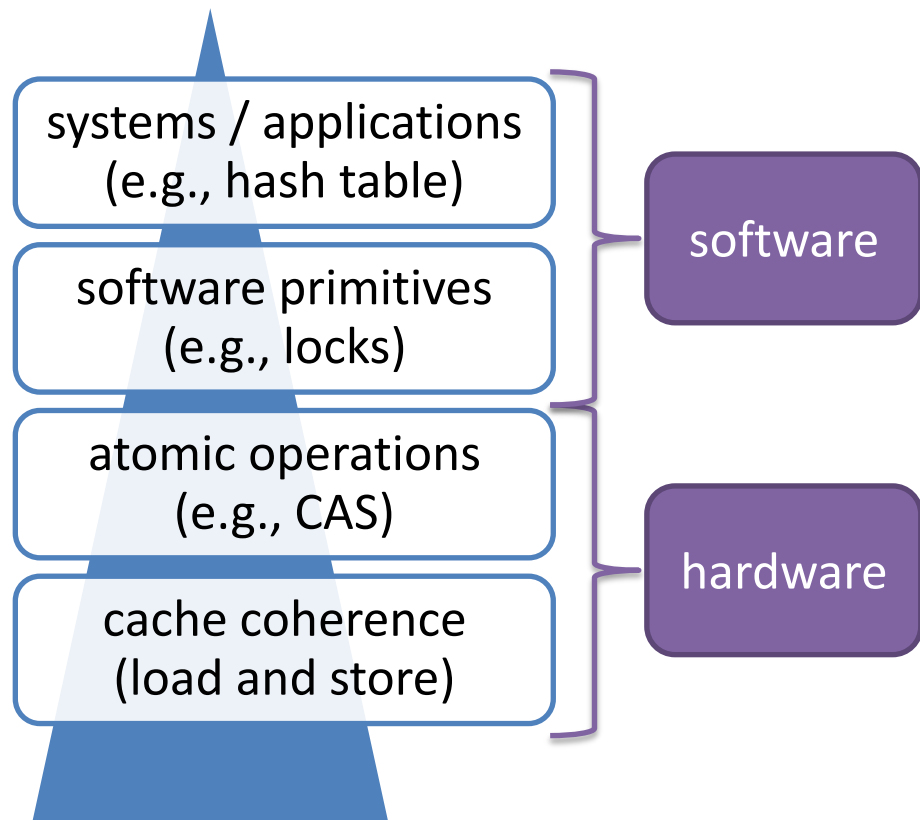
- AMD Opteron (4x 6172 - 48 cores)
- Intel Xeon (8x E7-8867L - 80 cores)

Single-sockets



- Sun Niagara 2 (8 cores)
- Tiler TILE-Gx36 (36 cores)

Synchronization layers



Outline

1. Crossing sockets
2. Sharing within a socket
3. Intra-socket uniformity
4. Atomic operations
5. Simple locks

applications

primitives

atomic ops

cache coherence

Distance on multi-sockets


applications

primitives

atomic ops


cache
coherence

Opteron



- Within socket: 40 ns
- Per hop: +40 ns
- Up to **3x** more

Xeon



- Within socket: 20 – 40 ns
- Per hop: +50 ns
- Up to **8x** more

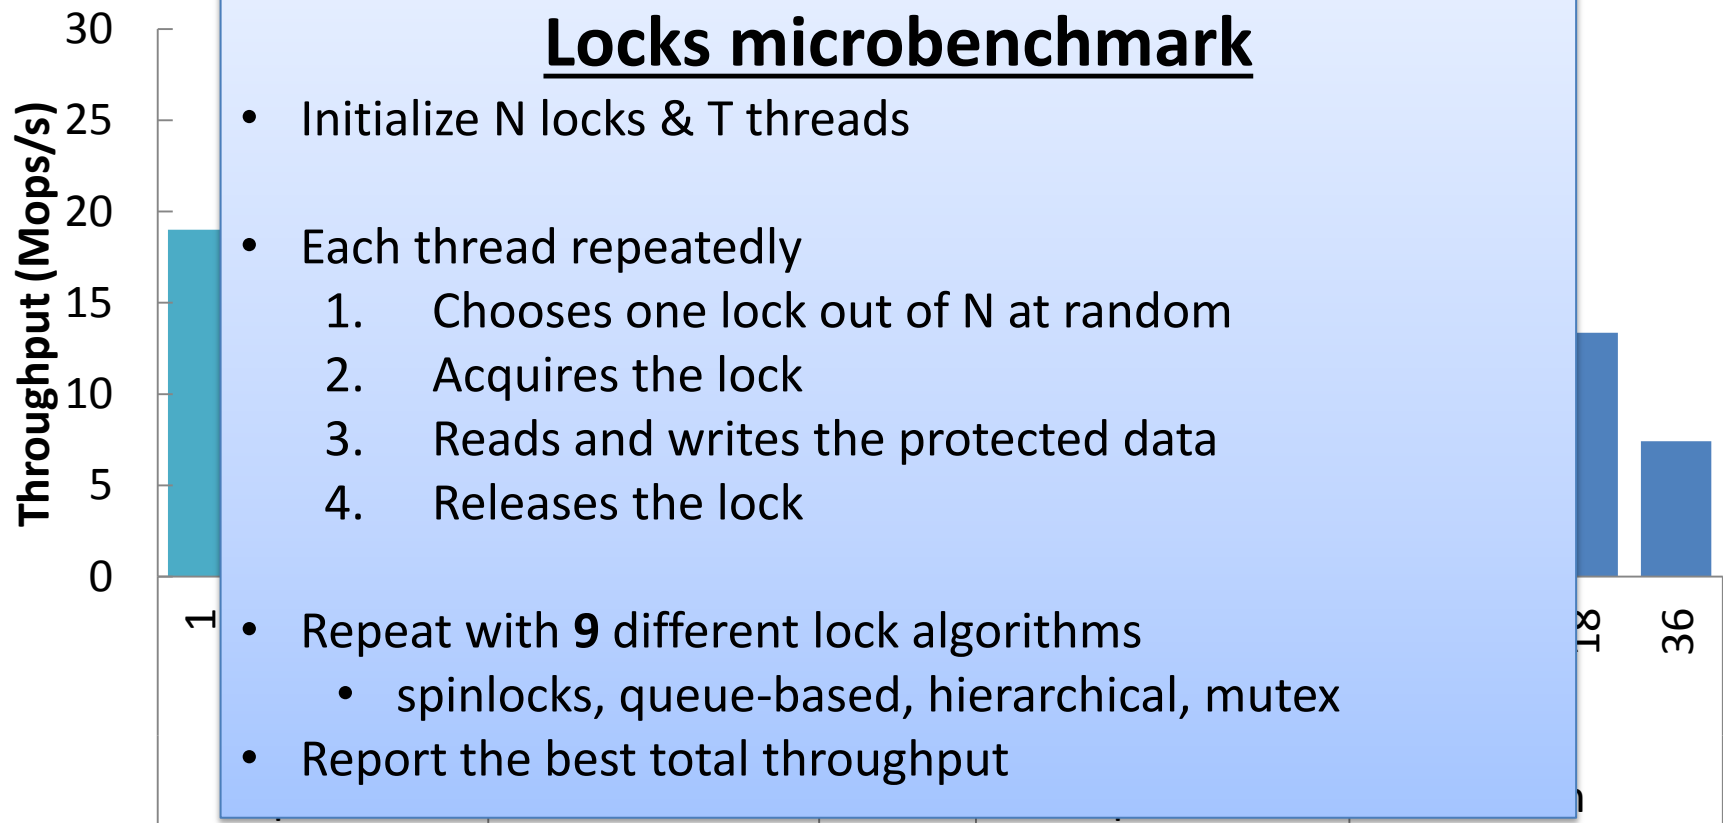
Crossing sockets is a killer: up to 8x more expensive

Locks on multi-sockets

**** Each point is the best result out of 9 lock algorithms**

High contention (4 locks)

Low contention (512 locks)



Locks on multi-sockets

applications

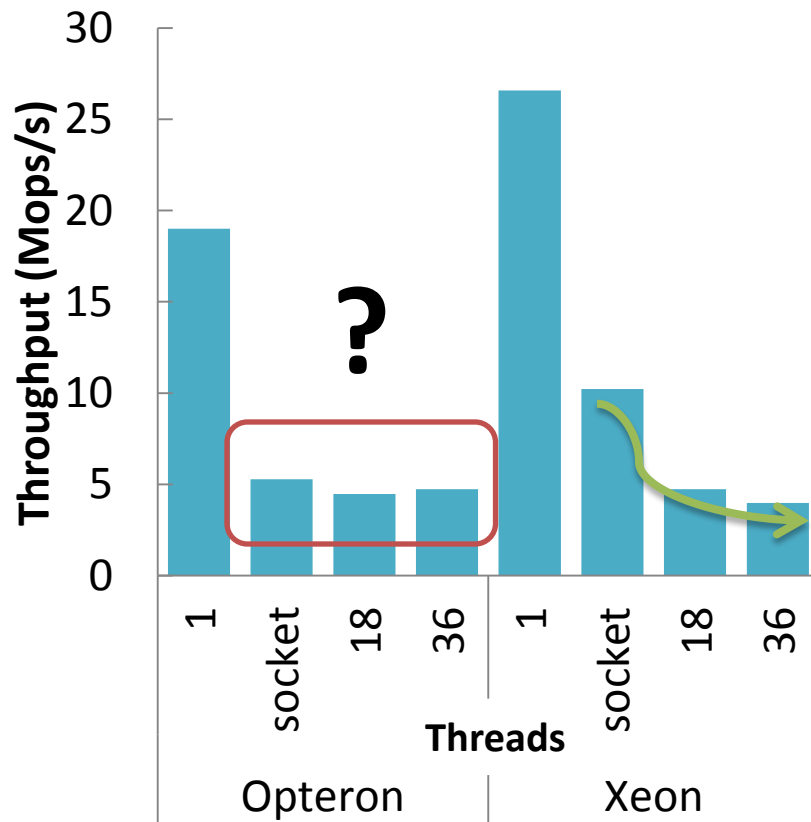
primitives

atomic ops

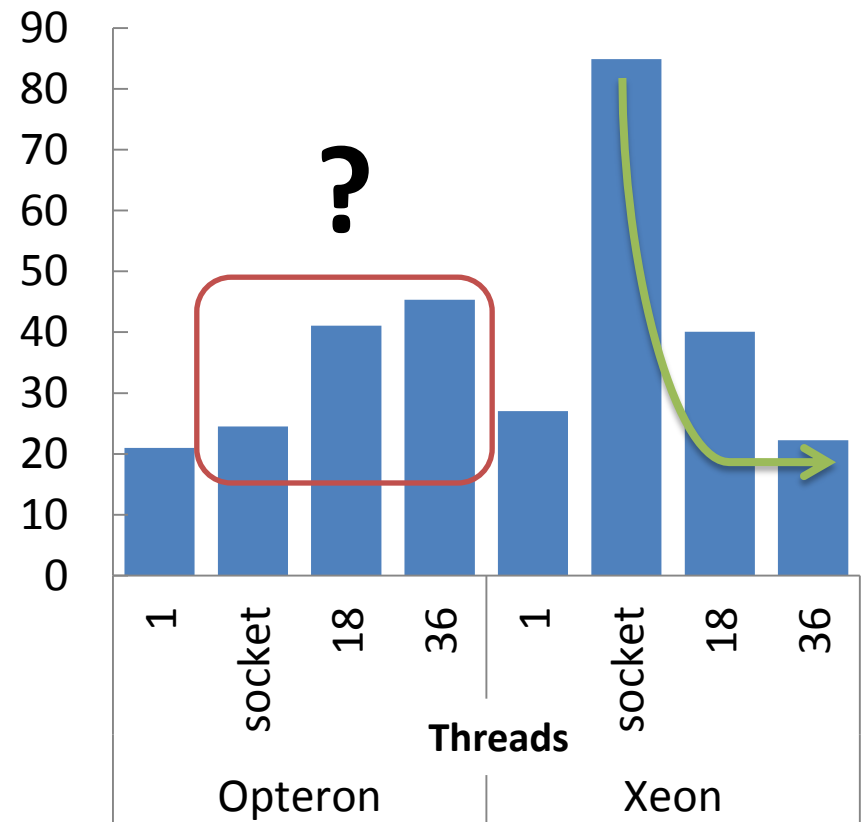
cache
coherence

** Each point is the best result out of 9 lock algorithms

High contention (4 locks)



Low contention (512 locks)

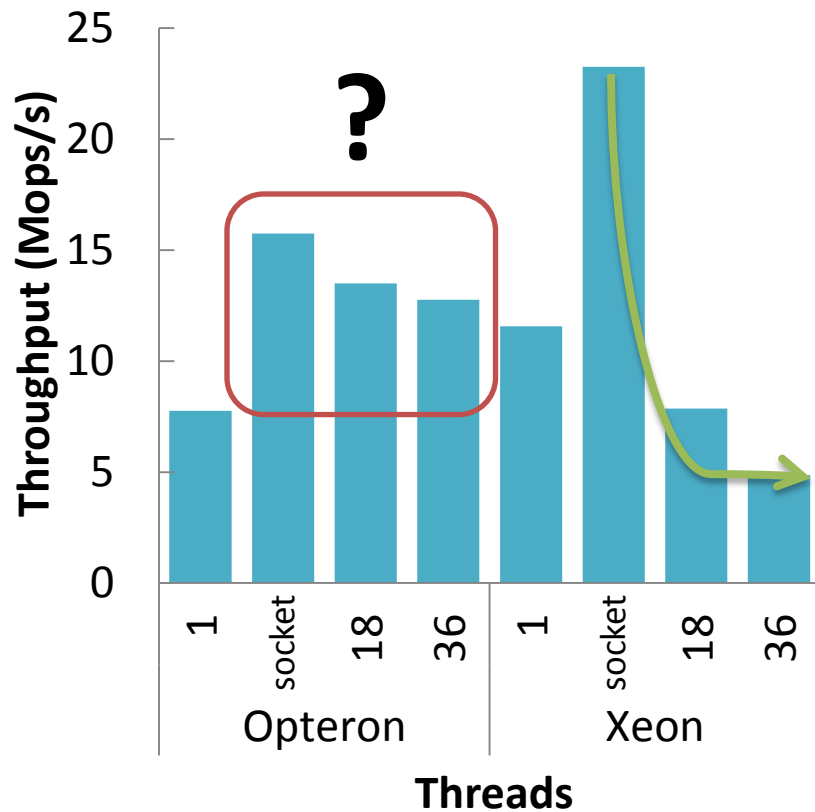


Crossing sockets is a killer: big decrease in performance

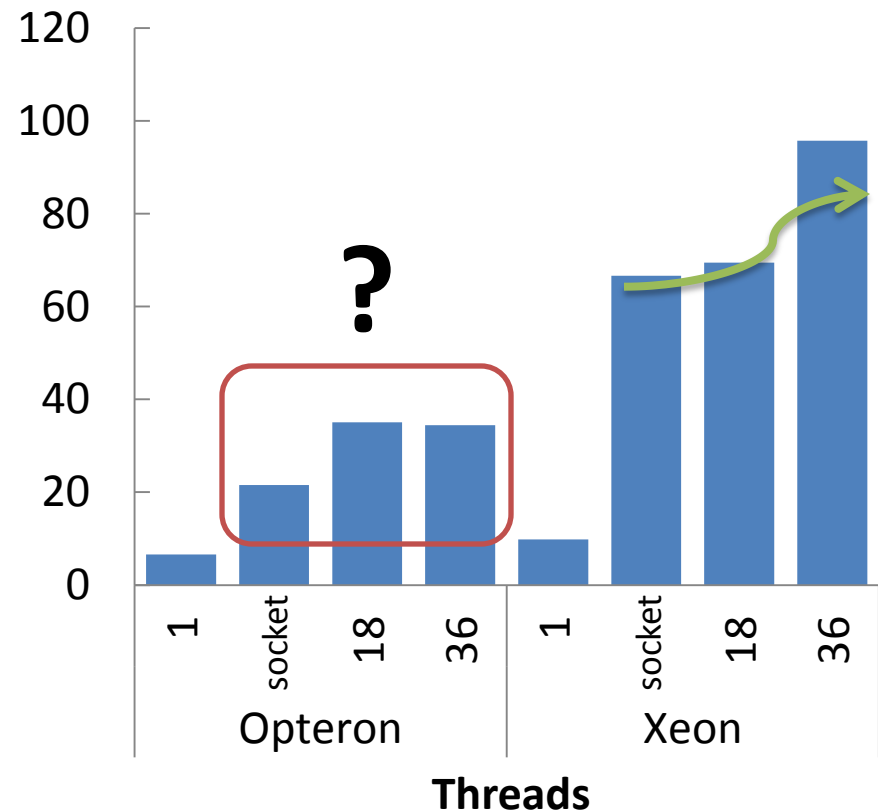
Hash table on multi-sockets

** Each point is the best result taken by any out of 9 lock algorithms

High contention (12 buckets)



Low contention (512 buckets)



Crossing sockets is a killer

Outline

1. Crossing sockets
- 2. Sharing within a socket**
3. Intra-socket uniformity
4. Atomic operations
5. Simple locks

applications

primitives

atomic ops

cache coherence

Coherence on multi-sockets

applications


primitives

atomic ops

cache


coherence

Opteron



Incomplete directory

Xeon



Broadcast requests

Locality on multi-sockets


applications

primitives

atomic ops


cache
coherence

Opteron



- **Within socket: 40 ns**
- Data within a socket
 - served locally (40 ns)
 - broadcast (**120 ns**)

Xeon



- **Within socket: 20 – 40 ns**
- Data within a socket
 - served by the LLC (20 – 40 ns)

Sharing within a socket is necessary but not sufficient

Locks on multi-sockets

applications

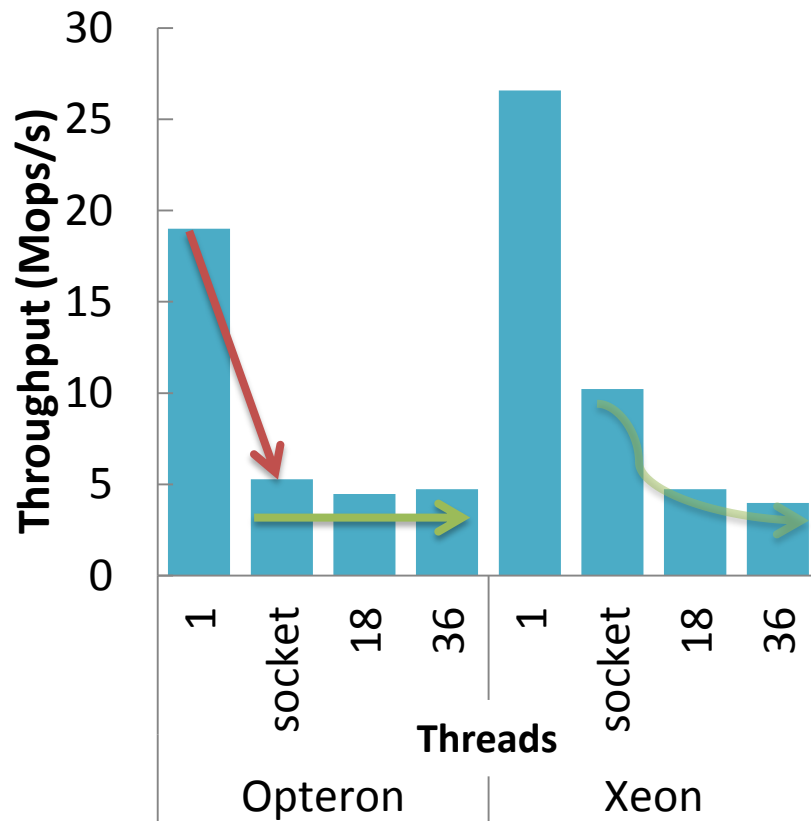
primitives

atomic ops

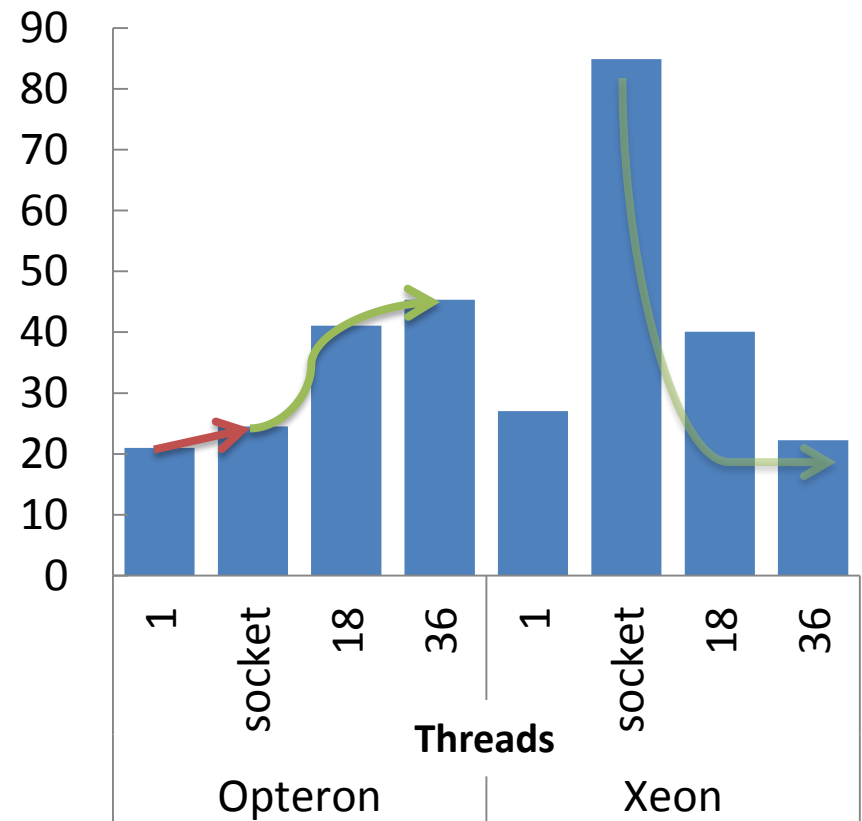
cache

coherence

High contention (4 locks)



Low contention (512 locks)



Sharing within a socket is not sufficient

applications

primitives

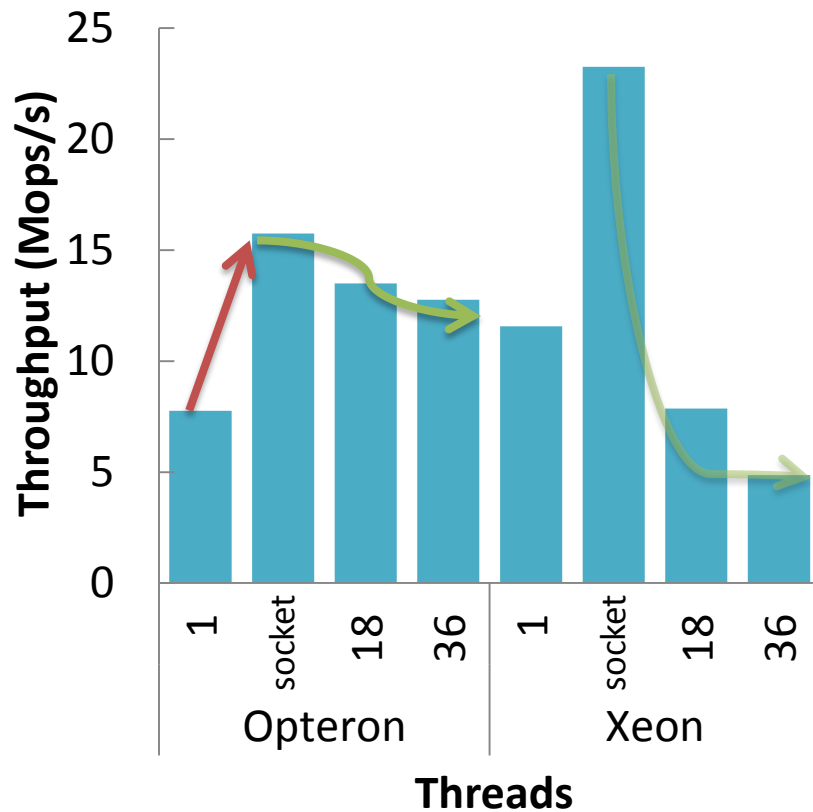
atomic ops

cache


coherence

Hash table on multi-sockets

High contention (12 buckets)



Low contention (512 buckets)



Sharing within a socket is necessary but not sufficient

Outline

1. Crossing sockets
2. Sharing within a socket
- 3. Intra-socket uniformity**
4. Atomic operations
5. Simple locks

applications

primitives

atomic ops

cache coherence

applications


primitives

atomic ops

cache
coherence

Distance on single-sockets

Niagara



- Uniform: 23 ns

Tilera



- 1 hop: 40 ns
- Per hop: +2 ns
- Up to **0.5x** more

Uniformity is expected to scale better

Locks on single-sockets

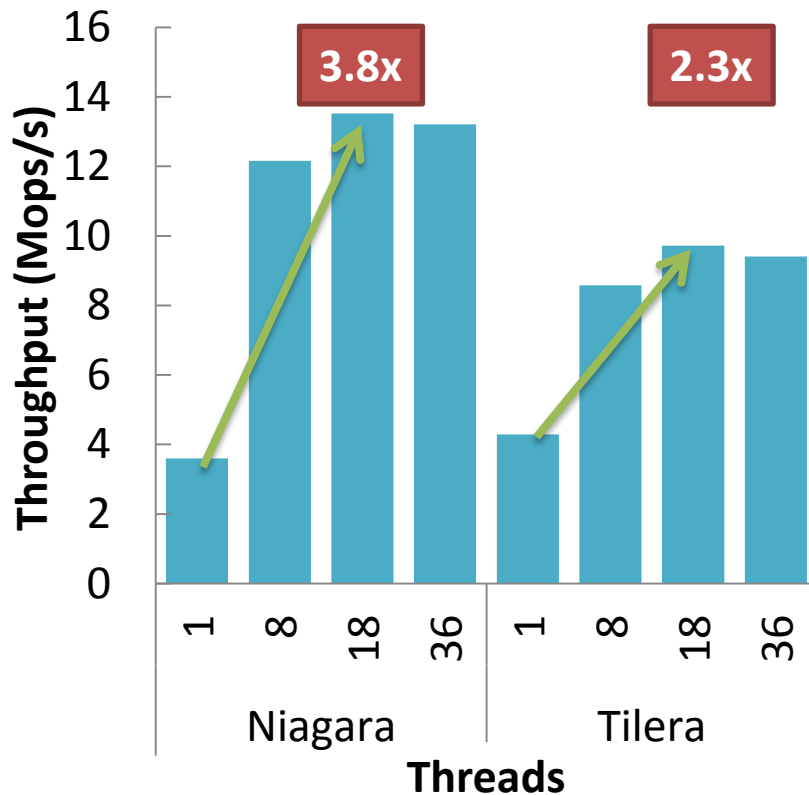
applications

primitives


atomic ops

cache
coherence

High contention (4 locks)



Low contention (512 locks)



Uniformity leads to up to 70% higher scalability

applications

primitives

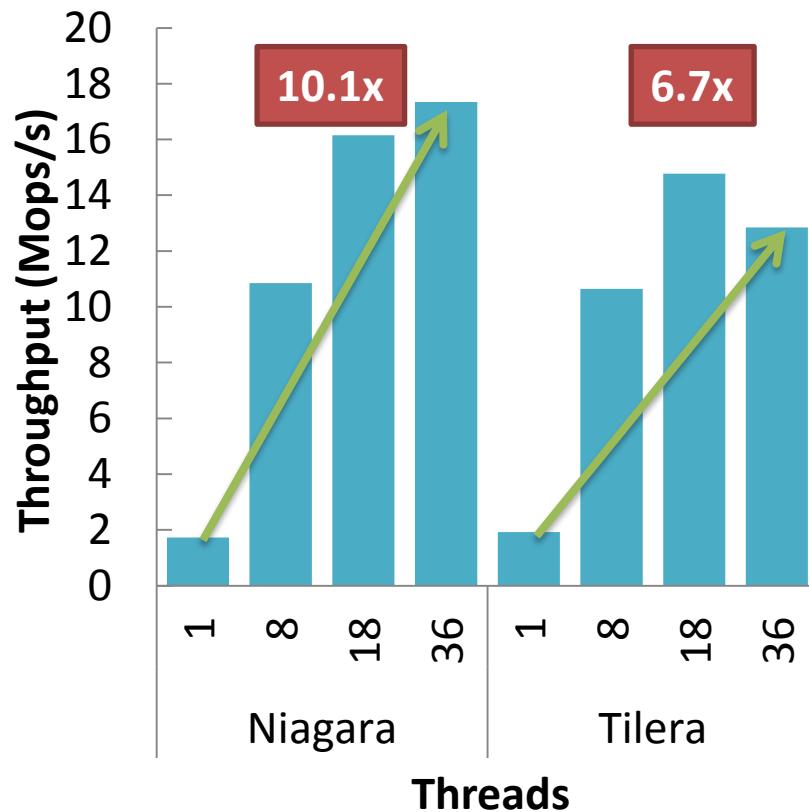
atomic ops

cache

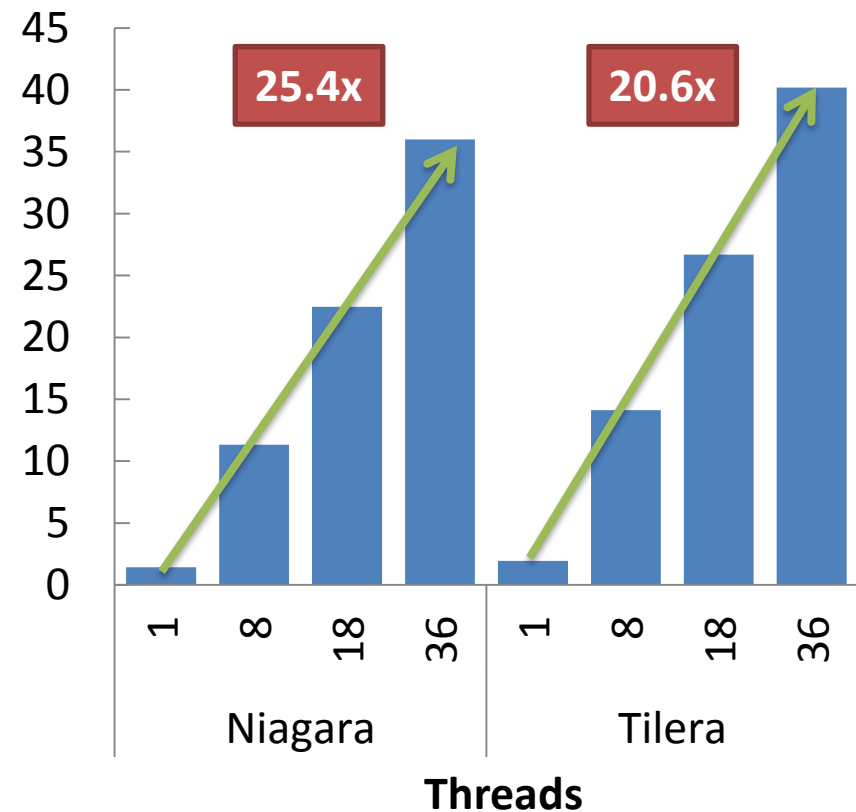
coherence

Hash table on single-sockets

High contention (12 buckets)



Low contention (512 buckets)



Uniformity leads to up to 50% higher scalability

Outline

1. Crossing sockets
2. Sharing within a socket
3. Intra-socket uniformity
- 4. Atomic operations**
5. Simple locks

applications

primitives

atomic ops

cache coherence

Atomic ops on local data

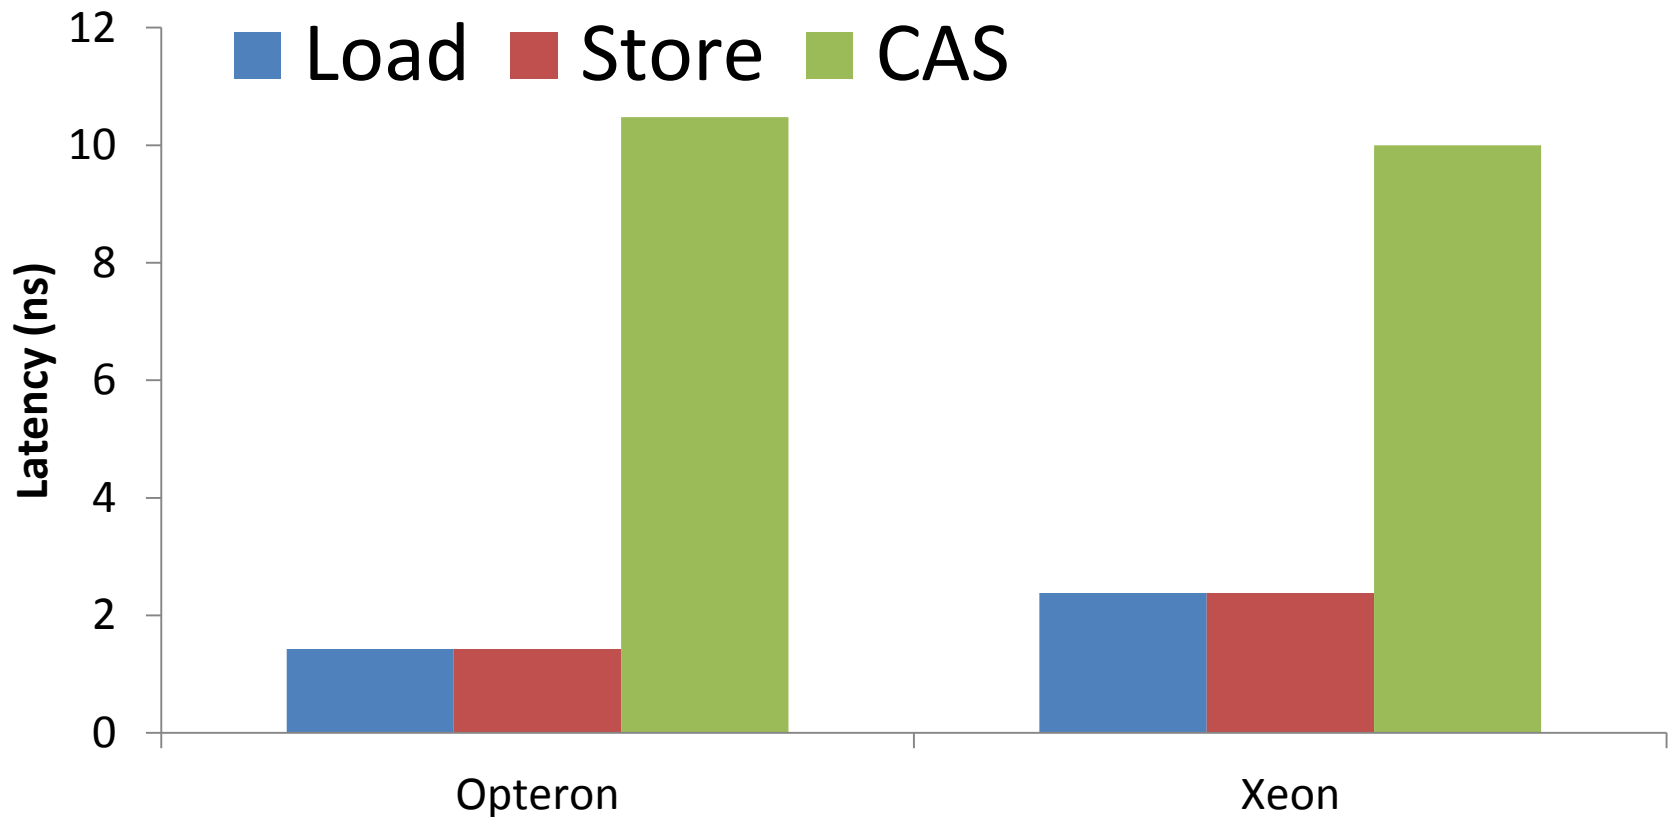
applications

primitives

atomic ops

cache

coherence



CAS: an order of magnitude more expensive on local data

Atomic ops on multi-sockets


applications

primitives

atomic ops

cache

coherence



Loads and stores can be as expensive as atomic operations

Outline

1. Crossing sockets
2. Sharing within a socket
3. Intra-socket uniformity
4. Atomic operations
- 5. Simple locks**

applications

primitives

atomic ops

cache coherence

Hash table – best locks

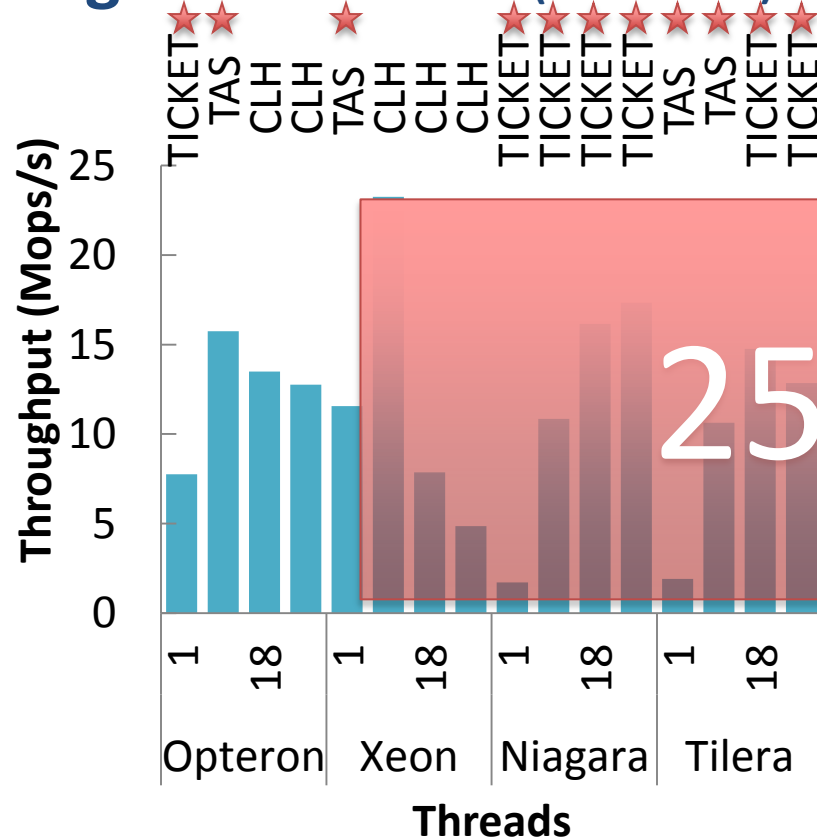
applications

primitives

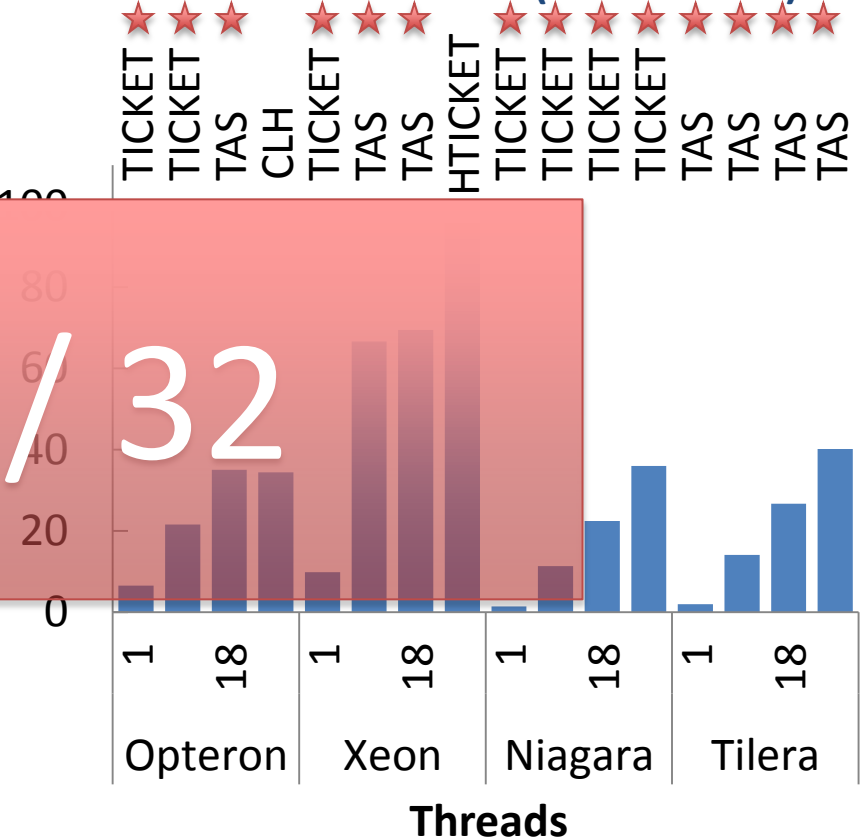
atomic ops

cache
coherence

High contention (12 buckets)



Low contention (512 buckets)



Simple locks are powerful

Lessons learned

1. Crossing sockets is a killer

→ up to 8x more expensive communication

2. Sharing within a socket is necessary but not sufficient

→ up to 3x more expensive communication

3. Intra-socket uniformity matters

→ up to 70% higher scalability

4. Loads & stores can be as expensive as atomic operations

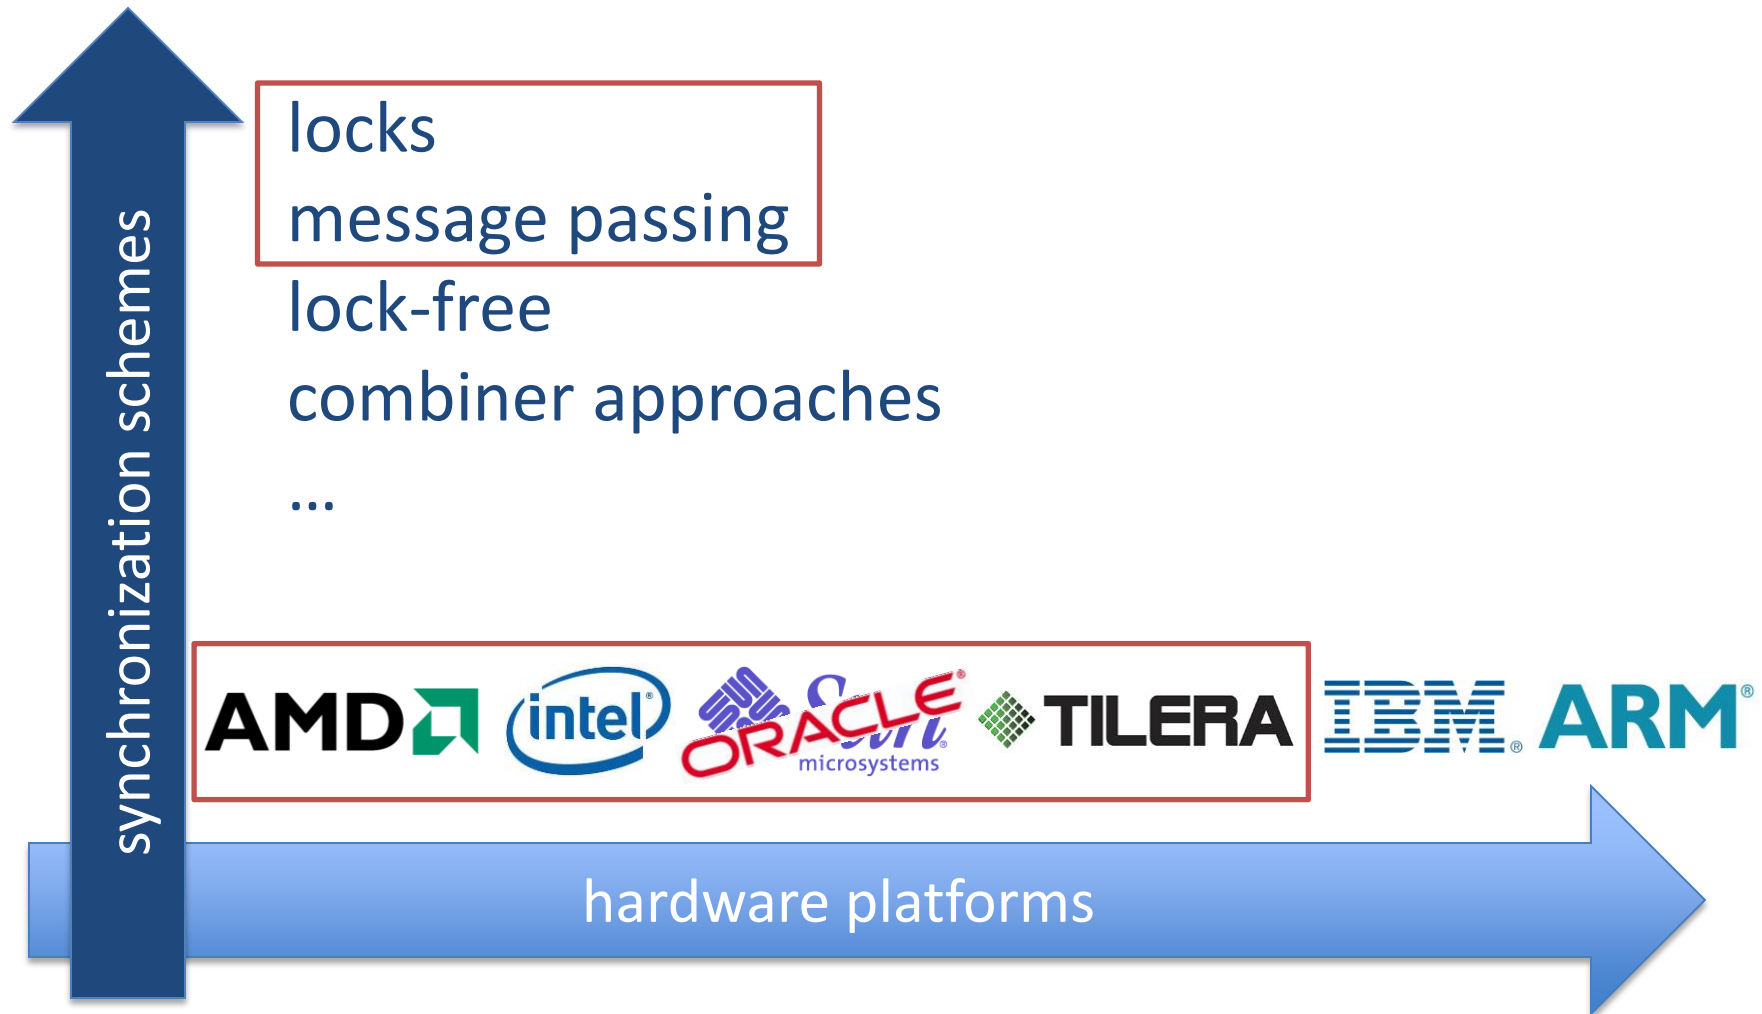
→ 8 - 35% more expensive on non-locally cached data

5. Simple locks are powerful

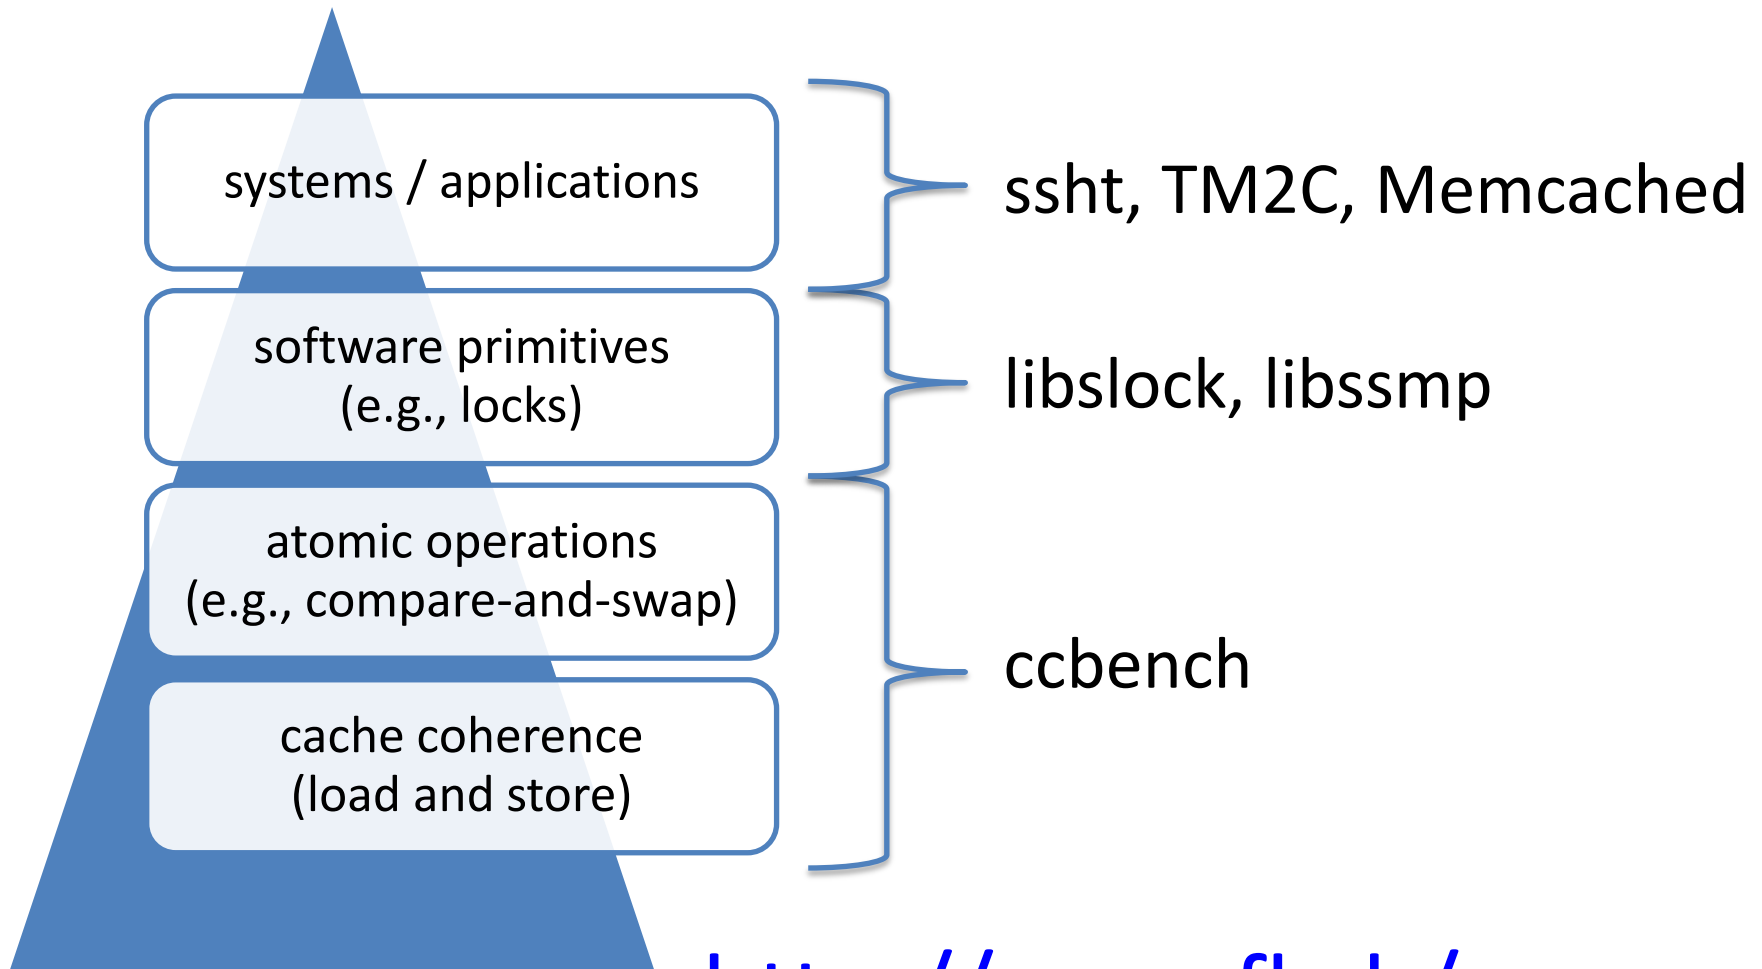
→ better in 25 out of 32 data-points on a hash table

Scalability of synchronization is mainly a property of the hardware

Analysis' space & limitations



SSYNC synchronization suite



<http://go.epfl.ch/ssync>

Thanks!

1. Crossing sockets is a killer
2. Sharing within a socket is necessary but not sufficient
3. Intra-socket uniformity matters
4. Loads & stores can be as expensive as atomic operations
5. Simple locks are powerful

Scalability of synchronization is mainly
a property of the hardware

<http://go.epfl.ch/ssync>