

Materials supplied by Microsoft Corporation may be used for internal review, analysis or research only. Any editing, reproduction, publication, rebroadcast, public showing, internet or public display is forbidden and may violate copyright law.

Secure Virtual Architecture: A Novel Foundation for Operating System Security

Vikram Adve

University of Illinois at Urbana-Champaign

Joint work with:

*John Criswell, Dinakar Dhurjati, Sumant Kowshik,
Andrew Lenharth, Balpreet Pankaj*

Thanks: NSF, SRC, DARPA, Motorola, Apple

Acknowledgements

Funding

- NSF (several programs), SRC, DARPA, Motorola, Apple

LLVM, pointer analysis (DSA), Automatic Pool Allocation

- Chris Lattner
- Other LLVM developers, past and present

SVA PrivOps design “consultants”

- Pierre Salverda
- David Raila

Other input and feedback

- Sam King, Roy Campbell
- Many reviewers

The Context - 1

Increasing threats to system software

US-CERT: 5198 O.S. vulnerabilities reported in 2005^(A)

- 812 on Windows, 2328 on Unix/Linux, 2058 on multiple systems

Month of Kernel Bugs (Nov. 2006)

- One new bug *every day*
- Linux: **8**, MacOS: **8**, FreeBSD: **2**, Solaris: **1**, Windows: **1**
- Code injection: **13**, DOS: **13**, Memory corruption: **4**

Situation could get dramatically worse

- Incentives are increasing: "*botnets*," online banking, identity theft
- Generations entering college are *much* more computer-savvy

The Context - 2

Vast majority of security-critical software is in C/C++

Existing Software

- “Commodity” OS: Windows, Linux, MacOS, BSD family, ...
- Special-purpose OS: Secure64, QNX, ...
- OS services: sshd, xinetd, named, sendmail, ...
- Servers: Apache, **FIND OTHER NAMES ...**

Existing Solutions

- Only ad-hoc, partial solutions are used in production systems
- E.g., non-executable stack/data, SFI, StackGuard, interface annotations, etc.
- Linux kernel has **2.5M** lines of code in ring 0, Windows has **5M** ^(B)

OS on a Safe Runtime: Vision

A complete OS in a safe execution environment

Eliminates an important class of security holes

Improves system reliability

Enables novel OS design techniques

- Application-specific extensions
- Run entire user processes in kernel address space
- Typed inter-process communication
- First-class HLL virtual machines
- Well-defined multithreading
- Novel compiler+run-time solutions to high-level security problems
- ...

OS on a Safe Runtime: So Far

To date, all rely on a **safe programming language**:

- Genera and others: **LISP**
- SPIN: **Modula 3**
- JavaOS, KaffeOS, ... : **Java**
- Singularity: **C#**
- ...

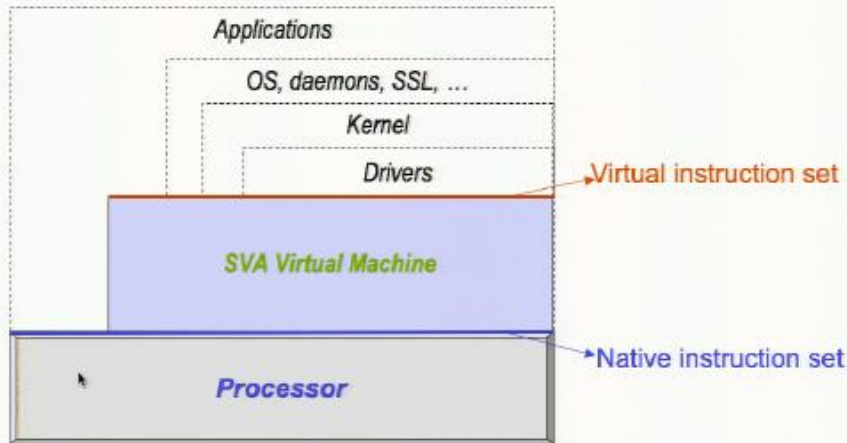
These projects led to many novel OS design techniques.

But what about today's commodity kernels?

Linux, MacOS, BSD family, Windows, ...

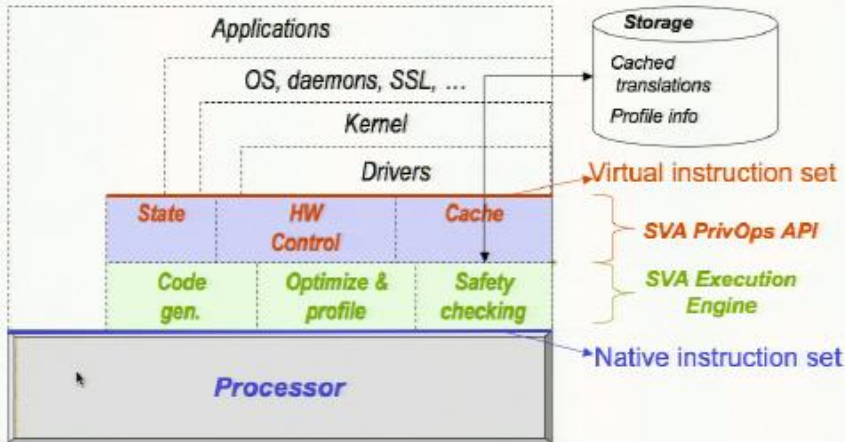
SVA: Secure Virtual Architecture

Designed to work with legacy kernels, e.g., Linux



SVA: Secure Virtual Architecture

Designed to work with legacy kernels, e.g., Linux



Categories of Security Problems

Solved with a safe
execution env.

Buffer overflows

Dangling pointers

Format string errors

Uninitialized pointers

Categories of Security Problems

New solutions enabled by a compiler-based VM

Excessive s/w privilege
Application data secrecy
Detecting rootkits
Information flow
Incorrect security checks
DOS by resource exhaustion
Dynamic code insertion
Race conditions

...

Categories of Security Problems

Need application-level solutions

Configuration errors
Excessive admin
privilege
Cross-site scripting
Network-level DOS
Email worms
Phishing attacks
Spam

...

Categories of Security Problems

Solved with a safe
execution env.

Buffer overflows

Dangling pointers

Format string errors

Uninitialized pointers

Categories of Security Problems

New solutions enabled by a compiler-based VM

Excessive s/w privilege
Application data secrecy
Detecting rootkits
Information flow
Incorrect security checks
DOS by resource exhaustion
Dynamic code insertion
Race conditions

...

Categories of Security Problems

Solved with a safe
execution env.

Buffer overflows

Dangling pointers

Format string errors

Uninitialized pointers

Categories of Security Problems

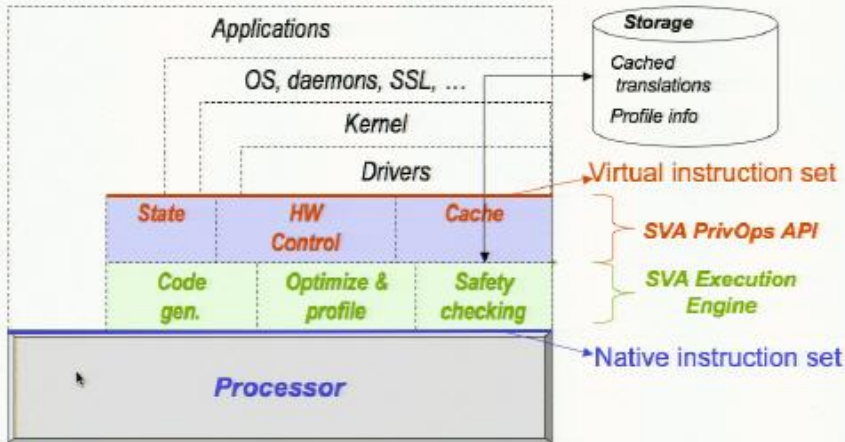
New solutions enabled by a compiler-based VM

Excessive s/w privilege
Application data secrecy
Detecting rootkits
Information flow
Incorrect security checks
DOS by resource exhaustion
Dynamic code insertion
Race conditions

...

SVA: Secure Virtual Architecture

Designed to work with legacy kernels, e.g., Linux



Outline

SVA: Secure Virtual Architecture

- ☛ **SVA Overview**
- ☛ **SVA Virtual Instruction Set and Compiler System**
- ☛ **SVA Safety Guarantees**
- ☛ **Future Work: Security Applications of SVA**


Porting an OS to SVA

E.g., This is how we ported Linux 2.4.22

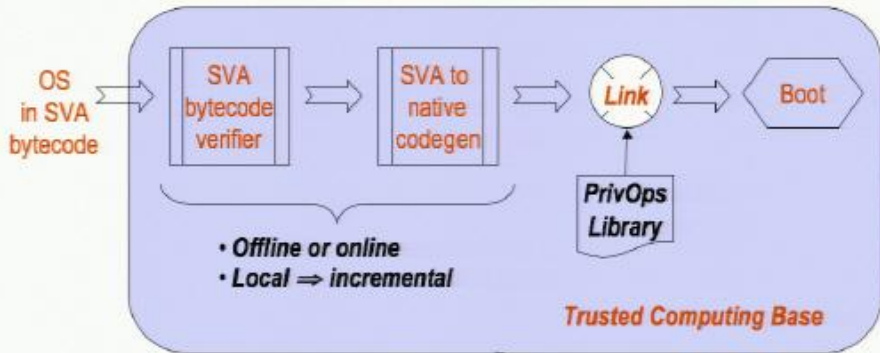


- *Port kernel to PrivOps API*
- *Port kernel allocators to checker API*

Running an OS on SVA



Running an OS on SVA



- TCB components are far simpler than the safety checking compiler
- Incremental ⇒ dynamic loading is easy

SVA Design Features

Comprehensive: All security-sensitive software can be protected

- Kernel, OS extensions, device drivers, daemons, SSL

‡Capability exists but not yet used

SVA Design Features

Comprehensive: All security-sensitive software can be protected

- Kernel, OS extensions, device drivers, daemons, SSL

Incontrovertible: Security criteria cannot be bypassed

- VM and OS can enforce safety policies at install time or load-time

‡Capability exists but not yet used

SVA Design Features

Comprehensive: All security-sensitive software can be protected

- Kernel, OS extensions, device drivers, daemons, SSL

Incontrovertible: Security criteria cannot be bypassed

- VM and OS can enforce safety policies at install time or load-time

Robust: Reduce trusted computing base (TCB):

- Complex safety-checking compiler is outside the TCB
- Extend to other security problems: Information flow, security automata, ...

‡Capability exists but not yet used

SVA Design Features

Comprehensive: All security-sensitive software can be protected

- Kernel, OS extensions, device drivers, daemons, SSL

Incontrovertible: Security criteria cannot be bypassed

- VM and OS can enforce safety policies at install time or load-time

Robust: Reduce trusted computing base (TCB):

- Complex safety-checking compiler is outside the TCB
- Extend to other security problems: Information flow, security automata, ...

Whole-system ‡: Security analysis, transforms across programs

- Can perform analysis, transformations *across application/kernel boundary*



‡Capability exists but not yet used

SVA Design Features

Comprehensive: All security-sensitive software can be protected

- Kernel, OS extensions, device drivers, daemons, SSL

Incontrovertible: Security criteria cannot be bypassed

- VM and OS can enforce safety policies at install time or load-time

Robust: Reduce trusted computing base (TCB):

- Complex safety-checking compiler is outside the TCB
- Extend to other security problems: Information flow, security automata, ...

Whole-system ‡: Security analysis, transforms across programs

- Can perform analysis, transformations *across application/kernel boundary*

Hardware assisted ‡: Can exploit upcoming hardware support

- TPM (secure boot, secure PKI), IOMMU (secure DMA)

‡Capability exists but not yet used

SVA Prototype

SVA PrivOps Library

- C and i386 assembly code for Pentium 3
- Compiled to native code library ahead of time

Safety Principles

- Same safety guarantees for standalone programs and kernel
- Untrusted safety-checking compiler, trusted type checker

Linux 2.4.22 port to SVA

- Like a port to a new architecture
- Assembly code replaced with SVA operations
- ***kmem_cache_alloc***: type homogeneous (TH) pools
- ***kmalloc***, others: non-TH pools

Outline

SVA: The Secure Virtual Architecture Project

☛ SVA Overview

☛ **SVA Virtual Instruction Set**

☛ **SVA Safety Principles**

☛ **Future Security Applications of SVA**

SVA Virtual Instruction Set

Unprivileged operations:

- Derived from *IR* of the LLVM compiler system
- Extended with security annotations:
 - *Current*: type system for safety properties

Privileged operations (PrivOps) API:

- API for kernel-hardware interactions
- Mechanisms, not policy

LLVM Provides Compiler Foundation

Framework for “lifelong compilation”

- Compile-time, link-time, install-time, load-time, run-time, “idle”-time
- IR: Compact, persistent, designed for effective optimization

Commercial-quality compiler infrastructure

- **JIT**: Apple (MacOS 10.5), Adobe, Hue AS
- **Static back end**: Cray, Ageia, Ascenium, Wind River
- **Silicon compilation**: AutoESL
- **Unknown**: Aerospace, Microchip Tech, Wind River
- Also many academic, open-source, users; many contributors

Available at llvm.org .

LLVM IR = Core Unprivileged Operations

```
/* C Source Code */
int SumArray(int Array[],
             int Num)
{
    int i, sum = 0;
    for (i = 0; i < Num; ++i)
        sum += Array[i];
    return sum;
}
```

- Architecture-neutral
- Low-level operations
- SSA representation
- Typed
- Mid-level type info

```
:: SVA Code
int SumArray(int* Array, int Num)
{
bb1:
    %cond = setgt int %Num, 0
    br bool %cond, label %bb2, label %bb3

bb2:
    %sum0 = phi int [%tmp10, %bb2], [0, %bb1]
    %i0 = phi int [%inc, %bb2], [0, %bb1]
    %tmp7 = cast int %i0 to long
    %tmp8 = getelementptr int* Array, long %tmp7
    %tmp9 = load int* %tmp8
    %tmp10 = add int %tmp9, %sum0
    %inc = add int %i0, 1
    %cond2 = setlt int %inc, %Num
    br bool %cond2, label %bb2, label %bb3

bb3:
    %sum1 = phi int [0, %bb1], [%tmp10, %bb2]
    ret int %sum1
}
```

SVA Privileged Operations API

Hardware Control

Syscall, interrupt, trap handlers

- all OS entries are monitored

Page table entries

- so are page mapping events

I/O operations

- ditto

State Manipulation

Context-switching

- save/restore native state

Signal delivery

Interrupt Context

- exploit hardware for fast interrupts

Outline

SVA: The Secure Virtual Architecture Project

- ☛ SVA Overview

- ☛ SVA Virtual Instruction Set

- ☛ **SVA Safety Principles**

 - ☛ Guarantees

 - ☛ Standalone programs

 - ☛ Complete kernel

- ☛ **Future Security Applications of SVA**

Safety Guarantees

Strong Guarantees: Close to a safe language, but not equal

- **Memory safety:** No uninitialized pointer uses, no array overflows
- **Type safety** for a *subset* of objects
- **Control flow integrity:** only follow compiler-predicted paths
- **Sound operational semantics** \Rightarrow sound static analysis

Safety Guarantees

Strong Guarantees: Close to a safe language, but not equal

- **Memory safety:** No uninitialized pointer uses, no array overflows
- **Type safety** for a *subset* of objects
- **Control flow integrity:** only follow compiler-predicted paths
- **Sound operational semantics** \Rightarrow sound static analysis

Primary Weakness: By Design

- **Tolerate but not eliminate dangling pointer errors:** avoid need for GC
- *Tolerate \Rightarrow Do not invalidate the above guarantees!*
- Option: detect *all* dangling pointer uses: low overhead for some programs

Safety Guarantees

Strong Guarantees: Close to a safe language, but not equal

- **Memory safety:** No uninitialized pointer uses, no array overflows
- **Type safety** for a *subset* of objects
- **Control flow integrity:** only follow compiler-predicted paths
- **Sound operational semantics** \Rightarrow sound static analysis

Primary Weakness: By Design

- Tolerate but not eliminate dangling pointer errors: avoid need for GC
- *Tolerate \Rightarrow Do not invalidate the above guarantees!*
- Option: detect *all* dangling pointer uses: low overhead for some programs

Practical

- Completely automatic, no wrappers, no GC
- Works for arbitrary C programs
- Very low overhead for C programs

Underlying Principles

Idea 1: Pool allocation \Leftrightarrow static points-to graph

- Generate “pool checks” on (some) loads, stores
- *Enforces partial type safety, sound pointer analysis*

Idea 2: Exploit type-homogeneous pools

- *Eliminate all checks except array-bounds checks*
- *Dangling pointers harmless (with careful object alignment)*

Idea 3: Efficient array bounds checks without metadata

- Jones & Kelly: Maintain lookup tables of allocated objects
- Use 1 table per pool: greatly reduces overhead
- *Avoid metadata on pointers*

Evaluation: Run-time Overhead

Olden, *Ptrdist*, 3 system
daemons [Full list in PLDI06]

Evaluation: Run-time Overhead

Olden, *Ptrdist*, 3 system
daemons [Full list in PLDI06]

No source changes necessary

Detected all attacks from Zitser's
suite.

Evaluation: Run-time Overhead

1.0 \equiv no pool allocation + no SAFECode passes

Olden, Ptrdist, 3 system
daemons [Full list in PLDI06]

No source changes necessary

Detected all attacks from Zitser's
suite.

| Program | SAFECode: No array checks | SAFECode with array checks |
|---------|------------------------------|-------------------------------|
| bh | 3% | 3% |
| bisort | 0 | -4 |
| em3d | 27 | 91 |
| treeadd | -1 | 0 |
| tsp | -1 | 1 |
| Yacr2 | 30 | 30 |
| Ks | 12 | 12 |
| anagram | 23 | 23 |
| ftpd | 0 | 4 |
| fingerd | 3 | 7 |
| ghttpd | 7 | -10 |

Evaluation: Run-time Overhead

1.0 \equiv no pool allocation + no SAFECode passes

Olden, *Ptrdist*, 3 system
daemons [Full list in PLDI06]

No source changes necessary

Detected all attacks from Zitser's
suite.

Compare CCured (*Olden*, *Ptrdist*)

- Much higher time overheads except em3d
- Much higher space overheads (1x-4x)
- Significant porting effort

| Program | SAFECode: No array checks | SAFECode with array checks |
|---------|---------------------------|----------------------------|
| bh | 3% | 3% |
| bisort | 0 | -4 |
| em3d | 27 | 91 |
| treeadd | -1 | 0 |
| tsp | -1 | 1 |
| Yacr2 | 30 | 30 |
| Ks | 12 | 12 |
| anagram | 23 | 23 |
| ftpd | 0 | 4 |
| fingerd | 3 | 7 |
| ghttpd | 7 | -10 |

Safety Challenges in the OS

Kernels have many custom pool allocators

- *Retain all existing kernel allocators*
- Define clean interface between allocator and checker: needs manual port
- Compiler: Infer mapping of pools to points-to graph nodes

Kernels have many external entry points bringing in pointers

- Can still use array indexing checks for externally-accessible pools


Kernel allocators not initialized early in boot sequence

- Reserve memory to use for metadata for early objects

Extensive use of non-type-safe code

- More aggressive static bounds-checking, run-time tuning

Verifying Security of SVA Bytecode



SVA type system

- Stack of "regions" (logical pools); every object registered in a pool
- Region annotation on every pointer

Bytecode Verifier

- Simple local, type checker
- Inserts run-time checks
- Local \Rightarrow easy to support dynamically loaded modules

This strategy can be used for any type system, e.g., security automata [Walker, POPL2000], information flow [Myers, POPL99]

Outline

SVA: The Secure Virtual Architecture Project

☞ SVA Overview

☞ SVA Virtual Instruction Set

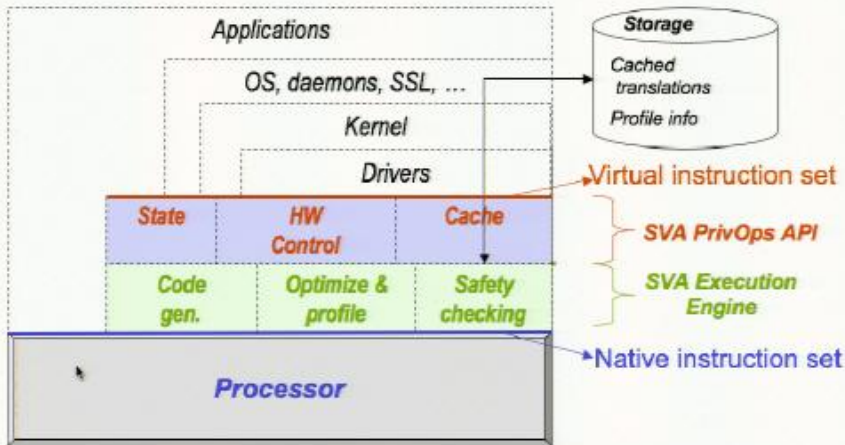
☞ SVA Safety Principles

☞ **Future Security Applications of SVA**

- ☞ Application data secrecy
- ☞ Reducing privilege
- ☞ Monitoring kernel-mode rootkits

SVA: Secure Virtual Architecture

Compiler + privileged run-time + type checker



Some Security Applications of SVA

Minimizing Privilege for Root Programs

- Compiler \Rightarrow Privilege bracketing + secure system call dispatch
- Compiler \Rightarrow "Authenticated system call" policies
- *Incontrovertible*: transformations at install-time or load-time

Protecting Application Data From OS

- Higher privilege VM \Rightarrow Private physical memory
- Higher privilege VM \Rightarrow Monitor all page mappings, I/O operations

Monitoring Kernel-mode Rootkits (with Sam King)

- Higher privilege VM \Rightarrow monitoring of kernel cannot be subverted
- Compiler \Rightarrow fine-grain monitoring of kernel operations

Summary

SVA: Secure Virtual Architecture

Safe environment for programs

- Fully automatic

Safe environment for entire OS

- Low porting effort

Higher-level security capabilities

- Wide-open research area
- Collaborations welcome!

Solved with a safe
execution env.

Buffer overflows

Dangling pointers

Format string
errors

Uninitialized
pointers

Questions?

Summary

SVA: Secure Virtual Architecture

Safe environment for programs

- Fully automatic

Safe environment for entire OS

- Low porting effort

Higher-level security capabilities

- Wide-open research area
- Collaborations welcome!


New solutions enabled by a compiler-based VM

Excessive s/w privilege
Application data secrecy
Detecting rootkits
Information flow
Incorrect security checks
DOS by resource
exhaustion
Dynamic code insertion
Race conditions

...

Questions?

Verifying Security of SVA Bytecode



SVA type system

- Stack of "regions" (logical pools); every object registered in a pool
- Region annotation on every pointer

Bytecode Verifier

- Simple local, type checker
- Inserts run-time checks
- Local \Rightarrow easy to support dynamically loaded modules

This strategy can be used for any type system, e.g., security automata [Walker, POPL2000], information flow [Myers, POPL99]

Summary

SVA: Secure Virtual Architecture

Safe environment for programs

- Fully automatic

Safe environment for entire OS

- Low porting effort

Higher-level security capabilities

- Wide-open research area
- Collaborations welcome!

New solutions enabled by a compiler-based VM

Excessive s/w privilege
Application data secrecy
Detecting rootkits
Information flow
Incorrect security checks
DOS by resource
exhaustion
Dynamic code insertion
Race conditions

...

Questions?