

Materials supplied by Microsoft Corporation may be used for internal review, analysis or research only. Any editing, reproduction, publication, rebroadcast, public showing, internet or public display is forbidden and may violate copyright law.

The Barrelfish OS for scalable and heterogeneous multicore systems

Andrew Baumann



Introducing myself

I'm interested in:

- ▶ Operating systems
- ▶ Distributed/networked systems

I've worked on:

- ▶ Mungi single-address-space OS
- ▶ Tracing and performance monitoring in K42
- ▶ Dynamic update to operating systems
- ▶ OS timer usage study
- ▶ Rhizoma runtime for self-hosting distributed systems
- ▶ Barrelfish OS for heterogeneous multicore systems

Introducing Barrelfish

The Problem

How should we structure an OS for future multicore systems?

Goals:

- ▶ Run a dynamic set of general-purpose applications
- ▶ Reduce code complexity to do this


Challenges:

- ▶ Scalability to many cores
- ▶ Heterogeneity and hardware diversity


System diversity



Sun Niagara T2




AMD Opteron (Istanbul)



Intel Nehalem (Beckton)


The interconnect matters

Today's 8-socket Opteron




The interconnect matters

Tomorrow's 8-socket Nehalem




The interconnect matters

Today's 8-socket Opteron




The interconnect matters

Tomorrow's 8-socket Nehalem



The interconnect matters

On-chip interconnects



Core diversity

- ▶ Within a system:
 - ▶ Programmable NICs
 - ▶ GPUs
 - ▶ FPGAs (in CPU sockets)
- ▶ On a single die:
 - ▶ Performance asymmetry
 - ▶ Streaming instructions (SIMD, SSE, etc.)
 - ▶ Virtualisation support

Your computer is already a distributed system

[Baumann et al., HotOS'09]

Summary:

- ▶ Increasing core counts, increasing diversity
- ▶ Unlike HPC systems, cannot optimise at design time

Your computer is already a distributed system

[Baumann et al., HotOS'09]


Summary:

- ▶ Increasing core counts, increasing diversity
- ▶ Unlike HPC systems, cannot optimise at design time

Two current research examples:


1. Intel SCC (Rock Creek)
2. MSR Beehive

Intel's single-chip cloud computer (SCC)



- ▶ 48 IA-32 cores on one die
- ▶ 2D mesh interconnect
- ▶ Message-passing architecture
- ▶ **No cache coherence**

MSR Beehive



- ▶ Ring interconnect
- ▶ Message passing
- ▶ No cache coherence
- ▶ Split-phase memory access

The multikernel model

[Baumann et al., SOSP'09]

- ▶ It's time to rethink the default structure of an OS
 - ▶ Shared-memory kernel on every core
 - ▶ Data structures protected by locks
 - ▶ Anything else is a device

The multikernel model

[Baumann et al., SOSP'09]

- ▶ It's time to rethink the default structure of an OS
 - ▶ Shared-memory kernel on every core
 - ▶ Data structures protected by locks
 - ▶ Anything else is a device
- ▶ Proposal: **structure the OS as a distributed system**
- ▶ Design principles:
 1. Make inter-core communication explicit
 2. Make OS structure hardware-neutral
 3. View state as replicated

Outline

Introduction

Architecture

Hardware diversity

The multikernel model

Design principles

The model

Barrelfish

Evaluation

Case study: Unmap (TLB shutdown)

Future work

1. Make inter-core communication explicit

- ▶ All communication with messages (no shared state)

1. Make inter-core communication explicit

- ▶ All communication with messages (no shared state)
- ▶ Decouples system structure from inter-core communication mechanism
 - ▶ Communication patterns explicitly expressed
- ▶ Naturally supports heterogeneous cores, non-coherent interconnects (PCIe)
- ▶ Better match for future hardware
 - ▶ ...with explicit message passing (e.g. Tile64, Beehive, SCC)
 - ▶ ...without cache-coherence (e.g. Beehive, SCC)
- ▶ Allows split-phase operations
 - ▶ Decouple requests and responses for concurrency
- ▶ We can reason about it

Message passing vs. shared memory: experiment

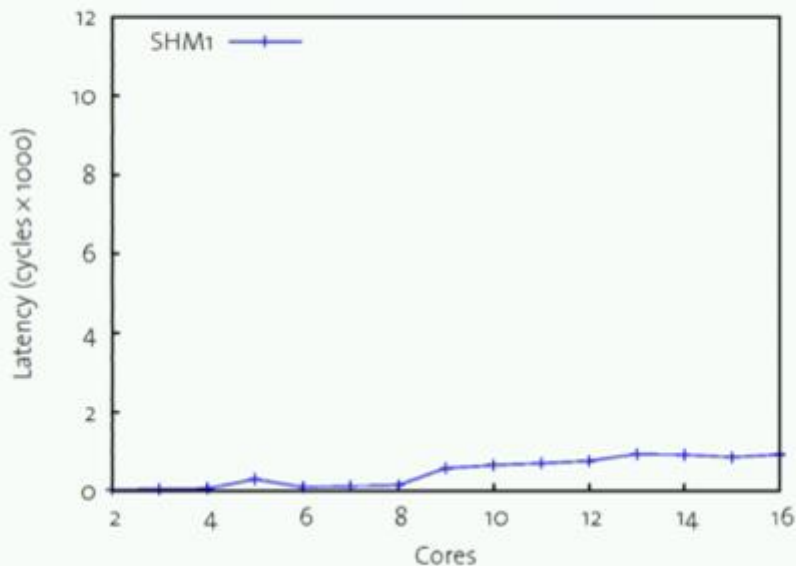
Shared memory (move the data to the operation):

- ▶ Each core updates the same memory locations (no locking)
- ▶ Cache-coherence protocol migrates modified cache lines
 - ▶ **Processor stalled** while line is fetched or invalidated
 - ▶ Limited by **latency** of interconnect round-trips
 - ▶ Performance depends on data size (cache lines) and contention (number of cores)




Shared memory results

4×4-core AMD system



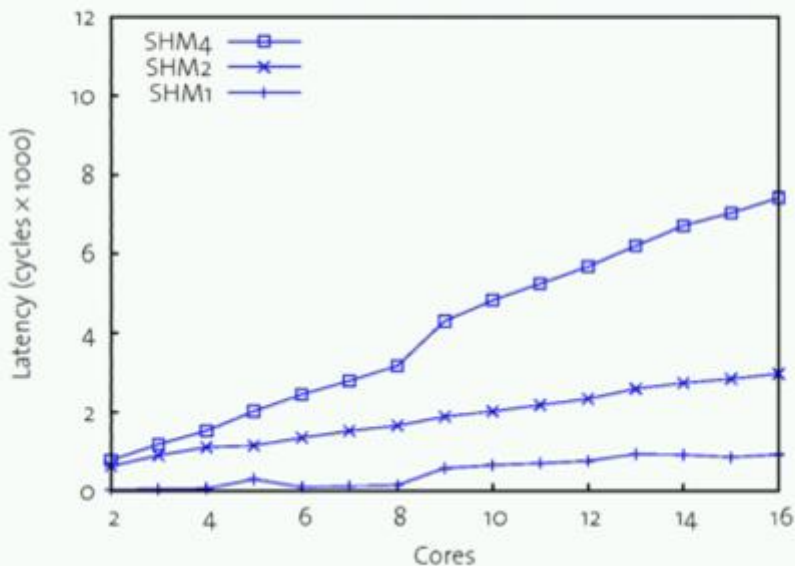
Shared memory results

4x4-core AMD system



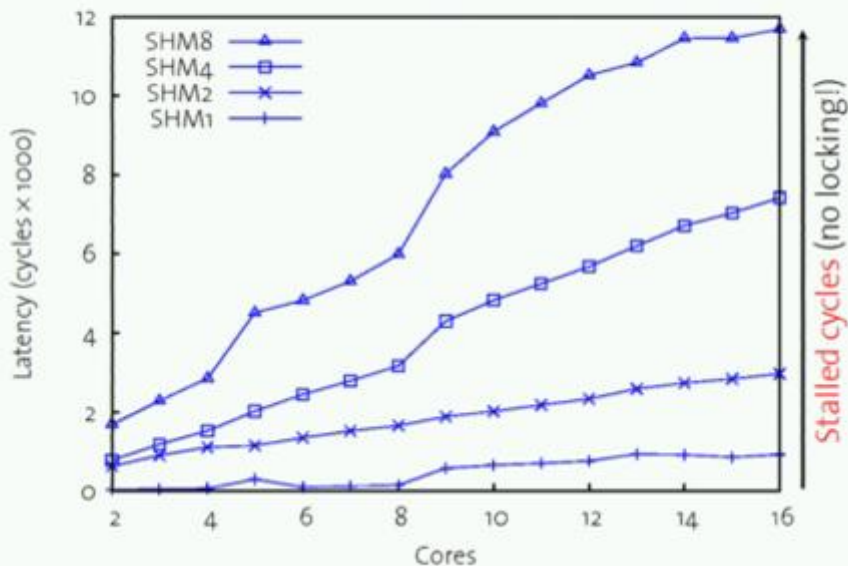
Shared memory results

4x4-core AMD system



Shared memory results


4x4-core AMD system



Message passing vs. shared memory: experiment


Message passing (move the operation to the data):

- ▶ A single server core updates the memory locations
- ▶ Each client core sends RPCs to the server
 - ▶ Operation and results described in a single cache line
 - ▶ Block while waiting for a response (in this experiment)



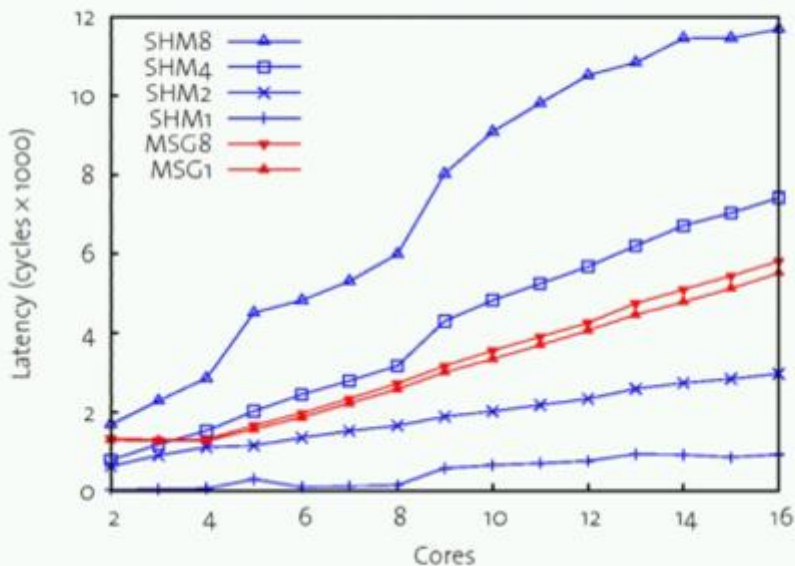
Message passing vs. shared memory: tradeoff

4x4-core AMD system



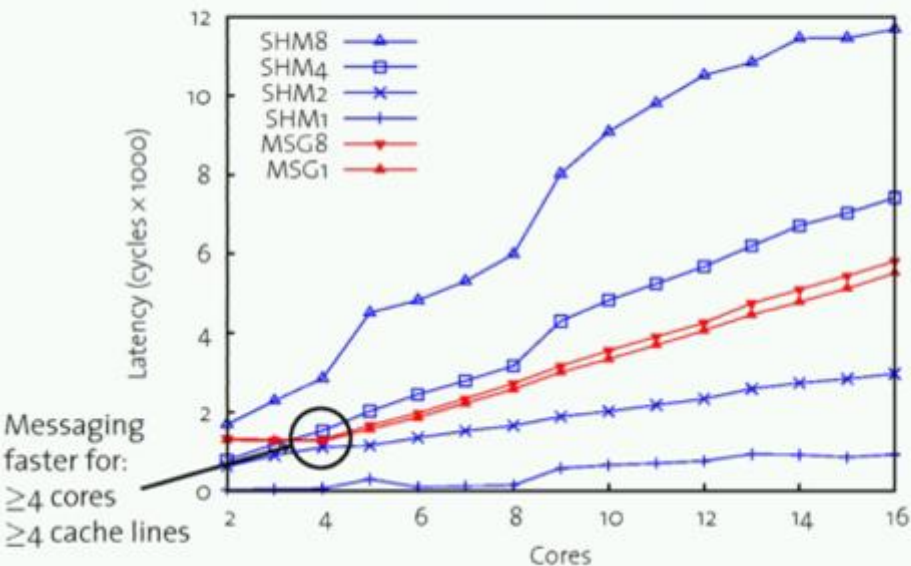
Message passing vs. shared memory: tradeoff

4x4-core AMD system




Message passing vs. shared memory: tradeoff

4x4-core AMD system




Message passing vs. shared memory: tradeoff

4x4-core AMD system



Message passing vs. shared memory: tradeoff

4x4-core AMD system



2. Make OS structure hardware-neutral

- ▶ **Separate OS structure from hardware**
- ▶ Only hardware-specific parts:
 - ▶ Message transports (highly optimised / specialised)
 - ▶ CPU / device drivers

2. Make OS structure hardware-neutral

- ▶ **Separate OS structure from hardware**
- ▶ Only hardware-specific parts:
 - ▶ Message transports (highly optimised / specialised)
 - ▶ CPU / device drivers
- ▶ Adaptability to changing performance characteristics
- ▶ Late-bind protocol and message transport implementations


3. View state as replicated

- ▶ Potentially-shared state accessed *as if* it were a local replica
 - ▶ Scheduler queues, process control blocks, etc.

3. View state as replicated

- ▶ Potentially-shared state accessed *as if* it were a local replica
 - ▶ Scheduler queues, process control blocks, etc.
- ▶ Required by message-passing model
- ▶ Naturally supports domains that do not share memory
- ▶ Naturally supports changes to the set of running cores
 - ▶ Hotplug, power management

Replication vs. sharing as default



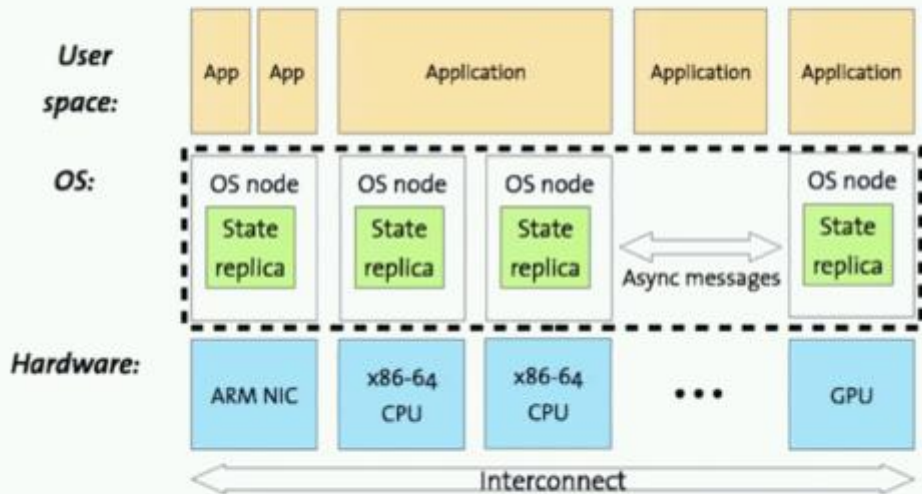
- ▶ Replicas used as an optimisation in previous systems:
 - Tornado, K42 clustered objects
 - Altix Linux read-only data, kernel text

Replication vs. sharing as default



- ▶ Replicas used as an optimisation in previous systems:
 - Tornado, K42 clustered objects
 - Altix Linux read-only data, kernel text
- ▶ In a multikernel, **sharing is a local optimisation**
 - ▶ Shared (locked) replica for threads or closely-coupled cores
 - ▶ Hidden, local
 - ▶ Only when faster, as *decided at runtime*
 - ▶ Basic model remains split-phase

The multikernel model



Outline

Introduction

Motivation

- Hardware diversity

The multitenant model

- Design principles

- The model

Barrelfish

Evaluation

- Case study: Unmap (TLB shootdown)

Future work

Barrelfish

- ▶ From-scratch implementation of a multikernel
- ▶ Open source (BSD licensed)
- ▶ Supports x86-32/64 multiprocessors
 - ▶ Beehive, SCC, ARM ports well on the way
 - ▶ Combinations of these, **in the same system**

Who's involved?



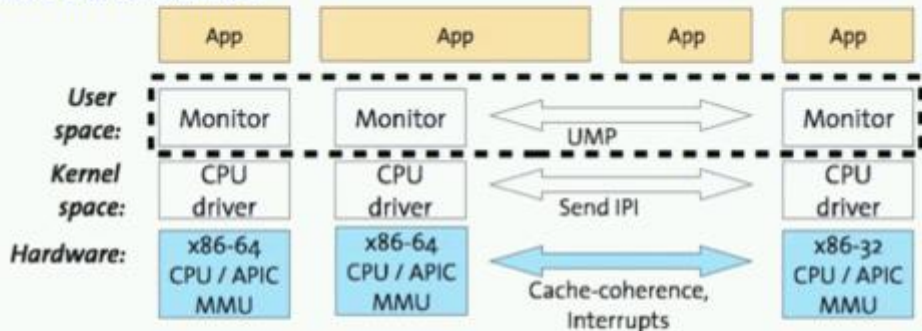
- ▶ Andrew Baumann
- ▶ Pierre-Evariste Dagand
- ▶ Simon Peter
- ▶ Timothy Roscoe
- ▶ Adrian Schüpbach
- ▶ Akhilesh Singhanian



- ▶ Paul Barham
- ▶ Richard Black
- ▶ Tim Harris
- ▶ Orion Hodson
- ▶ Rebecca Isaacs
- ▶ Ross McIlroy

Barrelfish structure

Monitors and CPU drivers



- ▶ CPU driver serially handles traps and exceptions
- ▶ Monitor mediates local operations on global state
- ▶ UMP inter-core message transport / interconnect driver
 - ▶ URPC on *current* (cache-coherent) x86 HW
- ▶ Other system facilities implemented as user-level services

Non-original ideas in Barrelfish

Multiprocessor techniques:

- ▶ Minimise shared state (Tornado, K42, Corey)
- ▶ User-space messaging decoupled from IPIs (URPC)
- ▶ Single-threaded non-preemptive kernel per core (K42)

Other ideas we liked:

- ▶ Capabilities for all resource management (seL4)
- ▶ Upcall processor dispatch (Psyche, Sched. Activations, K42)
- ▶ Push policy into application domains (Exokernel, Nemesis)
- ▶ Lots of information (Infokernel)
- ▶ Run drivers in their own domains (μ kernels)
- ▶ EDF as per-core CPU scheduler (RBED)
- ▶ Specify device registers in a little language (Devil)

Applications running on Barrelfish

- ▶ Slide viewer (this one!)
- ▶ Webserver (www.barrelfish.org)
- ▶ Virtual machine monitor (runs unmodified Linux)
- ▶ SPLASH-2, OpenMP, Phoenix MapReduce (benchmarks)
- ▶ SQLite
- ▶ ECLⁱPS^e (constraint engine)
- ▶ more...

Outline

Introduction

Motivation

Hardware diversity

The multi-armed model

Design principles

The model

Barrelfish

Evaluation

Case study: Unmap (TLB shutdown)

Future work

Evaluation goals

How do we evaluate an alternative OS structure?

- ▶ Good baseline performance
 - ▶ Comparable to existing systems on current hardware
- ▶ Scalability with cores
- ▶ Adapability to different hardware
- ▶ Ability to exploit message-passing for performance

UMP performance (one-way)


System	Cache	Latency cycles	Throughput msgs/kcycle
2×4-core Intel	shared	180	11.97
	non-shared	570	3.78
2×2-core AMD	same die	450	3.42
	one-hop	532	3.19
4×4-core AMD	shared	448	3.57
	one-hop	545	3.53
	two-hop	659	3.19
8×4-core AMD	shared	538	2.77
	one-hop	613	2.79
	two-hop	682	2.71

- ▶ Two HyperTransport requests on AMD
- ▶ Batching/pipelining for free



Case study: Unmap (TLB shutdown)

- ▶ Send a message to every core with a mapping, wait for all to be acknowledged
- ▶ Linux/Windows:
 1. Kernel sends IPIs
 2. Spins on shared acknowledgement count/event
- ▶ Barrelfish:
 1. User request to local monitor domain
 2. Single-phase commit to remote cores
- ▶ How to implement communication?

Unmap communication protocols




Unmap communication protocols




Unmap communication protocols

Raw messaging cost




Why use multicast

8×4-core AMD system




Why use multicast

8×4-core AMD system




Multicast communication




Why use multicast

8×4-core AMD system




Multicast communication



- “NUMA-aware” multicast

Unmap communication protocols

Raw messaging cost




System knowledge base

- ▶ Constructing multicast tree requires hardware knowledge
 - ▶ Mapping of cores to sockets (CPUID data)
 - ▶ Messaging latency (online measurements)
- ▶ More generally, Barrelfish needs a way to reason about diverse system resources

System knowledge base

- ▶ Constructing multicast tree requires hardware knowledge
 - ▶ Mapping of cores to sockets (CPUID data)
 - ▶ Messaging latency (online measurements)
- ▶ More generally, Barrelfish needs a way to reason about diverse system resources
- ▶ We tackle this with **constraint logic programming** [Schüpbach et al., MMCS'o8]
- ▶ **System knowledge base** stores rich, detailed representation of hardware, performs online reasoning
 - ▶ Initial implementation: port of the ECLⁱPS^e constraint solver
- ▶ Prolog query used to construct multicast routing tree

Unmap latency




Compute-bound (shared memory) workloads


4×4-core AMD system

- ▶ Barrelfish support for shared-memory applications
- ▶ Benchmarks from NAS OpenMP and SPLASH-2 suites
- ▶ User-level synchronisation primitives (e.g. spinlocks)
- ▶ Mainly exercise hardware coherence mechanisms

SPLASH-2 Barnes-Hut



NAS OpenMP fast Fourier transform



IO workloads

- ▶ Faster IP loopback through efficient use of interconnect
- ▶ Network throughput: 951.7Mbit/s (same as Linux)
- ▶ Pipelined web server



- ▶ Static: 640 Mbit/s vs. 316 Mbit/s for lighttpd/Linux
- ▶ Dynamic: 3417 requests/s (17.1Mbit/s) bottlenecked on SQL

Outline

Introduction

Motivation

- Hardware diversity

The multi-tenant model

- Design principles

- The model

Barrelfish

Evaluation

- Case study: Linmap (TLB shootdown)

Future work

Open questions and future work

- ▶ What are the right protocols for replication and agreement?
 - ▶ How can local sharing be used to optimise them?
 - ▶ Tradeoffs very different to classic distributed systems

Open questions and future work

- ▶ What are the right protocols for replication and agreement?
 - ▶ How can local sharing be used to optimise them?
 - ▶ Tradeoffs very different to classic distributed systems
- ▶ What does a native Barrelfish application look like?
 - ▶ Integration of high-level languages and multicore runtimes
 - ▶ New models for concurrent programming

Open questions and future work

- ▶ What are the right protocols for replication and agreement?
 - ▶ How can local sharing be used to optimise them?
 - ▶ Tradeoffs very different to classic distributed systems
- ▶ What does a native Barrelfish application look like?
 - ▶ Integration of high-level languages and multicore runtimes
 - ▶ New models for concurrent programming
- ▶ How do you construct system services above this model?
 - ▶ Networking architecture
 - ▶ Replicated file system based on cluster/HPC techniques

Open questions and future work

- ▶ What are the right protocols for replication and agreement?
 - ▶ How can local sharing be used to optimise them?
 - ▶ Tradeoffs very different to classic distributed systems
- ▶ What does a native Barrelfish application look like?
 - ▶ Integration of high-level languages and multicore runtimes
 - ▶ New models for concurrent programming
- ▶ How do you construct system services above this model?
 - ▶ Networking architecture
 - ▶ Replicated file system based on cluster/HPC techniques
- ▶ Heterogeneity-aware scheduling and resource allocation [Peter et al., HotPar'10]
 - ▶ Queuing effects, scheduling based on queue lengths
- ▶ Timeout behaviour in the OS [Peter et al., EuroSys'08]

Does the model extend outside a machine?

- ▶ A strong boundary at the machine edge is likely to remain:
 - ▶ Message guarantees (reliable, in-order)
 - ▶ Performance tradeoffs (latency / bandwidth)
 - ▶ Hardware support for shared memory
- ▶ However, a multikernel lowers the gap:
 - ▶ Explicit message-passing
 - ▶ Explicit split-phase replica-maintenance APIs

Does the model extend outside a machine?

- ▶ A strong boundary at the machine edge is likely to remain:
 - ▶ Message guarantees (reliable, in-order)
 - ▶ Performance tradeoffs (latency / bandwidth)
 - ▶ Hardware support for shared memory
- ▶ However, a multikernel lowers the gap:
 - ▶ Explicit message-passing
 - ▶ Explicit split-phase replica-maintenance APIs

Potential scenarios:

- ▶ Models for using resources inside one machine the same way as many machines
 - ▶ MapReduce, Dryad, etc.
- ▶ Extending your PC with cloud resources

Conclusion

- ▶ Modern computers are inherently distributed systems
- ▶ It's time to rethink OS structure to match
- ▶ **The Multikernel:** model of the OS as a distributed system
 1. Explicit communication, replicated state
 2. Hardware-neutral OS structure

Conclusion

- ▶ Modern computers are inherently distributed systems
- ▶ It's time to rethink OS structure to match
- ▶ **The Multikernel**: model of the OS as a distributed system
 1. Explicit communication, replicated state
 2. Hardware-neutral OS structure
- ▶ **Barrelfish**: concrete implementation
 - ▶ Real system
 - ▶ Reasonable performance on current hardware
 - ▶ Better scalability/adaptability for future hardware
 - ▶ Promising approach



www.barrelfish.org