


Power Management with DVFS

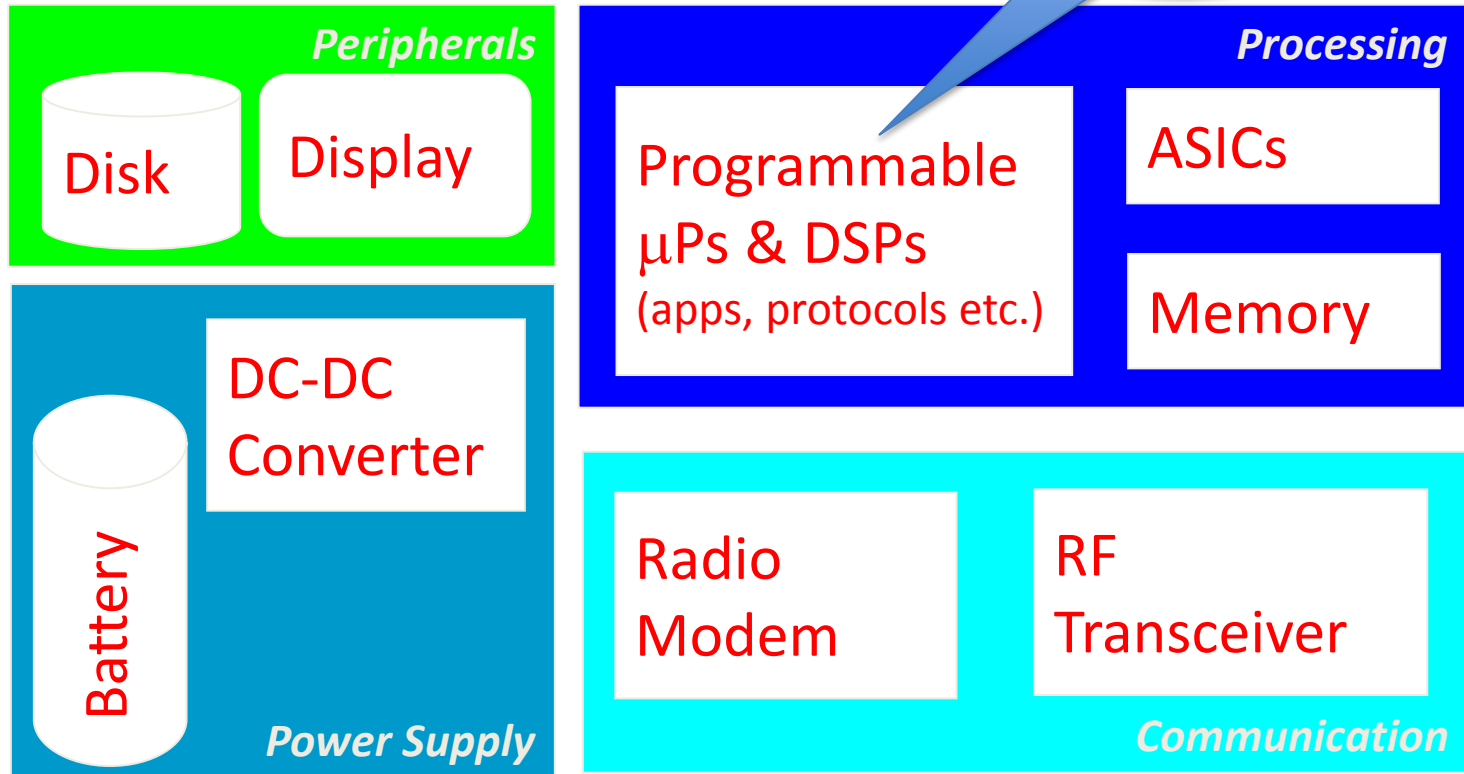
Instructor:
Zonghua Gu

The Case for Power Awareness



Where does the Power Go?

We focus
on CPU
power



Leakage vs. Dynamic


- Leakage power (static power) is consumed even if the circuit is idle (standby mode). The only way to avoid it is to turn off the CPU.
 - The only way to reduce leakage power is to turn off the power supply
- Dynamic power is still the main source of power consumption.
 - In the past, Leakage power was considered negligible compared to dynamic power.
 - Today, Leakage power is approaching or exceeding size of dynamic power.

Dynamic Power/Energy

- Power consumption: $P = \alpha * C * V^2 * f$
 - α : switching factor
 - C : effective capacitance
 - V : operating voltage
 - f : operating frequency
- Energy consumption for task with execution time t_{exe} (it is obvious that $t_{\text{exe}} \propto 1/f$):
 - $E = P * t_{\text{exe}} = \alpha * C * V^2 * f * t_{\text{exe}} = \alpha * C * V^2 * N_{\text{cycles}}$
 - N_{cycles} : total execution cycles
- Assuming $f \propto V$ (approximation):
 - $P \propto V^3 \propto f^3$; $E \propto V^2 \propto f^2$

Reducing Power Consumption

- Given power consumption $P = \alpha * C * V^2 * f$, we can reduce P by:
 - At device level:
 - reduce capacitance C
 - At system level:
 - Reduce CPU voltage V
 - Reduce CPU frequency f



Power vs. Energy

- Which is more important?
- Sometimes we care about power consumption:
 - heat dissipation, cooling
 - physical deterioration due to temperature.
- Sometimes we care about total energy consumption:
 - battery life for portable devices
 - power bills (money) for large servers

Dynamic Voltage and Frequency Scaling (DVFS)

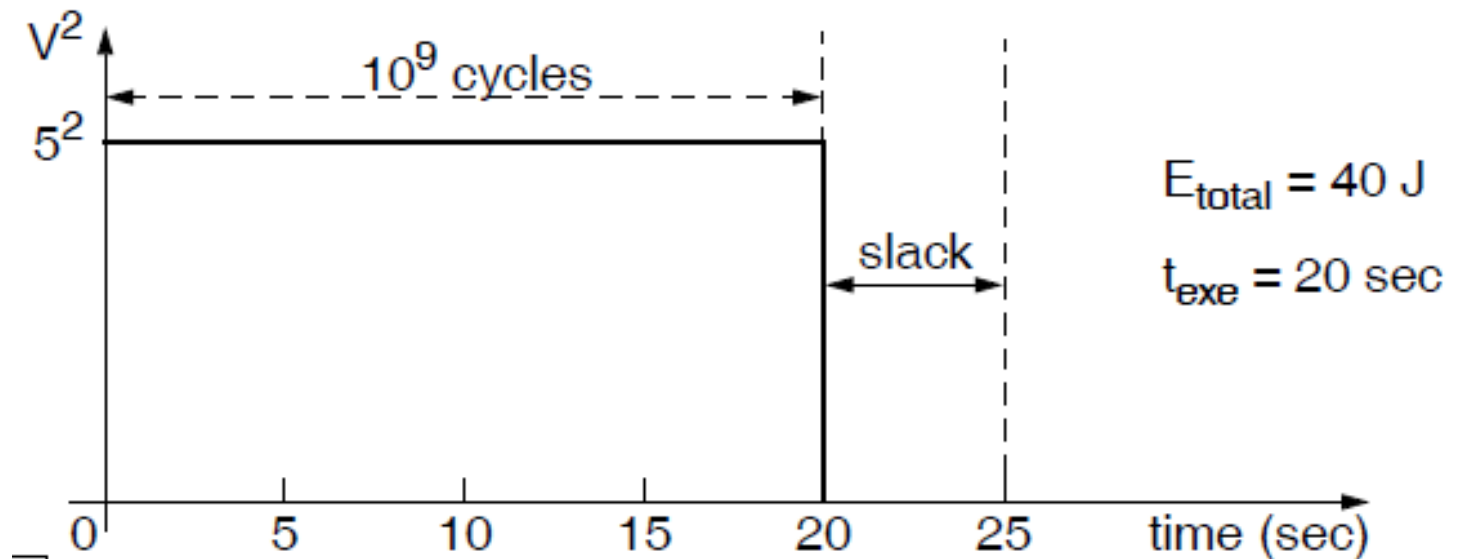
- A device can be run at different speeds at several different (supply voltage, CPU frequency) pairs
- Supply voltage can be chosen by the OS through execution of particular instructions. CPU frequency is automatically adjusted to fit the current supply voltage.
- Execution of jobs can be slowed down to save power as long as all jobs are completed by their deadline.

Scheduling w/ DVFS

- Reducing supply voltage increases execution time!
- The scheduling problem:
 - Which task to execute at a certain moment on a certain processor, *and at which voltage level, so that time constraints are fulfilled and energy consumption is minimized?*

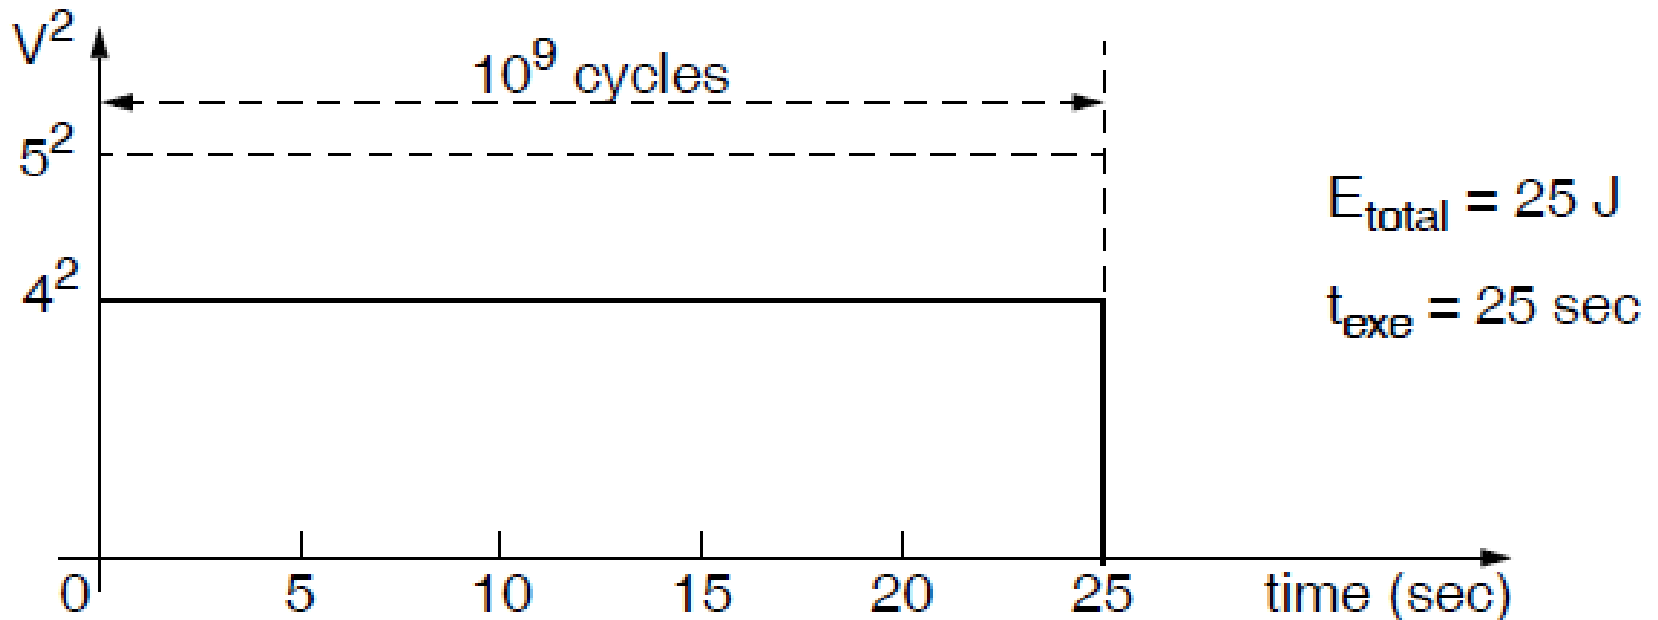
Example

- Consider a task τ , and CPU runs at full speed until τ finishes:
 - Processor nominal (maximum) voltage: 5V.
 - Total computation: 10^9 execution cycles; deadline: 25 seconds.
 - CPU frequency: 50MHz (50 million cycles/sec) at nominal voltage; Energy per cycle: 40 nJ/cycle at nominal voltage.
 - Execution time $t_{\text{exe}} = 10^9 \text{ (cycles)} / 50 * 10^6 \text{ (cycles/sec)} = 20 \text{ s}$
 - Total energy $E_{\text{total}} = 10^9 \text{ (cycles)} * 40 \text{ (nJ/cycle)} = 40 \text{ J}$




Reduced Voltage & Frequency

- Reduce voltage at time 0 to $V = 4V$
 - CPU frequency: $50 * (4/5) = 40\text{MHz}$; Energy per cycle: $40 * (4/5)^2 = 25\text{nJ/cycle}$.
 - Execution time $t_{\text{exe}} = 10^9(\text{cycles}) / 40 * 10^6 (\text{cycles/sec}) = 25 \text{ s}$
 - Total energy $E_{\text{total}} = 10^9 (\text{cycles}) * 25 (\text{nJ/cycle}) = 25 \text{ J}$



Reduced Voltage & Frequency

- Reduce voltage at time 15 to $V = 2.5V$
 - CPU frequency: $50 * (2.5/5) = 25\text{MHz}$; Energy per cycle: $40 * (2.5/5)^2 = 10\text{nJ/cycle}$.
 - Execution time $t_{\text{exe}} = 750 * 10^6 (\text{cycles}) / 50 * 10^6 (\text{cycles/sec}) + 250 * 10^6 (\text{cycles}) / 25 * 10^6 (\text{cycles/sec}) = 15 + 10 = 25 \text{ s}$
 - Total energy $E_{\text{total}} = 750 * 10^6 (\text{cycles}) * 40 (\text{nJ/cycle}) + 250 * 10^6 (\text{cycles}) * 10 (\text{nJ/cycle}) = 32.5 \text{ J}$



Basic Principle

- If a processor uses a single supply voltage and completes a program at exactly the deadline moment, the energy consumption is minimized.

[Zhang10]

- Xiao Zhang et al, An Evaluation of Per-Chip Non-Uniform Frequency Scaling on Multicores, 2010
- *Similarity grouping*:
 - On a multi-chip system, group applications with similar last level cache (LLC) miss ratios to run on the same multicore processor chip, in order to lower the voltage/frequency of the chip running memory-intensive tasks w/ high LLC miss ratios.
- Question: this seems opposite to [Blagodurov11] paper, which tries to distribute memory-intensive tasks on different chips?
 - [Blagodurov11] addresses NUMA, while [Zhang10] addresses UMA. But even for UMA, since each chip has its private LLC cache, the cache contention issue is similar.
 - [Zhang10]'s approach exacerbates cache contention on the chip running mem-intensive tasks, but power consumption is reduced by reducing voltage/frequency of the the chip running mem-intensive tasks.
 - [Blagodurov11] focuses on performance and does not consider power consumption.

[Zhang10]

- Question from audience
 - On a multi-chip system, tasks running on different chips may interfere with each other. Even though they do not share CPU, they share the bus, interconnect, memory controller, etc.
- The Frequency-to-Performance Model did not consider this interference, hence may not be very accurate.