

Data Sharing or Resource Contention: Toward Performance Transparency on Multicore Systems


Sharanyan Srikanthan
Sandhya Dwarkadas
Kai Shen

*Department of Computer Science
University of Rochester*




The Performance Transparency Challenge

Modern multicore systems...




Qualcomm Snapdragon...



IBM Power8's 12-core...



Intel's 10-core (20-thread), ...



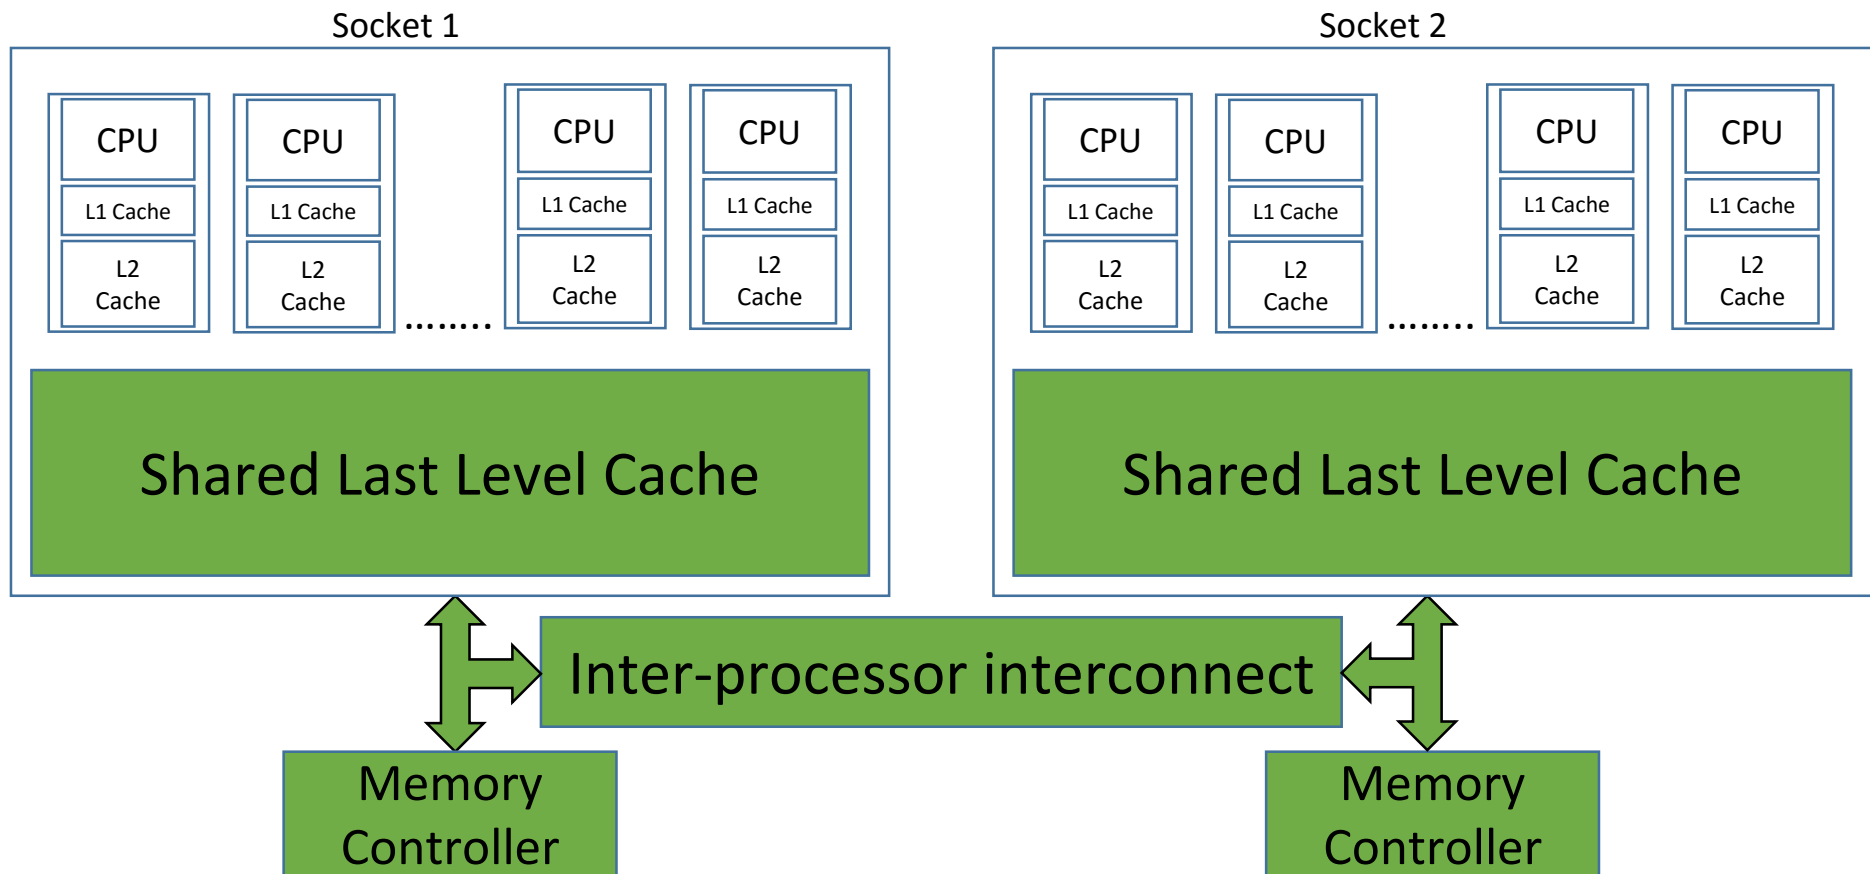
AMD's 16-core, ...

- Resource sharing
 - caches, on- and off-chip interconnects, memory
- Non-uniform access latencies



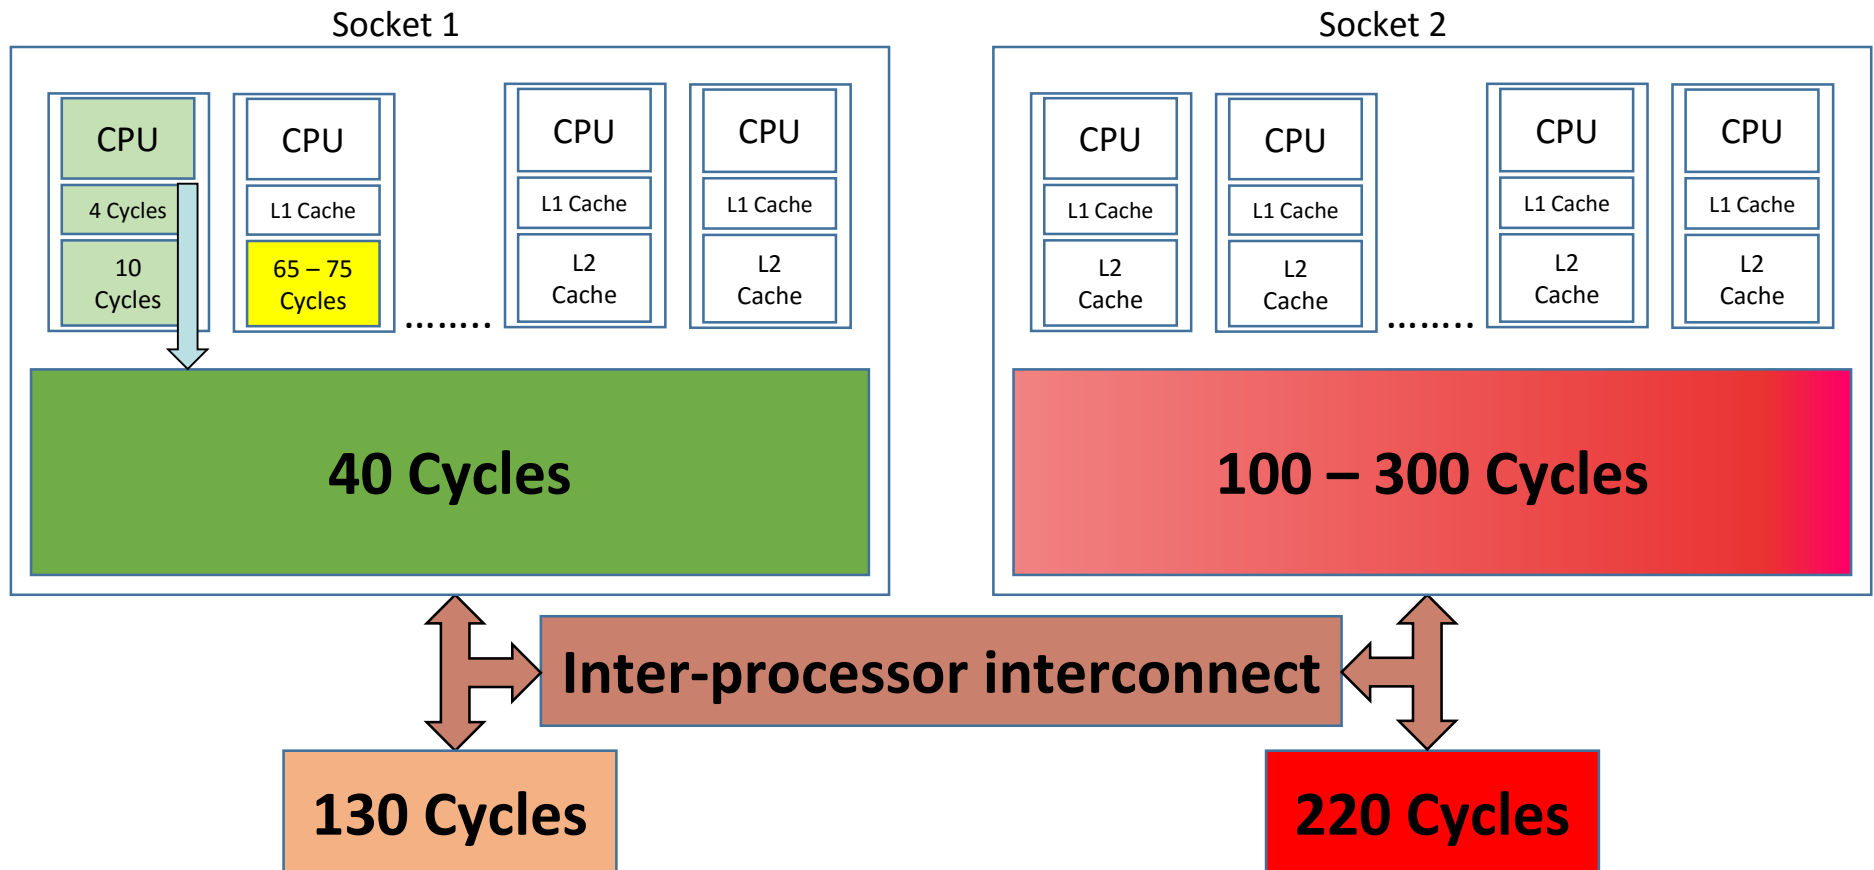
Modern Multicore Systems: Resource Sharing

Problem: Resource contention



Modern Multicore Systems: Non-Uniform Access Latencies

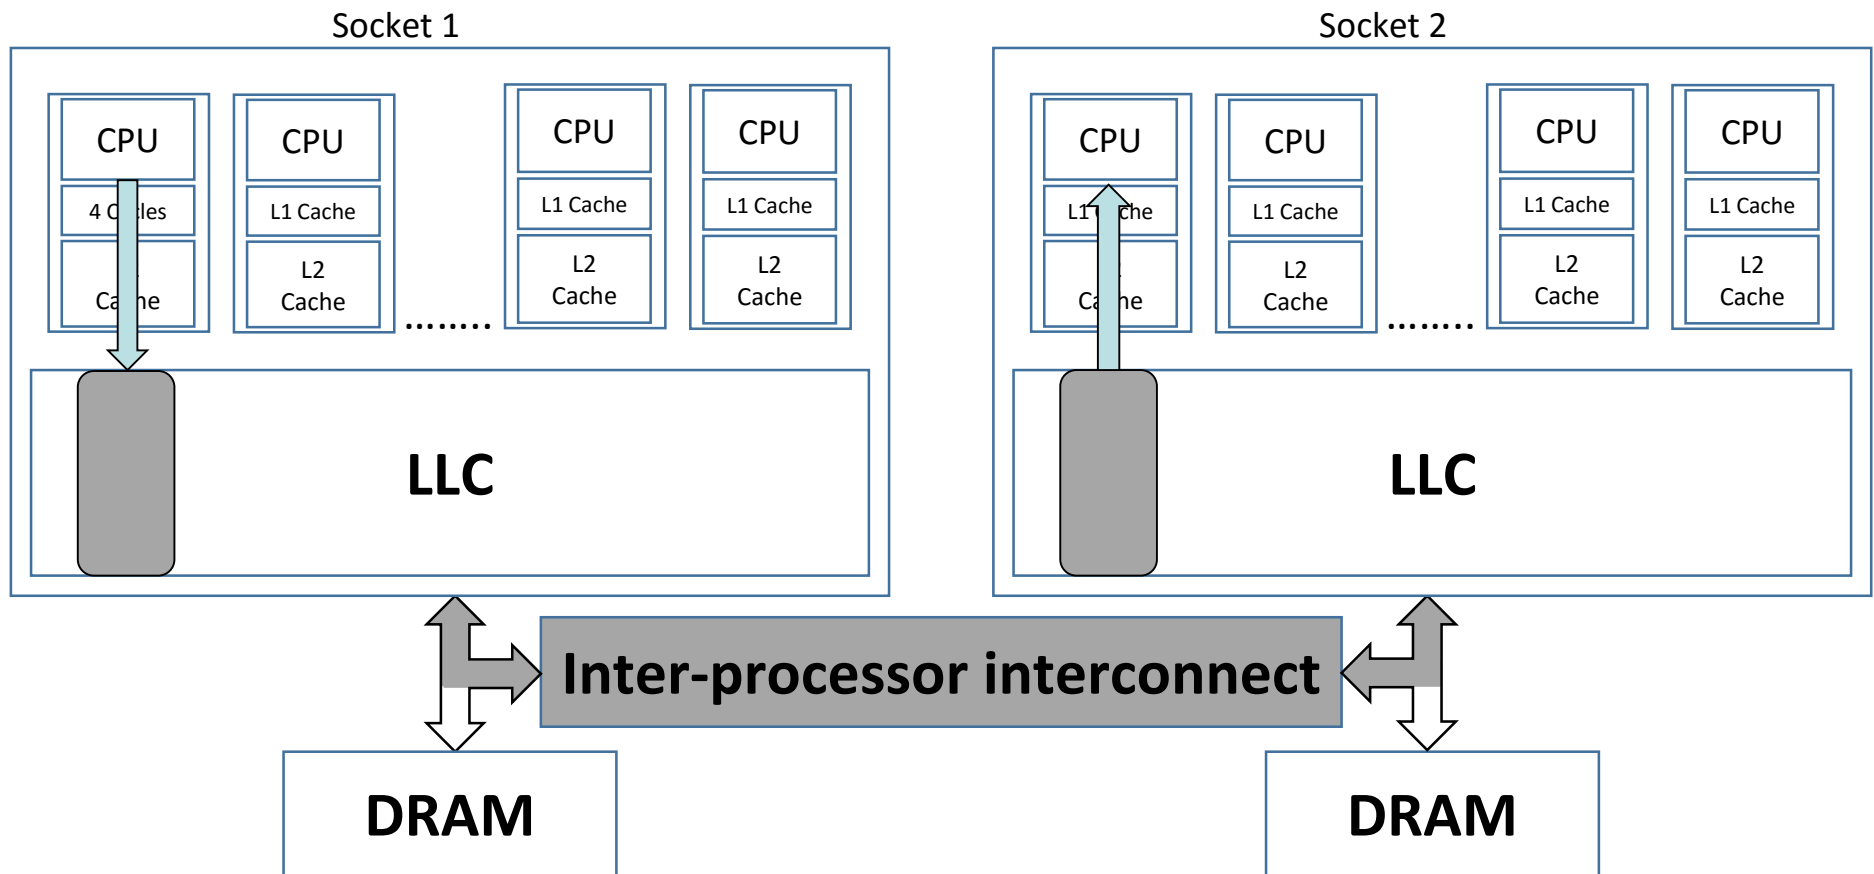
Problem: Communication costs a function of thread/data placement



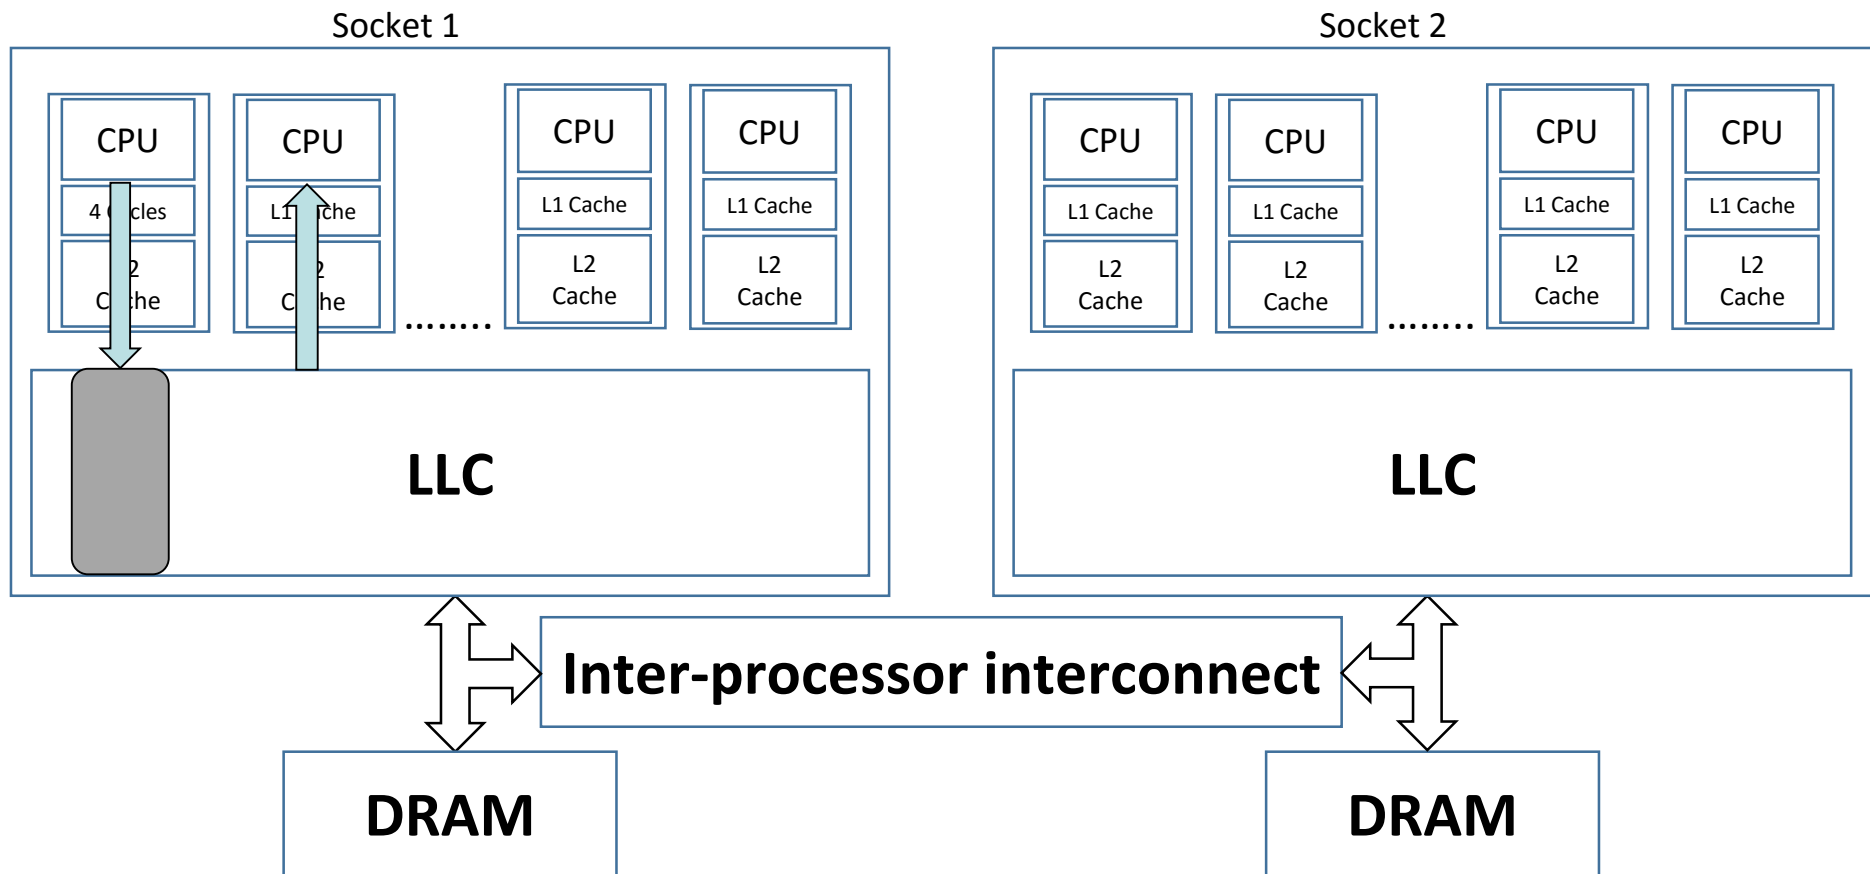
Ref: https://software.intel.com/sites/products/collateral/hpc/vtune/performance_analysis_guide.pdf



Impact of Thread Placement on Data Sharing Costs



Impact of Thread Placement on Data Sharing Costs



Solutions?

- Tam et al. [Eurosys 2007]: address sampling to identify data sharing, costly and specific to some processors
- Tang et al. [ISCA 2011]: performance counters to identify resource contention, does not consider non-uniform communication
- Lepers et al. [USENIX ATC, 2015]: *non-uniform memory bandwidth* utilization, does not consider data sharing



Our approach: **Sharing Aware Mapper (SAM)**:
Separates data sharing from resource contention using performance counters and considers non-uniformity when mapping parallel and multiprogrammed workloads



Sharing-Aware Mapper (SAM)

- Monitoring
- Identifying bottlenecks
- Taking remedial action



Sharing-Aware Mapper (SAM)

- Monitoring: periodically reading hardware performance counters to characterize application behavior
- Identifying bottlenecks
- Taking remedial action



Monitoring Using Performance Counters

- 4 metrics identified: 8 counter events need to be monitored
 - Inter-socket coherence activity
 - Last level cache misses served by remote cache
 - Intra-socket coherence activity
 - Last private level cache misses – (sum of hits and misses in LLC)
 - Local Memory Accesses
 - Approximated by LLC misses
 - Remote Memory Accesses
 - LLC misses served by remote DRAM
- 4 hardware programmable counters available: requires multiplexing



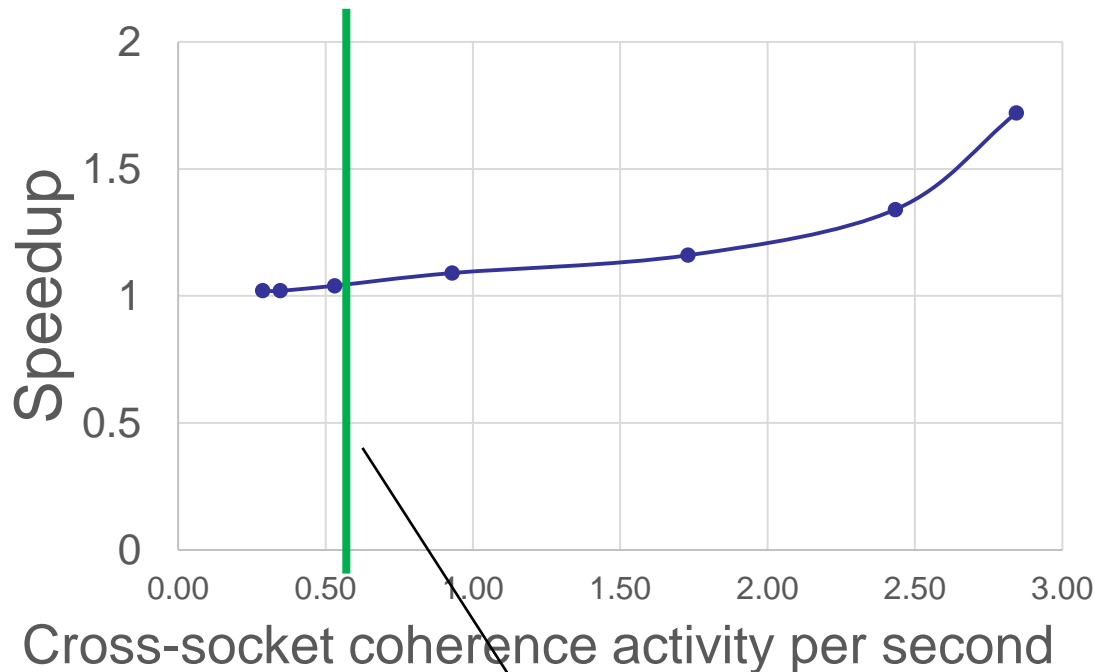
Sharing-Aware Mapper (SAM)

- Monitoring: periodically reading hardware performance counters to characterize application behavior
- Identifying bottlenecks: using pre-characterized thresholds
- Taking remedial action



Identifying Bottlenecks: Coherence Activity Threshold


Comparing performance of two tasks running on same socket against running on different sockets



- Microbenchmark forces data to move from one task to the other, generating coherence activity
- Rate of coherence activity varied by varying ratio of private to shared variable access



Identifying Bottlenecks: Memory Bandwidth Threshold



- Variable number of memory intensive threads run on the same processor
- ~10% difference in available bandwidth at saturation
- Conservative estimate can sufficiently identify bottlenecks

Threshold for flagging memory bandwidth bottleneck



Sharing-Aware Mapper (SAM)

- Monitoring: periodically reading hardware performance counters to characterize application behavior
- Identifying bottlenecks: using pre-characterized thresholds
- Taking remedial action: periodically remapping to manipulate process-processor affinity



SAM Decision Making

- Prioritize reduction of
 - Inter-socket coherence >
 - Remote memory access >
 - Socket bandwidth contention
- Co-locate threads with inter-socket coherence activity if
 - Idle cores are available >
 - CPU bound tasks can be moved >
 - Memory bound tasks can be moved
 - But not at the expense of separating threads with intra-coherence activity
- Return threads with remote memory accesses to original socket if possible
- Balance memory accesses across sockets
 - Transfer excess memory intensive threads to other sockets
 - Idle processors > CPU bound




SAM Decision Making

- Prioritize reduction of
 - Inter-socket coherence >
 - Remote memory access >
 - Socket bandwidth contention
- Co-locate threads with inter-socket coherence activity if
 - Idle cores are available >
 - CPU bound tasks can be moved >
 - *Memory bound tasks can be moved*
 - But not at the expense of separating threads with intra-coherence activity
- Return threads with remote memory accesses to original socket if possible
- Balance memory accesses across sockets
 - Transfer excess memory intensive threads to other sockets
 - Idle processors > CPU bound



Example: Data Sharing and Resource Contention




Tasks T1, T3 share data. Tasks T2 and T4 use memory.

- **There are 16 ways to map tasks to cores**
- Linux's default load balancing can result in
 - Inter-socket communication
 - Imbalanced memory bandwidth utilization
 - Remote memory accesses



Example: Data Sharing and Resource Contention

Step 1: Reduce inter-socket coherence




- Co-locate T1 and T3 to reduce cross-socket coherence

Tasks T1, T3 share data. Tasks T2 and T4 use memory.



Example: Data Sharing and Resource Contention

Step 2: Reduce remote memory accesses




- Swap T2 and T4 to eliminate remote memory accesses

Tasks T1, T3 share data. Tasks T2 and T4 use memory.



Example: Data Sharing and Resource Contention

Step 3: Balance memory bandwidth utilization

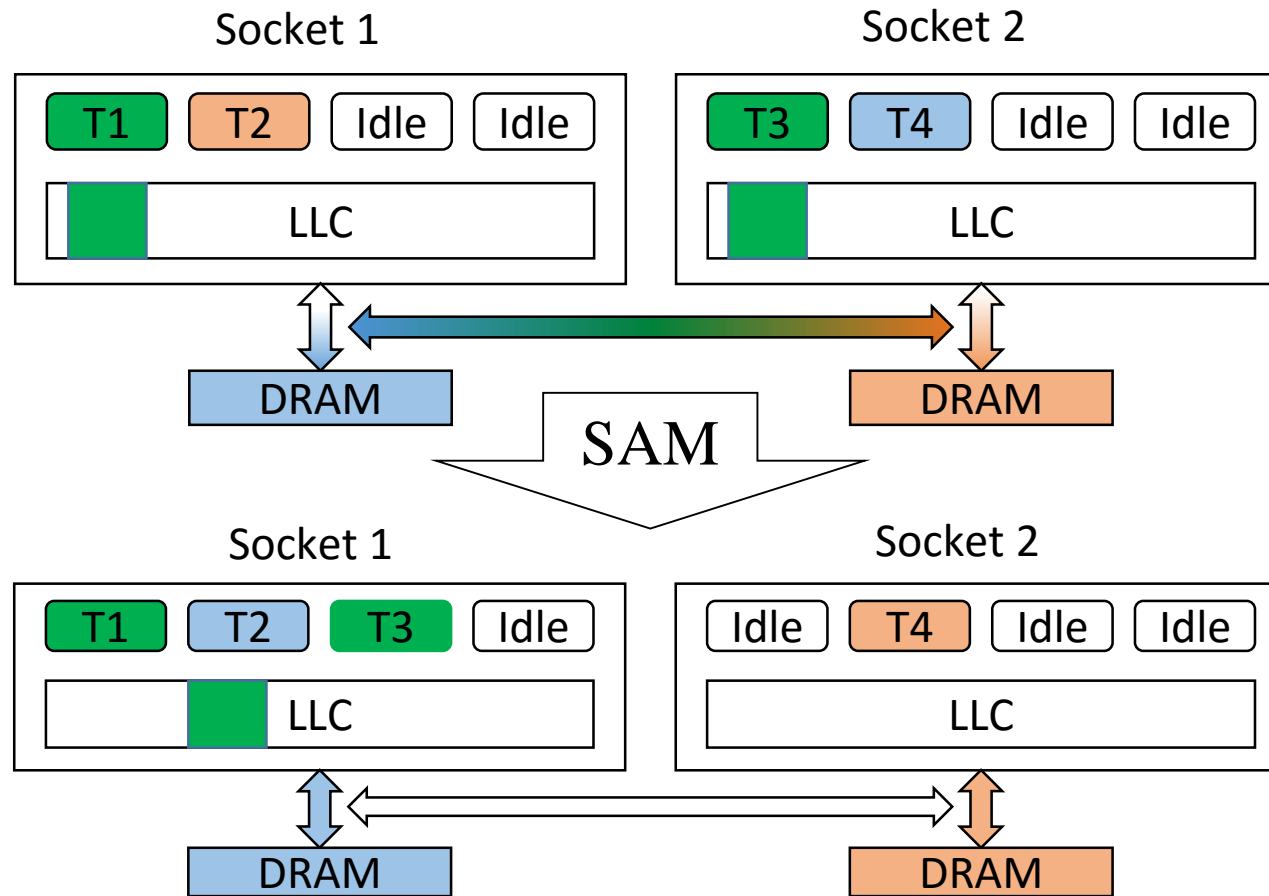


Tasks T1, T3 share data. Tasks T2 and T4 use memory.

- Memory bandwidth is already balanced so no additional action is performed



Example: Data Sharing and Resource Contention



Optimal Solution

- SAM identifies task characteristics with performance counters
- Minimizes inter-processor communication, remote memory accesses while maximizing memory bandwidth utilization



Experimental Environment

- Fedora 19, Linux 3.14.8
- Dual socket machine – Intel Xeon E5-2660 v2 “IvyBridge” processor (10 physical cores, 2 contexts each, 2.20 GHz, 25MB of L3 cache)
- Benchmarks
 - Microbenchmarks (Used for thresholding)
 - SPEC CPU '06 (CPU & Memory bound workloads)
 - PARSEC 3.0 (Parallel workloads – light on data sharing)
 - ALS, Stochastic Gradient Descent, Single Value Decomposition and other parallel workloads (Parallel workloads – stronger data sharing, dynamic behavior, and emerging workloads)




Implementation Context

Implemented as a daemon running periodically
using CPU affinity masks to control placement

Map threads to cores
(Sharing Aware Mapper – SAM)


Select which applications run
together (Linux Scheduler)



(Tasks to be scheduled)



Standalone Applications

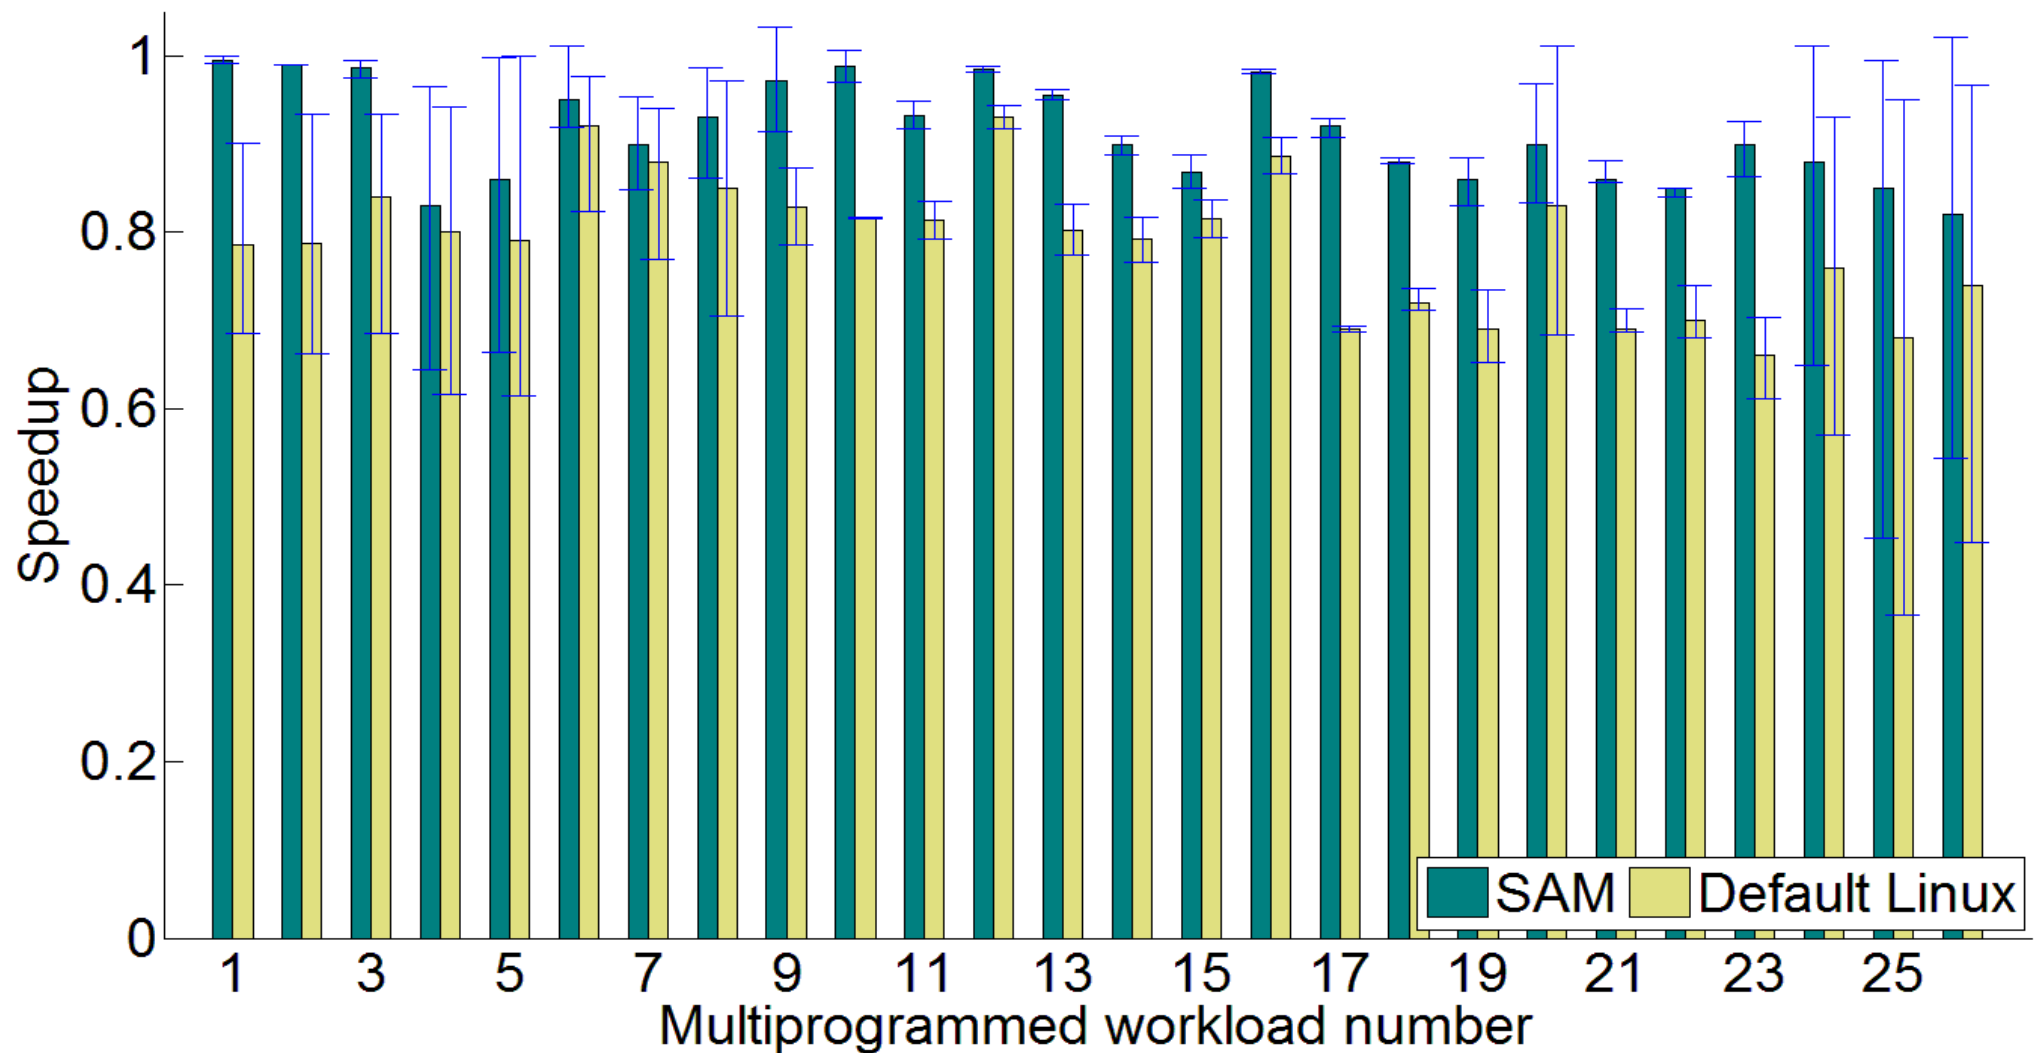


Baseline (Normalization factor): Best static mapping determined by exhaustive search

Speedup relative to Linux: Mean = 1.21, Min = 0.99, Max = 1.72



Multiple Applications



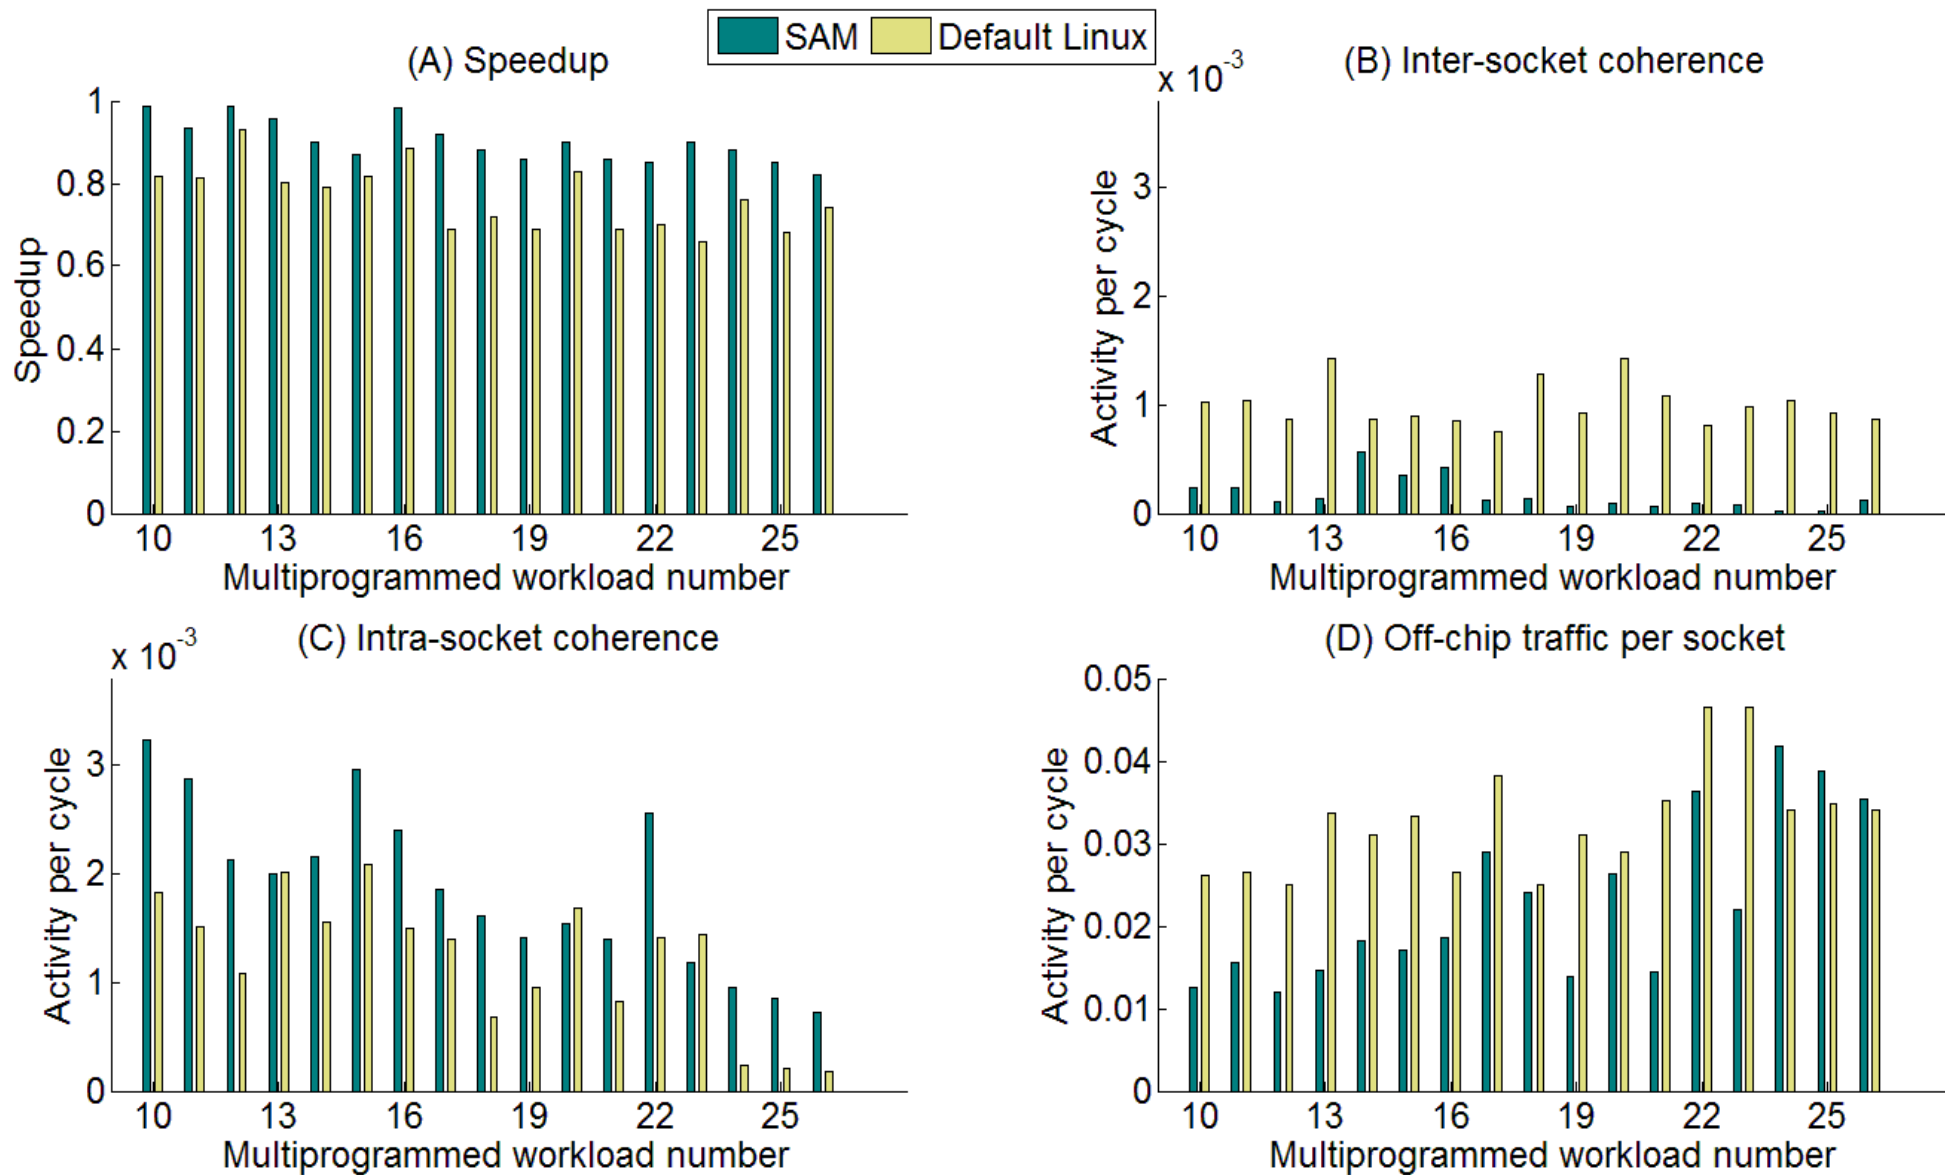
Speedup relative to Linux: Mean = 1.16, Min = 1.02, Max = 1.36

SAM improves fairness:

Average speedup spread per workload: SAM – 0.12, Linux – 0.18



Detailed Analysis



Overheads and Scaling

- SAM's overhead <1% (40 cores)
 - Performance counter reading
 - Invoked every tick – 8.9 μ s per tick = 8.9 msec per second
 - Constant time overhead
 - Data consolidation and bottleneck identification
 - Invoked every 100ms – 9.9 μ s per call = 99 μ s per second
 - Scales linearly with number processing cores
 - Decision making
 - Invoked every 100ms – negligible time related to data consolidation
 - $O(n^2)$ complexity, but for $n \leq 40$, time spent is within measurement error
- SAM's data consolidation and decision making is centralized



Conclusions

- Performance counters provide sufficient information to identify and separate data sharing from resource contention
- SAM improves performance and reduces performance disparities across applications by:
 - Minimizing cross-socket coherence
 - Minimizing remote memory accesses
 - Maximizing memory bandwidth utilization
- Future work to improve data sharing aware scheduling
 - Benefits of hyper-threading
 - Distributed decision making for better scalability



Thank you

Questions?

Data Sharing or Resource Contention: Toward Performance Transparency on Multicore Systems

Sharanyan Srikanthan
Sandhya Dwarkadas
Kai Shen

*Department of Computer Science
University of Rochester*

