

# **CMPT 300**

## **Introduction to Operating Systems**

---

Course Organization


# CMPT 300: Operating Systems

⦿ Time: Wed 5:30-8:20pm

🌀 Location: RCB 8100

 Textbook:

Modern Operating Systems,  
**Third edition**, Andrew S.  
Tanenbaum, Prentice Hall 2008



# Contact Information

---

- ⌘ Instructor: Zonghua Gu
- ⌘ Office: TASC 1, Room 9215
- ⌘ Office Hours: Wed 3:00-5:30 pm or by appointment
- ⌘
- ⌘ Email/MSN: zonghua(at)gmail.com
- ⌘ QQ: 59331972

# Web-site

---

- ❁ All the information discussed today and more can always be found on the class web-site
- ❁ To find the class web site go to <http://www.cs.sfu.ca/CourseCentral>
  - 💧 Select Course Home pages
  - 💧 Select the Homepage for CMPT 300
    - 2 sections: Dr. Evans (300D) and Dr. Gu (300E)
    - <http://www.cs.sfu.ca/CourseCentral/300/csguestk> is 300E

# Topics

---

- ⌘ History, Evolution, and Philosophies
- ⌘ The User's View of Operating System Services
- ⌘ Tasking and Processes
- ⌘ Inter-process Communication, Concurrency Control and Resource Allocation
- ⌘ Scheduling and Dispatch
- ⌘ Physical and Virtual Memory Organization
- ⌘ File Systems
- ⌘ Security and Protection

# **CMPT 300**

## **Introduction to Operating Systems**

---

Introduction / Review

# Hardware and Software


---

- ❁ A **computer** is a machine designed to perform operations specified with a set of instructions called a **program**.
- ❁ **Hardware** refers to the computer equipment.
  - 🟢 keyboard, mouse, terminal, hard disk, printer, CPU
- ❁ **Software** refers to the programs that describe the steps we want the computer to perform.
- ❁ The software that manages the hardware and shares the hardware between different application programs is called the **operating system**.

# Computer Hardware

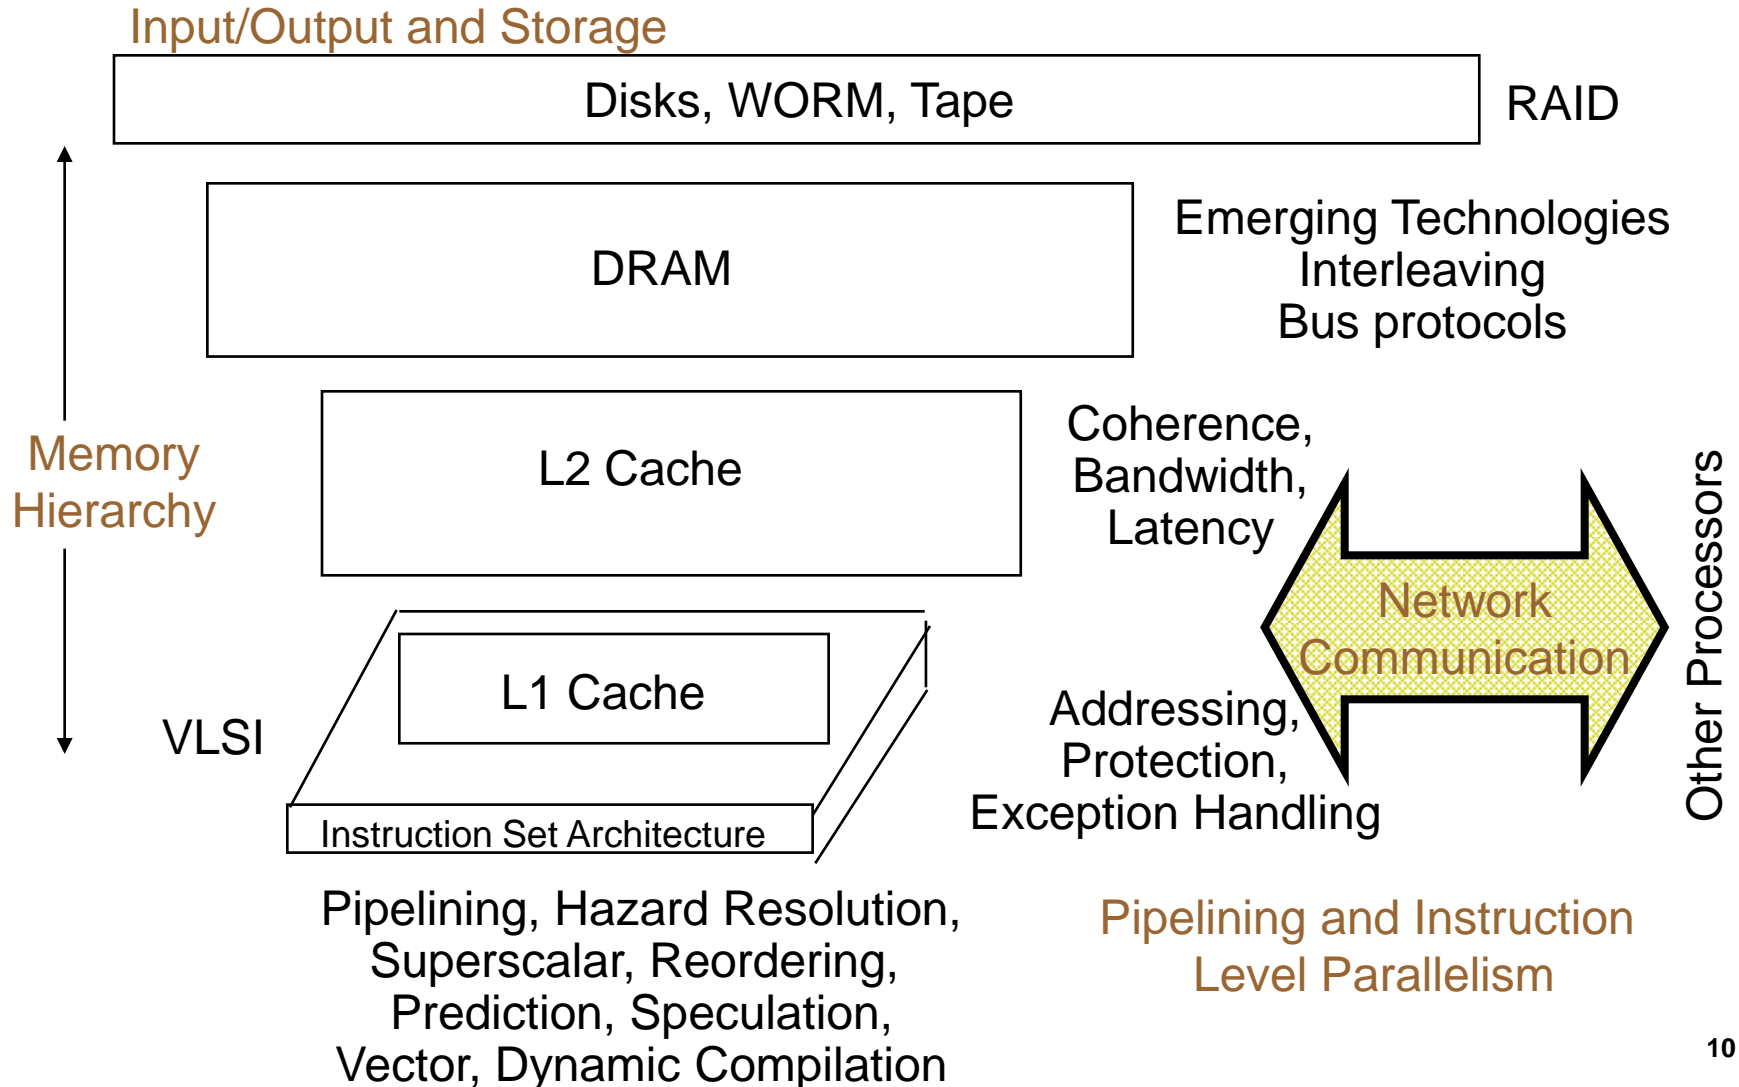
## ❁ Computer-system operation

- 💧 One or more CPUs, device controllers connect through common bus providing access to shared memory
- 💧 Concurrent execution of CPUs and devices competing for memory cycles






# Sample of Computer Architecture



# Increasing Software Complexity

Millions of lines of  
source code



From MIT's 6.033 course

# How do we tame complexity?

---

- ⚙ Every piece of computer hardware different
  - 💧 Different CPU
    - Pentium, PowerPC, ColdFire, ARM, MIPS
  - 💧 Different amounts of memory, disk, ...
  - 💧 Different types of devices
    - Mice, Keyboards, Sensors, Cameras, Fingerprint readers, touch screen
  - 💧 Different networking environment
    - Cable, DSL, Wireless, Firewalls,...
- ⚙ Questions:
  - 💧 Does the programmer need to write a single program that performs many independent activities?
  - 💧 Does every program have to be altered for every piece of hardware?
  - 💧 Does a faulty program crash everything?
  - 💧 Does every program have access to all hardware?

# Virtual Machine Abstraction

---

Application

---

Virtual Machine Interface

Operating System

---

Physical Machine Interface

Hardware

- ❖ Software Engineering Problem:
  - 💧 Turn hardware/software quirks  $\Rightarrow$  what programmers want/need
  - 💧 Optimize for convenience, utilization, security, reliability, etc...
- ❖ For any OS area (e.g. file systems, virtual memory, networking, scheduling):
  - 💧 What's the hardware interface? (physical reality)
  - 💧 What's the application interface? (nicer abstraction)

# Virtual Machines

## Software emulation of an abstract machine

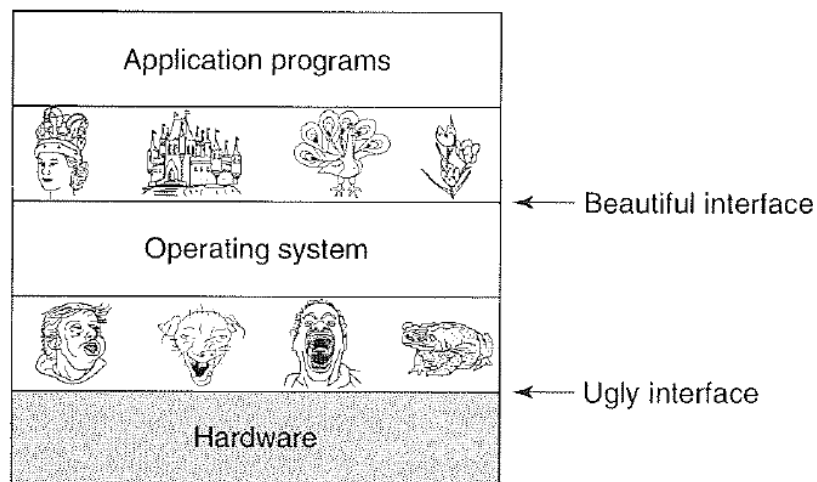
- Make it look like hardware has features you want
- Programs from one hardware & OS on another one

## Programming simplicity

- Each process thinks it has all memory/CPU time
- Different Devices appear to have same interface
- Device Interfaces more powerful than raw hardware
  - Bitmapped display  $\Rightarrow$  windowing system
  - Ethernet card  $\Rightarrow$  reliable, ordered, networking (TCP/IP)

## Fault Isolation

- Processes unable to directly impact other processes
- Bugs cannot crash whole machine



# What does an OS do?

---

- ❁ Silberschatz and Gavin: “An OS is Similar to a government”
  - ◆ Begs the question: does a government do anything useful by itself?
- ❁ Coordinator and Traffic Cop:
  - ◆ Manages all resources
  - ◆ Settles conflicting requests for resources
  - ◆ Prevents errors and improper use of the computer
- ❁ Facilitator (“useful” abstractions):
  - ◆ Provides facilities/services that everyone needs
  - ◆ Standard Libraries like Windowing systems
  - ◆ Make application programming easier, faster, less error-prone
- ❁ Some features reflect both tasks:
  - ◆ File system is needed by everyone (Facilitator) ...
  - ◆ ... but File system must be protected (Traffic Cop)

# What is an Operating System,... Really?

---

- ⌘ Most Likely:
  - 💧 Memory Management
  - 💧 I/O Management
  - 💧 CPU Scheduling
  - 💧 Synchronization / Mutual exclusion primitives
  - 💧 Communications? (Does Email belong in OS?)
  - 💧 Multitasking/multiprogramming?
  
- ⌘ What about?
  - 💧 File System?
  - 💧 Multimedia Support?
  - 💧 User Interface?
  - 💧 Internet Browser? 😊

# Operating System Definition (Cont'd)

---

- ❁ No universally accepted definition
- ❁ “*Everything a vendor ships when you order an operating system*” is good approximation
  - But varies wildly
- ❁ “*The one program running at all times on the computer*” is the OS kernel
  - Everything else is either a system program (ships with the operating system) or an application program



# Summary

---

- ❖ Operating systems provide a virtual machine abstraction to handle diverse hardware
- ❖ Operating systems coordinate resources and protect users from each other
- ❖ Operating systems simplify application development by providing standard services and abstractions
- ❖ Operating systems can provide an array of fault containment, fault tolerance, and fault recovery

# Machine language

---

- ⌘ Each type of processor (like Pentium 4, Athalon, Z80, ...) has its own instruction set
- ⌘ Each instruction in an instruction set does a single thing like access a piece of data, add two pieces of data, compare two pieces of data ...
- ⌘ Each instruction is represented by a unique number .This # may be different for different instruction sets, but no two instructions in the same instruction set should have the same #

# Machine Language programs

---

- ❁ A machine language program is a list of instructions
  - 💧 Each instruction is represented by a number
  - 💧 Inside the memory of the computer, each number is represented in binary (as a string of 1's and 0's)
  - 💧 The long string of 0's and 1's is easy for the computer to understand, but difficult for a human to read or write

# Assembly

---

- ❁ Assembly languages make it easier for the programmer.
  - 💧 Assembly is easier for humans to read/write
  - 💧 Use mnemonics like ADD, CMP, ... to replace the numbers that identify each of the instructions in the instruction set
  - 💧 The code for an Assembly program is written into a text file, which is translated into machine language program and executed.

# Computer Software: Languages

- ❁ Some Computer Languages
  - ◆ Machine language (machine instruction set)
  - ◆ assembly language
  - ◆ high level languages (Compilers/Interpreters)
    - C, C++, Ada, Fortran, Basic, Java
    - Do YOU know of any others?
    - mathematical computation tools (MATLAB, Mathematica, ...)
- ❁ Application software is written using computer languages.
- ❁ Operating systems are also written using computer languages (often C, some assembly)

# Computer Software: Applications

## ❁ Application Software (Software Tools)

- 💧 Word processors (Microsoft Word, WordPerfect, ...)
- 💧 Spreadsheet programs (Excel, Lotus1-2-3, ...)
- 💧 Computer games
- 💧 Communication software (email, chat, web browser...)
- 💧 Telecommunication software (VOIP, ...)
- 💧 Integrated programming environments


# User mode / kernel mode

---

- ❁ Most application software runs in user mode.  
Applications have access to a subset of the instruction set that does not include most direct hardware access
- ❁ Operating systems run in kernel mode (supervisor mode) and have access to the complete instruction set, including the instructions used to directly manage the hardware
- ❁ Application software running in user mode can use system calls to access hardware managed by the Operating System
- ❁ User mode programs may perform duties for the OS

# Modes

---




For some operating systems there may not be a separation between kernel mode and user mode (embedded systems, interpreted systems)



# Basic computer configuration

---



# Registers in CPU

---

- ⌘ Data registers
- ⌘ Accumulator
- ⌘ Address registers
- ⌘ Control/Status registers
  - Program counter
  - Stack pointer
  - Instruction register
  - Status registers

# Controller


---

- ❁ Fetch, Decode, Execute cycle (each instruction)
  - ◆ Fetch next instruction: Instruction contains op-code and possibly data
  - ◆ Decode op-code
  - ◆ Execute op-code (using data if necessary)
- ❁ Instructions
  - ◆ access data, moving it from memory (or disk or cache) to/from registers, and between registers
  - ◆ Complete arithmetic and logical manipulations of data in registers

# Executing an instruction (1)

---

- ⌘ Examine program counter
- ⌘ **Fetch** instruction indicated by program counter



- ⌘ Increment program counter to point at next instruction to be executed
- ⌘ Place fetched instruction in instruction register

# Executing an instruction (2)

---

- ⌘ Decode the instruction



- ⌘ Determine what is to be done
- ⌘ If needed, load address into an address register

# Executing an instruction (3)

---

- ❁ Execute the instruction in the instruction register, may result in one or more of the following
  - 💧 Fetch any data from memory (locations given in instruction) and place into the appropriate data registers
  - 💧 Place results from a data register or the accumulator into a memory location indicated in the instruction
  - 💧 Operate (arithmetic or logical operation) on data in data registers and save the result in the indicated register
  - 💧 Control the flow of the program (for example change the value in the program counter register)

# Adding 2 numbers ( $Z=X+Y$ )

---

- ❁ Program counter points to instruction to load value at location X
- ❁ Instruction is fetched into the instruction register, decoded and executed to load the first number into data register A
- ❁ Program counter is incremented and now points to an instruction to load value at location Y
- ❁ Instruction is fetched into the instruction register, decoded and executed to load the second number into data register B
- ❁ Program counter is incremented and now points to an instruction to add the values in data register A and B
- ❁ Instruction is fetched into the instruction register, decoded and executed to add the two numbers and place the result in the accumulator register.
- ❁ Program counter is incremented and now points to an instruction to place the value in the accumulator register into memory location Z
- ❁ Instruction is fetched, decoded and executed to place result in Z
- ❁ ([See animation](#))

# Memory Hierarchy

---

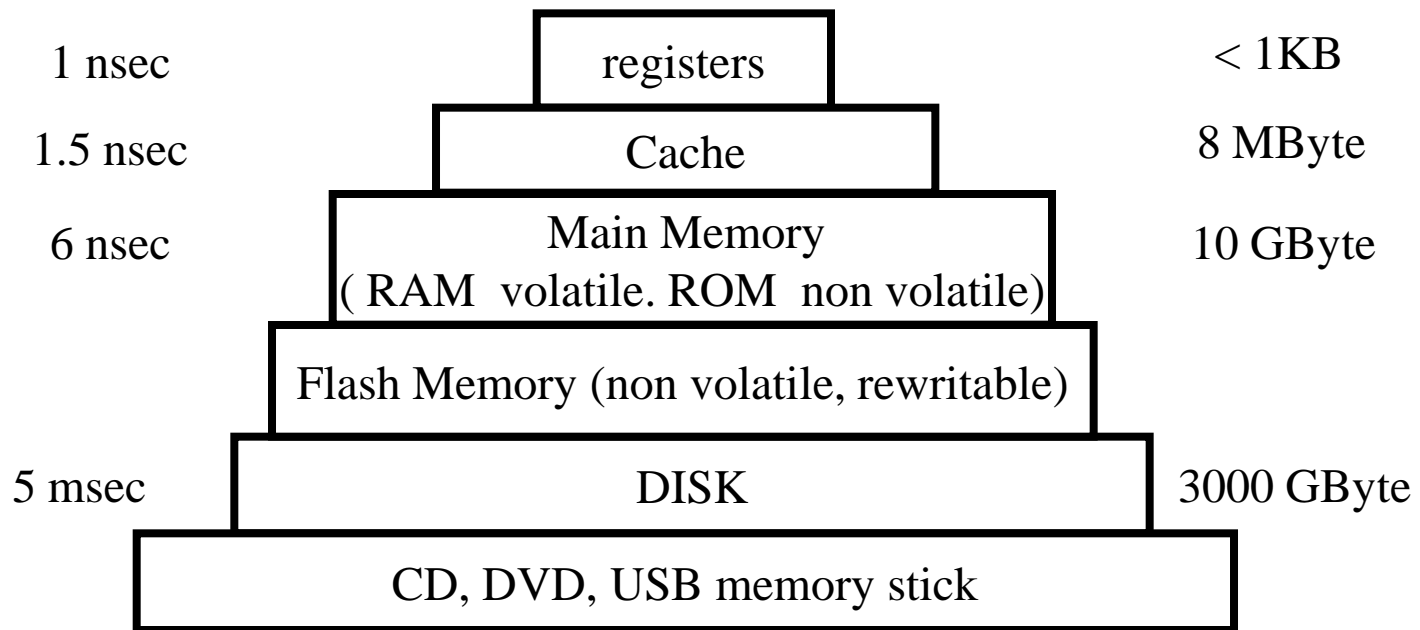
- ❁ Different types of memory have different access speeds and costs
- ❁ Faster access speed implies higher cost
- ❁ Greater capacity often implies lower access speed
- ❁ From fastest access to slowest access
  - ◆ Registers
  - ◆ Cache
  - ◆ Memory
  - ◆ Disk
  - ◆ Tapes



# Memory

---

❁ Modern computers use several kinds of storage



# Memory Hierarchy

---

- ⚙ As you go down the pyramid
  - a) Decreasing cost per bit, Increasing capacity
  - b) Increasing access time, Decreasing frequency of access
- 💧 Note that the fastest memory, sometimes referred to as **primary memory**, is usually **volatile** (register, cache, main memory)
- 💧 Non-volatile (continues to store information when the power is off) memory is usually slower. Referred to as **secondary** or **auxiliary memory**. Examples flash memory (flash that holds the BIOS, or removable flash), internal and external hard drives, CD, tape, ...

# Registers and cache

---

- ⌘ Parts of the CPU
- ⌘ Register access speed comparable to CPU clock speed
- ⌘ Cache memory may be as fast or as much as several times slower
- ⌘ Registers
  - Usually 64x64 for 64-bit machine, 32x32 for 32-bit machine
  - Usually < 1 Kbyte
- ⌘ Cache
  - As much as 8Mbytes


# Concept of Cache

---

- ❁ Provide memory on the CPU that is slower than the registers, cheaper and larger than registers, but can be accessed much faster than main memory
- ❁ The next few instructions, and data that will be needed will be loaded into cache in anticipation of faster access by the processor.

# Cache and main memory

---



# Using Cache

---

- ❁ Instructions (and data) are generally moved from main memory to cache in blocks; one such block ( $N$  bytes of memory) is called a cache line
- ❁ Cache has a series of slots, each  $N$  bytes long, and can hold a copy of one cache line.
- ❁ The main memory can contain many more cache lines than there are slots in the cache.
- ❁ Each time a copy of a new cache line is loaded into a cache slot, the original content of that slot is overwritten

# Cache design

---

- ❁ Cache size and Cache line size
  - 💧 Determined to optimize access time
- ❁ Mapping function
  - 💧 Which cache lines may be loaded into which cache slots
    - can any line go in any slot, or is there a mapping function to define rules governing which line can be place in which slot
- ❁ Replacement algorithm
  - 💧 When is a cache line in a cache slot replaced by another cache line

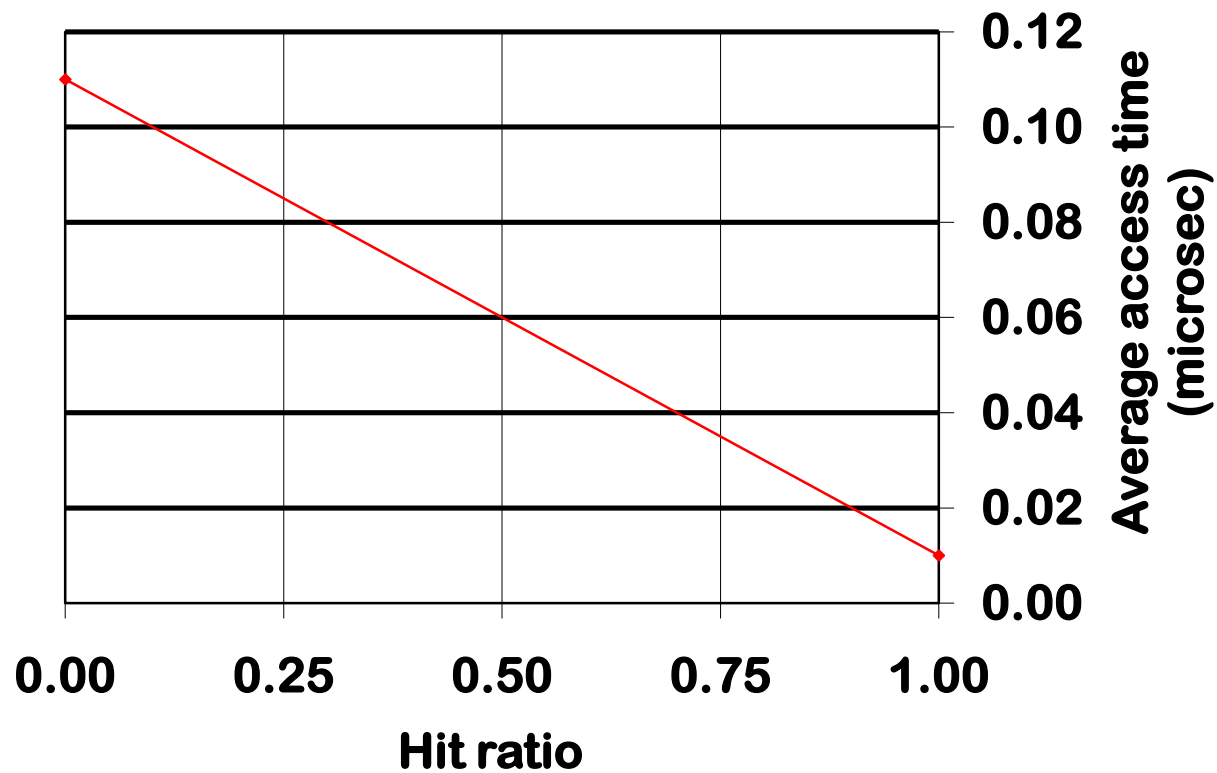
# Hit ratio

---

- ❁ A hit occurs when a memory access finds its information in the cache.
- ❁ A miss occurs when it is necessary to access the slower main memory (or lower level cache) to find the information.
- ❁ The proportion of accesses that are hits is called the hit ratio
- ❁ Consider an example,
  - ◆ Assume that any access to the cache takes  $.01\mu\text{s}$ , any access to main memory takes  $0.1\mu\text{s}$  (100 nsec)
  - ◆ Any instruction or piece of data not in cache will have to be accessed in main memory and moved to cache before being accessed  $(0.1+0.01)\mu\text{s}$
  - ◆ For a hit ratio of  $x\%$ , the average memory access time
    - $0.01 \cdot x\% + 0.11 \cdot (1-x\%)$



# Hit ratio




# Cache Line size

---

- ❁ As cache line size increases from a single byte, the hit ratio will increase at first.
  - ◆ It is very likely that bytes near a needed byte will be accessed in the near future (principle of locality (spatial & temporal))
- ❁ But as cache line size increases, the number of lines decreases
  - ◆ As cache line size increases past it's optimal value then the hit ratio will begin to decrease
  - ◆ This happens when it becomes more probable that the next access will involve the cache line that was just removed, i.e., the useful cache line was kicked out prematurely.
- ❁ Performance heavily depends on the application workload, mapping function and replacement algorithm, hence difficult to generalize.

# Effect of Line size (example)

Sisoft Sandra 2007 - Cache/Memory Speed (1 of 2)



NOTE: MB/s = 1/nsec \* 1000


# Cache specifications on common systems

---

- ❁ Most modern CPUs have an on-chip cache called an L1 cache
- ❁ Many modern systems also have a 2<sup>nd</sup> , and even 3<sup>rd</sup> level of cache between the L1 cache and the main memory called the L2 (L3) cache.
  - 🟢 L2 cache can be on-chip or off-chip (connected to the CPU via a bus)
  - 🟢 L3 cache is typically off-chip


# Multiple levels of cache: L2

---




# Multiple levels of cache: L3

---



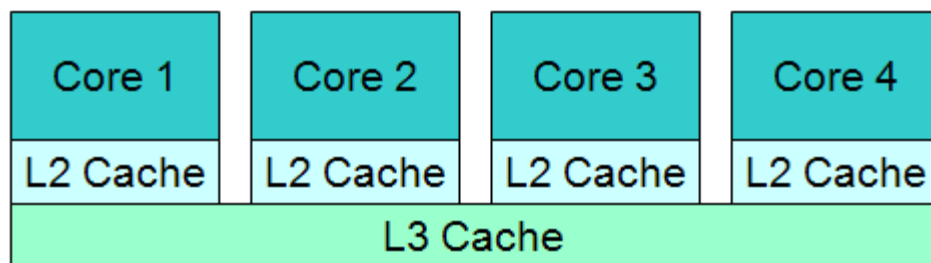
# Modern Cache Architectures

---



Core 2 Duo    Core2 Quad

Shared cache requires more complicated cache controller



Nehalem (i5, i7)    AMD K10 (Phenom 9)

Individual caches are more difficult to keep coherent (properly synchronized)

On-chip L1 caches are omitted from the figure

# Memory

---

- ❁ Main memory is typically DRAM (Dynamic Random Access Memory)
- ❁ Cache is typically SRAM (Static Random Access Memory)
  - 💧 Smaller and faster than DRAM
- ❁ Both are volatile: contents lost when power is turned off



# Disk

---

- ⌘ Hard disk
- ⌘ CD, DVD, Blu-Ray

Disk storage is much cheaper than memory

- ⌘ (3GB memory or 2000GB disk about the same cost)
- ⌘ Access time for disk is at least one order of magnitude slower than for memory

# Input / Output

---

- ❌ Reading or writing data from a peripheral device is not simple
  - 💧 Each device has its own controller (hardware)
  - 💧 Each device is managed by a device driver (software to use the controller)
    - Device drivers are specific to hardware and to the operating system using the device
  - 💧 Input and output is SLOW in comparison to CPU operations.

# Busy Waiting

---

- ❁ Reading or writing data to a device is SLOW in comparison to the time it takes to complete one CPU operation
- ❁ The CPU must send one or more instructions to the controller to make the I/O device begin to read or write. These instructions tell the controller where to find the data and/or where to store the data
- ❁ The CPU can then wait until the I/O operation is finishes.
  - ◆ While it waits the CPU will be in a loop
  - ◆ Each time through the loop the CPU will check a register in the controller to see if the I/O operation is complete
  - ◆ This is called **busy waiting**.

# Alternatives to Busy waiting

---

- ❁ Busy waiting does not use CPU resources efficiently
- ❁ Want to use the CPU to execute other instructions while the I/O operation is being completed by the I/O device (instead of executing the busy waiting loop)
- ❁ To use the CPU for other calculations while the I/O device controller completes the I/O operation, we need to use interrupts (a mechanism to tell the CPU when the controller completes the I/O)

# Interrupts

---

- ❁ Interrupts are a mechanism by which other modules (memory, I/O, timers ...) may interrupt the normal sequence of instructions being executed by the processor
- ❁ Interrupts are a critical component of the operation of spooling and multiprogramming (more later).
- ❁ Interrupts allow the transfer of control between different programs (remember the OS is also a program)
- ❁ Interrupts are generated by hardware (asynchronous)
  - ◆ Exceptions are generated by particular instructions in software (synchronous), e.g., divide by 0, overflow, illegal instruction or address...

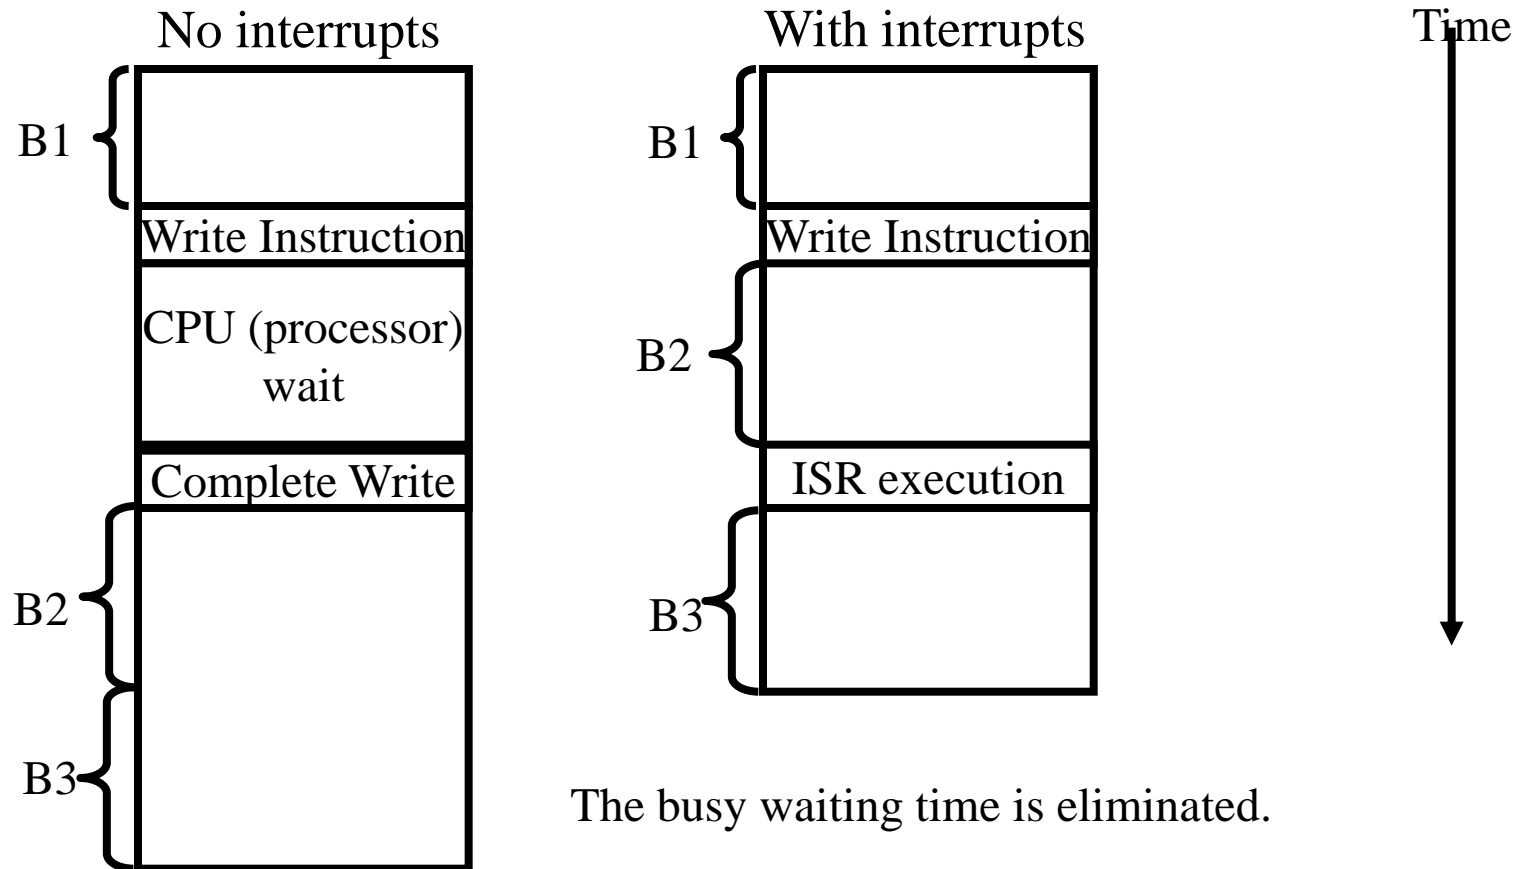
# Some types of interrupts

---

- ⌘ I/O
  - ◆ Signaling normal completion of an operation (read or write)
  - ◆ Signaling error during operation
- ⌘ Timer expiry
  - ◆ Begin regularly scheduled task
  - ◆ End task that has exceeded allocated time
- ⌘ Hardware failure

# Increase in efficiency

---



# Interrupt example: Output (1)

---

- ⌘ Program executes until it reaches a write instruction
- ⌘ The write instruction sets up the hardware output device operation then leaves the output device controller (not the CPU) processing the output
- ⌘ The program continues execution of additional instructions (N+3 to N+7 next slide)
- ⌘ When the hardware device completes the output operation, it generates and sends an interrupt to the CPU, signaling normal completion of output

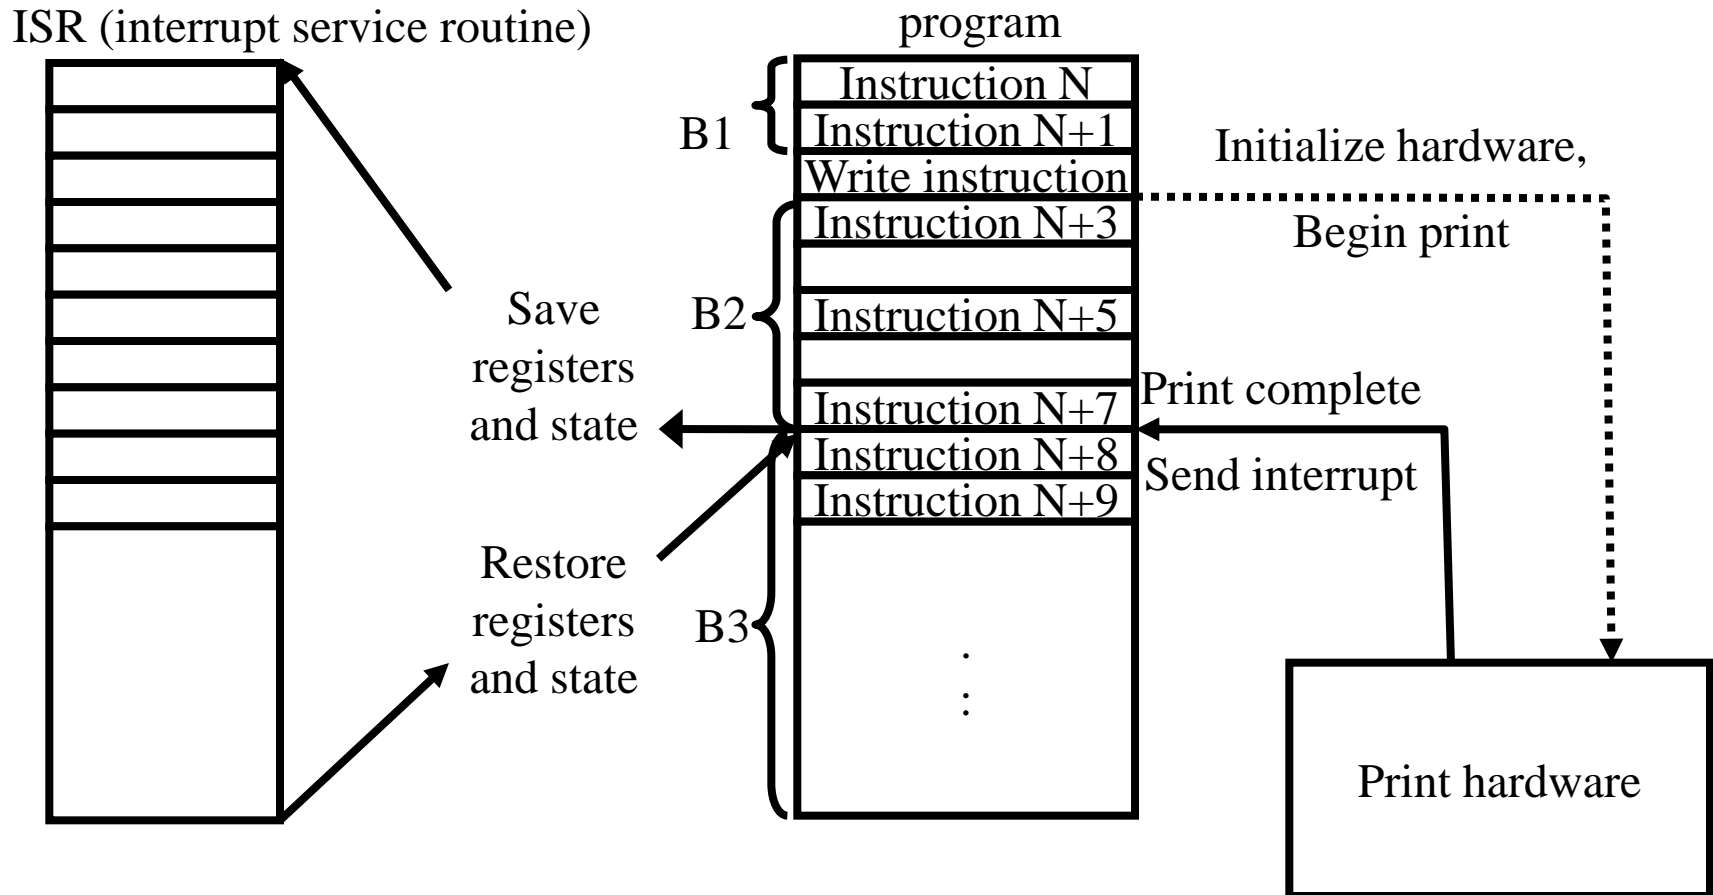


# Interrupt example Output (2)

---

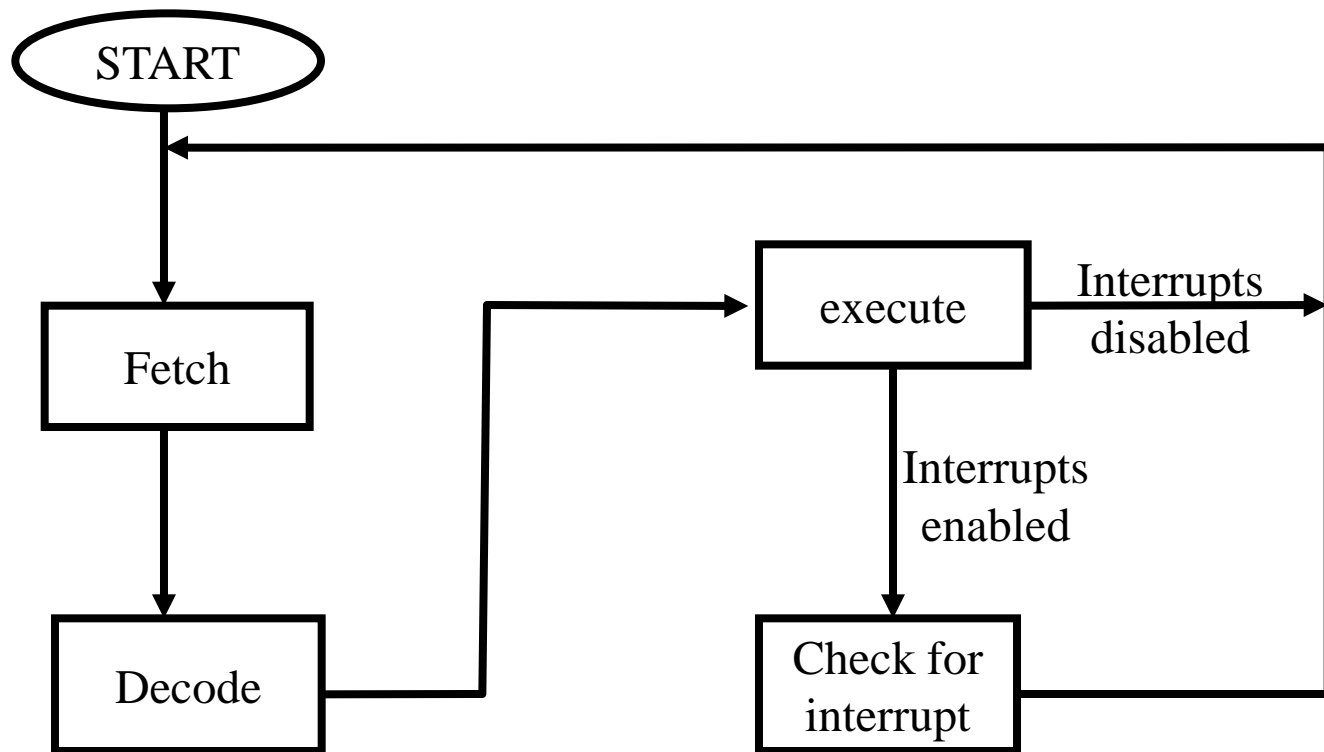
- ⌘ When the currently executing instruction completes, the program is interrupted
- ⌘ The program's registers and state are saved
- ⌘ An ISR (interrupt service routine) completes the output operation
- ⌘ Registers and state are restored
- ⌘ The original program continues executing

# Interrupt operation



# Instruction cycle with interrupts

---



# Interrupt processing (1)

---

- ⌘ A device issues an interrupt request
- ⌘ The CPU (processor) finishes execution of the present instruction
- ⌘ The CPU tests to see if there is a pending interrupt, determines there is, and sends a acknowledgement to the device requesting the interrupt
- ⌘ The CPU saves registers and state to the stack (including the program counter register value)
- ⌘ The CPU loads the address of the appropriate ISR into the address register

# Interrupt processing (2)

---

- ⌘ The CPU executes the ISR
- ⌘ When the ISR finishes, the saved register and state information is restored to the CPU registers
- ⌘ The program counter is reset to point to the next instruction
- ⌘ The original program continues execution

REMEMBER: the time when an interrupt occurs is not known in advance!! Interrupts are asynchronous

# **CMPT 300**

## **Introduction to Operating Systems**

---

Operating Systems  
Overview Part 2: History

# History of Operating Systems

- ⌘ First generation 1945 - 1955
  - 💧 vacuum tubes, plug boards
- ⌘ Second generation 1955 - 1965
  - 💧 transistors, batch systems
- ⌘ Third generation 1965 – 1980
  - 💧 ICs and multiprogramming
- ⌘ Fourth generation 1980 – present
  - 💧 personal computers

# The earliest computers (1945-55)

---

- ❖ Built of relays, vacuum tubes
- ❖ Very large, Very slow by today's standards
- ❖ Built, programmed and maintained by the same people
- ❖ Programmed by using switches, paper tape, etc
- ❖ No operating system, single operation, single problem, sequential access



# The next generation (1955-65)

---

- ❁ Transistor-based, increased reliability
- ❁ The first commercial mainframes, still very large and very expensive
- ❁ Used assembly or even early high level languages like Fortran or ALGOL
- ❁ Rudimentary operating system, one program at a time, with control commands to compile, load, execute, terminate, basic compilers
- ❁ Input using cards, paper tape, magnetic tape ...

# Single Job to Batch

---

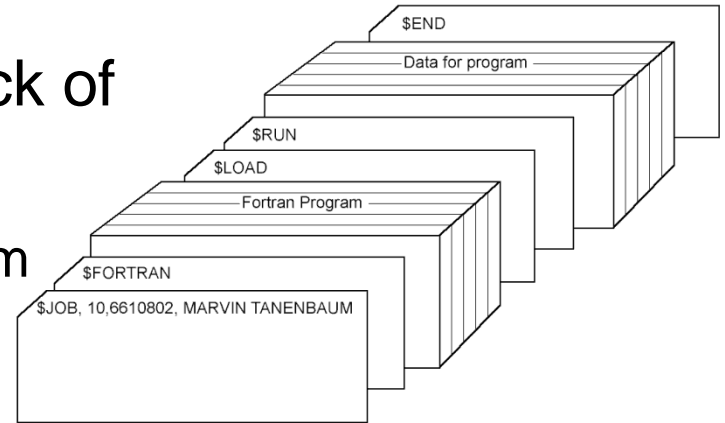
- ❁ Earliest machines had very rudimentary OS.
  - 💧 To run a job, needed to load the compiler as well as the code for the job.
  - 💧 Required a great deal of operator intervention
  - 💧 CPU not efficiently used
- ❁ Batch processing evolved to reduce the amount of time wasted setting up single jobs

# Early Batch processing

## ❁ Collect a group of jobs

💧 Each job was submitted as a stack of punched cards.

- Job card, language definition card
- One card per line of code in program
- Cards for load and run instructions
- Cards containing data for program
- End card indicating end of job




❁ A group of jobs was submitted as a batch to the card reader

❁ Each job was read in, executed, produced its output, terminated, then the next job took over the machine

# Early Batch Processing

---



# Operating System

---

- ❁ Commands to
  - 💧 Read a single card to memory
  - 💧 Compile to machine language
  - 💧 Place machine language code in memory
  - 💧 Start execution (load address of first instruction in program in Program counter then begin execution)
  - 💧 Write output to the line printer (or other output device)
  - 💧 Trap condition switches control from program to OS
    - END card being executed
    - Illegal opcode, divide by zero, ...

# Problems with early batch processing

- ❖ Input and output, particularly from peripheral I/O devices (card reader, line printer), are very slow when compared to the execution of other instructions
- ❖ When input and output is happening the CPU is mostly idle. An expensive resource is being wasted
- ❖ A program trapped in a infinite loop would never terminate

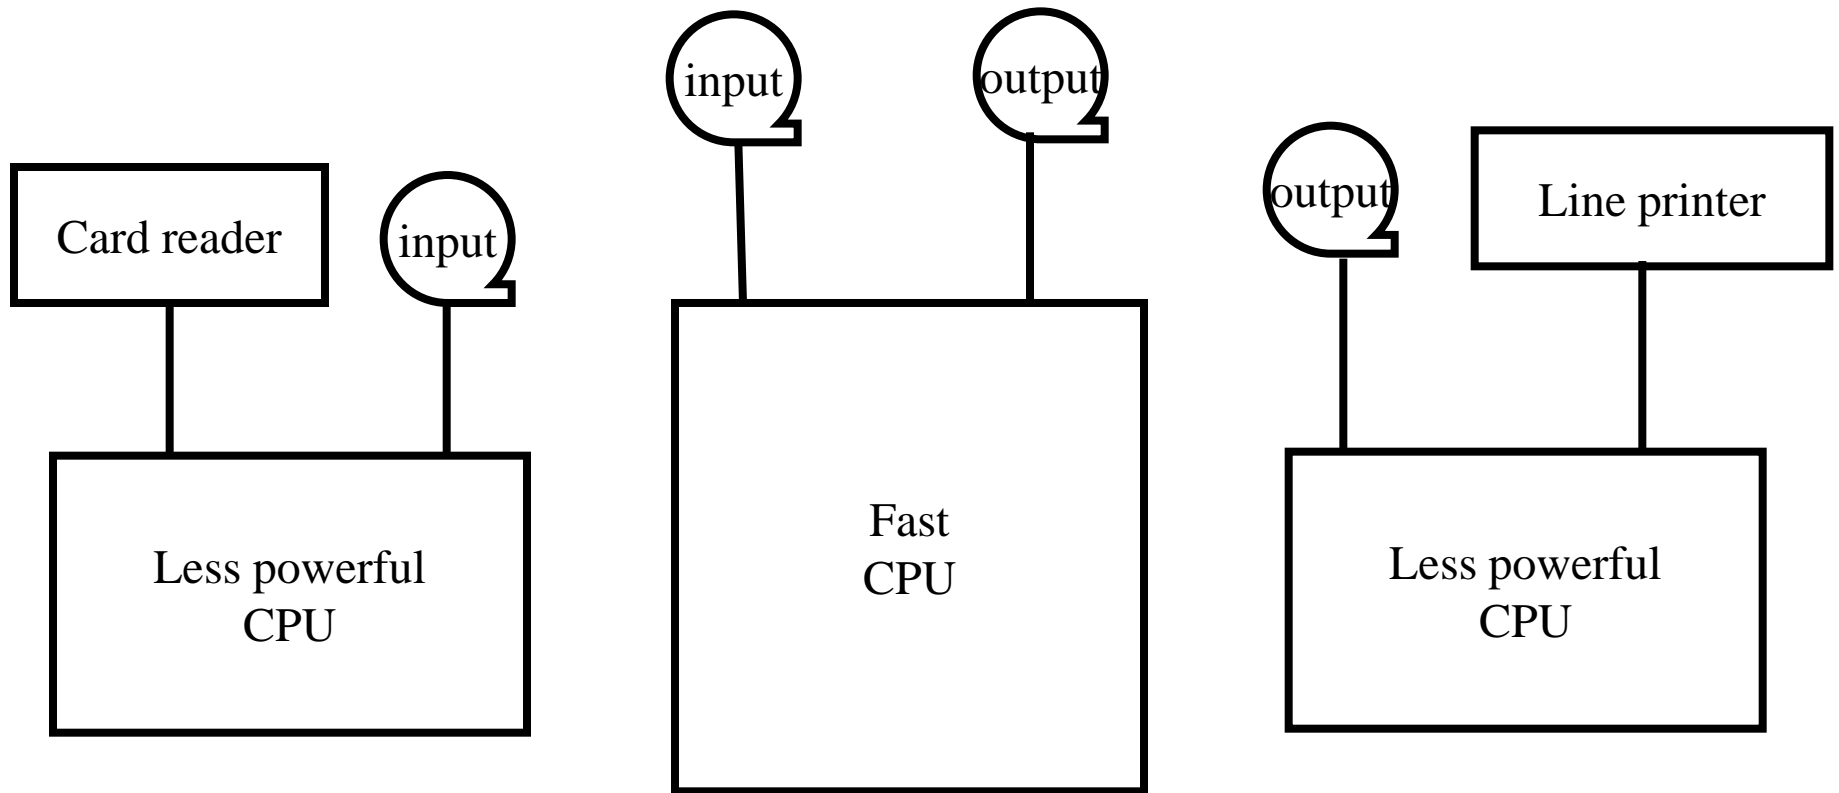
# Improving batch processing

---

- ❁ Offload the slow I/O tasks to less costly and powerful computers
- ❁ Use faster I/O media (like tapes) for input to fast powerful machine
- ❁ Add timers, if your time runs out an interrupt is generated which terminates your program (deals with infinite loops)
- ❁ Adds complexity, improves efficiency

# Improving batch processing

---






# The next generation (1965-1980)

---

- ❁ More complicated OS
  - 💧 Deal both with I/O intensive and CPU intensive jobs efficiently
  - 💧 Multi-programming (with partitioned memory)
  - 💧 Switches between tasks
    - Load and run jobs
    - Read cards to job queue on disk (whenever card reader is ready)
    - Print results to printer from printer queue on disk (only when printer is available and job is complete)
  - 💧 Changes enabled by going from tape to disk
    - Tape is sequential-access
    - Disk is random-access, hence faster for general workload

# The next generation (1965-1980)

---



# Simultaneous Peripheral Operation On Line (Spooling)


---

- ❁ **Spooling** refers to the process of placing data in a temporary working area for another program to process.
  - 💧 Send jobs from card reader to program queue on disk
  - 💧 Send job outputs into printer queue on disk
  - 💧 When card reader or printer are available, it can add to/print from queue
  - 💧 When a job is finished, the next job is loaded from the queue


# Multiprogramming

---

- ❁ Partition memory into pieces (often of different sizes)
- ❁ Load one job into each partition (choose partition according to needs of job)
- ❁ Have multiple jobs executing simultaneously, one in each partition
- ❁ When one job is doing I/O, another can be using the CPU



# Recall: Interrupt operation



# Interrupts for input

---

- ⌘ We have seen an example of using an interrupt to facilitate output
- ⌘ What about input? Can we do the same?
- ⌘ If we do the same we have problems
  - ◆ If we execute the code following the read, the value being read may be used in that code
  - ◆ If the value being read has not yet been placed in the variable, results will not be correct
- ⌘ How to solve the problem?
  - ◆ Instead of executing the next block of code in the same program we let a different program run until the I/O has completed.

# Multiprogramming:

## Interrupt example (1)

---

- ⌘ Program executes until it reaches a read instruction
- ⌘ The read routine sets up the input operation, then returns leaving the input hardware device (not the CPU) processing the input
- ⌘ Because successive instructions in the program may use the input value, the program itself cannot continue until the read is complete.
- ⌘ The read routine then interrupts program 1, saving its registers and state. Then the ISR loads the registers and state for program 2 and continues execution of its instructions (P to P+7 next slide)

# Multiprogramming:

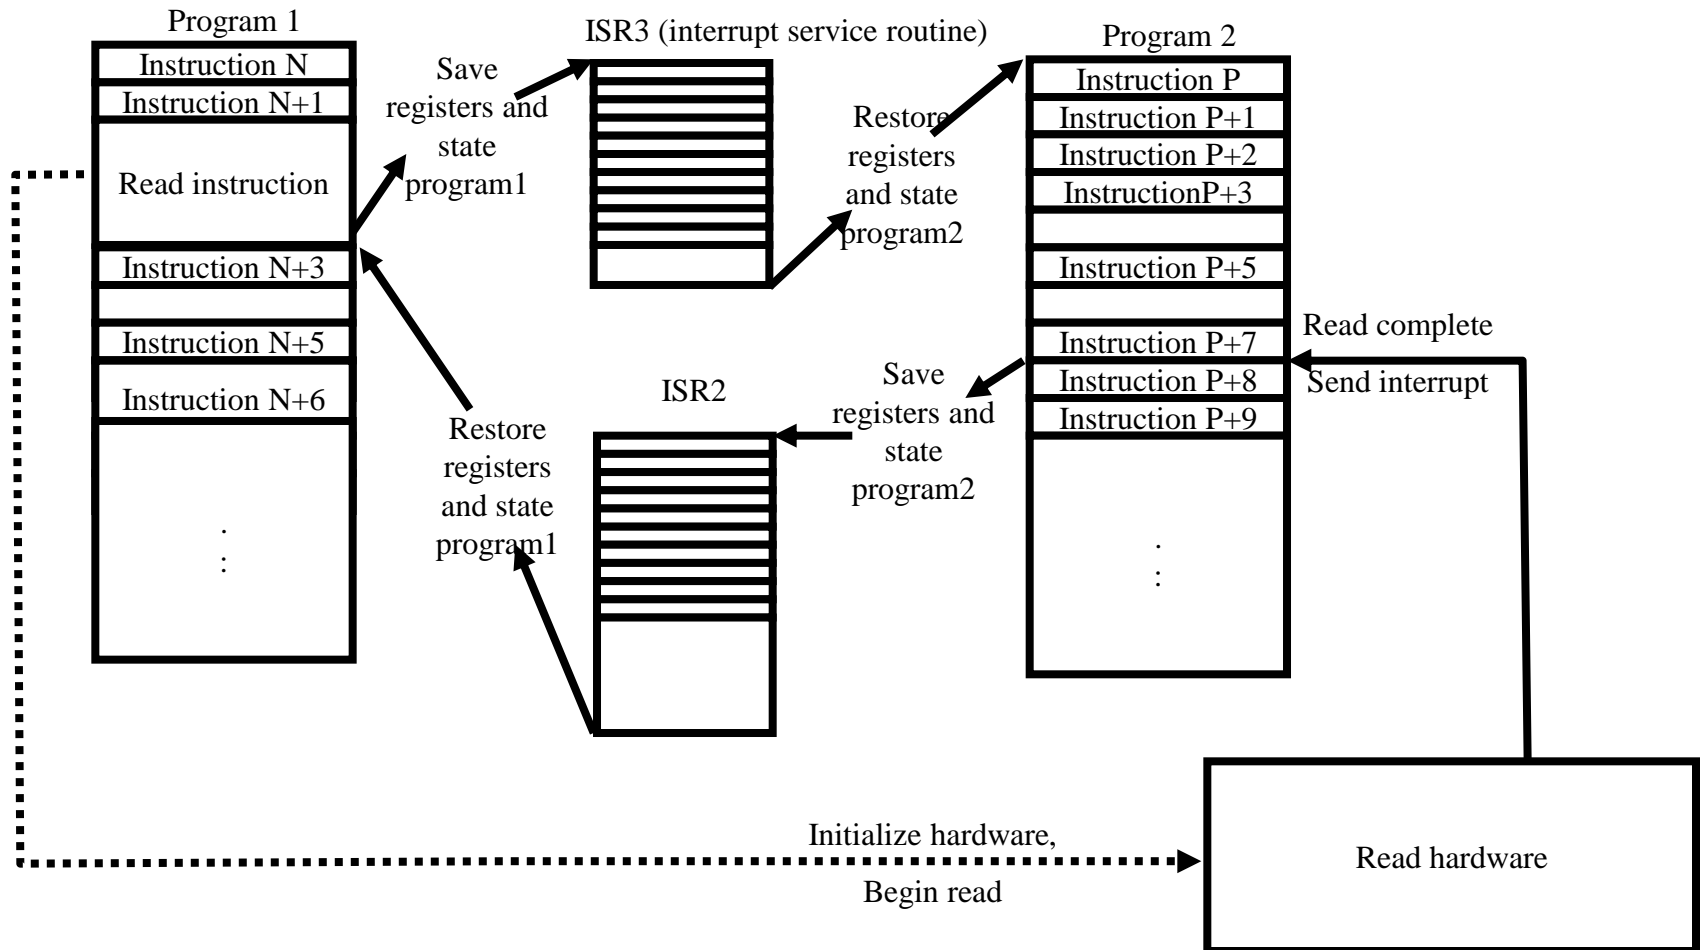
## Interrupt example (2)

---

- ⌘ When the input hardware device completes operation, it generates and sends an interrupt to the system (signaling normal completion)
- ⌘ When the currently-executing instruction in program 2 completes, program 2 is interrupted
- ⌘ The registers and state of program 2 are saved
- ⌘ An interrupt service routine completes the input operation
- ⌘ The registers and state of program 1 are restored
- ⌘ The original program continues executing




# Multiprogramming operation



# Multiprogramming example

---


Program A, uniprocessing



Followed by Program B, uniprocessing



multiprocessing



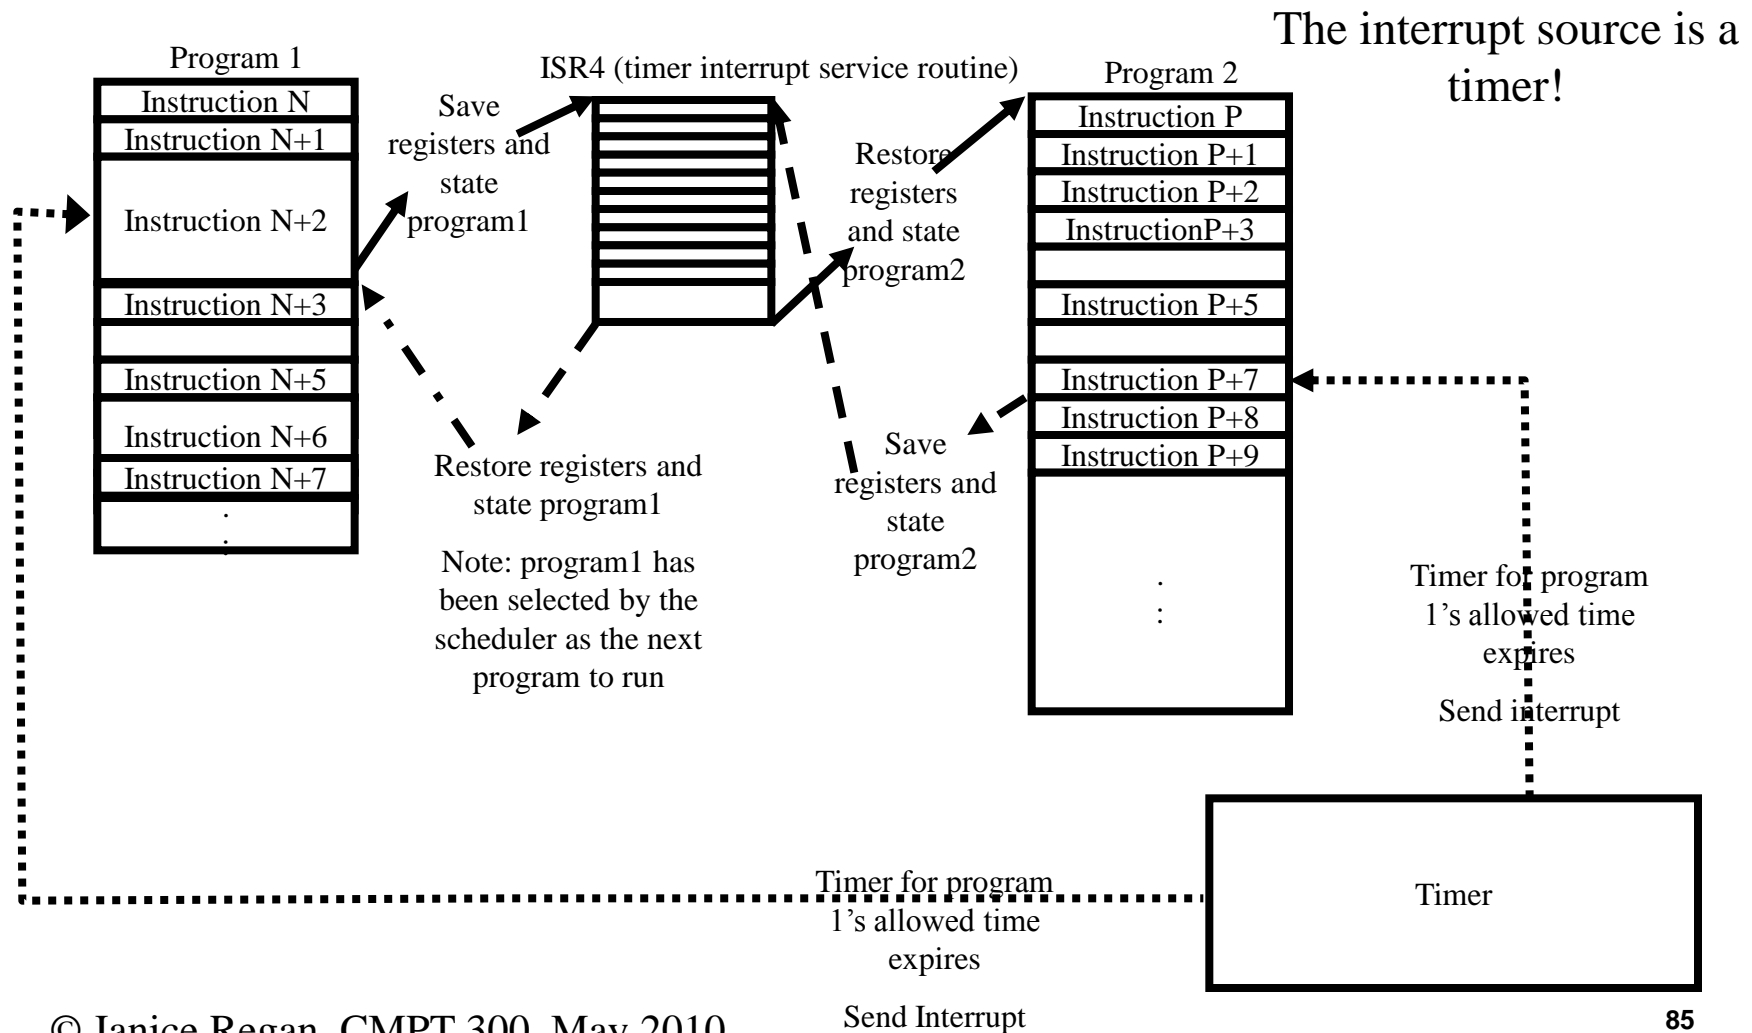
# Time sharing:

## Scheduling, fair sharing

---

- ❁ Once multiple jobs can share the CPU (sequentially, not at the same time) it becomes necessary to determine how time is shared between the processes
- ❁ The simplest approach to time sharing(taken by CTSS (Compatible Time-Sharing System) and some later OS's) is time-sliced round-robin
  - 💧 Each process is given N seconds of CPU, after N seconds the next processes takes its turn
- ❁ There are many other variants of scheduling, some of which we will discuss later.

# Simple time sharing operation



# Time sharing and scheduling

---

- ❁ In the previous example the simplest case (only two programs) was considered
- ❁ In a real system there will be many programs running.
  - 💧 Each time an ISR runs, at the end of the interrupt servicing, the OS scheduler must run to determine which program to give the CPU to next, and/or what to set the timer to

# Minicomputers

---

- ❁ Near the end of this generation (1965-1980) small, less expensive machines came into common use (for example the DEC PDP and VAX series)
- ❁ Costs were reduced from millions to 100's of thousands (about a factor of 20).
- ❁ Memory in Kbytes
- ❁ Used to develop UNIX operating system (multi-user)
- ❁ Problems as mini-computers proliferated: each vendor had their own flavor of UNIX (BSD, system 5, POSIX ...) or their own proprietary OS (VMS ... ), compatibility was an issue
- ❁ Mainframes and supercomputers were still necessary for computationally intensive applications.

# The next generation(1980- now)

---

- ❁ Use VLSI (very large scale integrated circuits) to build microcomputers
- ❁ Reduced price (thousands not 100's of thousands)
- ❁ First used early operating systems like CP/M (control program for microcomputers) or DOS (disk operating system)
- ❁ First uses of user friendly GUIs

# Commonly used OSs

---

- ❁ Versions of commonly used OSs like Windows, Unix, Linux are available for different types of platforms (PCs, Handhelds, Embedded systems)
- ❁ Specialized OSs for purposes such as real time embedded systems or large servers